

Ambient Occlusion

Pere-Pau Vázquez

ViRVIG – UPC

Outline

- Concept
- Derivation
- Algorithms
- Practical issues
- Implementation in mobile

Outline

- **Introduction**
- Derivation
- Algorithms
- Practical issues
- Implementation in mobile

Introduction

- <https://www.youtube.com/watch?v=1MuCGwoAH-U>

Introduction. Concept

- Ambient occlusion (AO) approximates global illumination
 - Determining how exposed a point on a surface is to ambient lighting
 - Relatively cheap approximation of rendering equation
 - Can be easily combined with other shading calculations (e.g. Phong shading, shadows...)


Introduction. Concept

- Global illumination
 - Light reaching to your eye (or surface) from any part of the scene (world)
 - Especially interesting on non-directly lit surfaces



Introduction. Concept

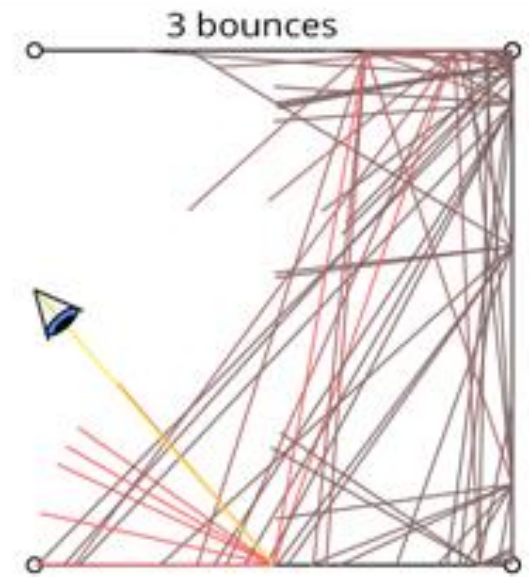

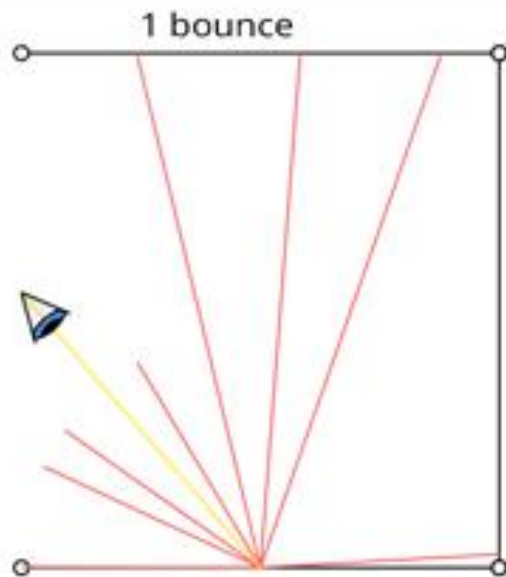
- Occlusion of incoming ambient light



Introduction. Concept

- Global illumination

- Result of multiple bounces of light across the objects in the scene



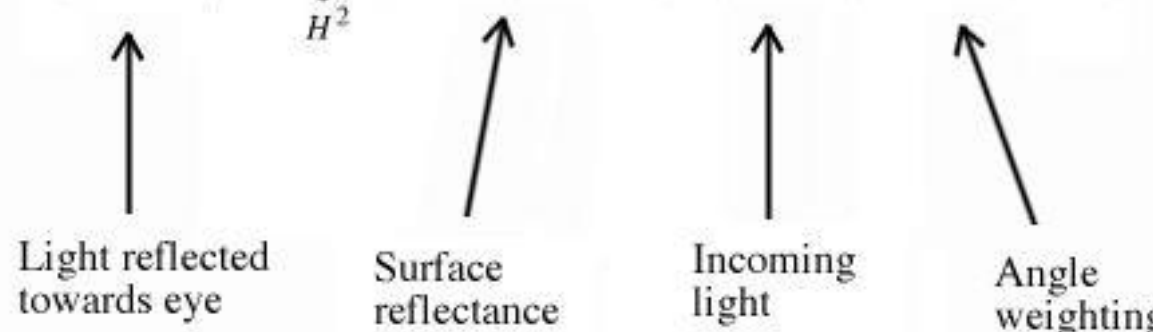
Introduction. Concept

- Ambient occlusion



Introduction. The Rendering Equation

- The interaction between light and objects is modelled by what is known as the rendering equation.
 - Describes how light reflects off a surface

$$L_r(x, \omega_r) = \int_{H^2} f_r(x, \omega_i \rightarrow \omega_r) L_i(x, \omega_i) \cos \theta'_i d\omega_i$$




The diagram illustrates the components of the rendering equation with arrows pointing from descriptive labels to specific parts of the equation:

- Light reflected towards eye**: An arrow points from this label to $L_r(x, \omega_r)$.
- Surface reflectance**: An arrow points from this label to $f_r(x, \omega_i \rightarrow \omega_r)$.
- Incoming light**: An arrow points from this label to $L_i(x, \omega_i)$.
- Angle weighting**: An arrow points from this label to $\cos \theta'_i$.

Introduction. The Rendering Equation

- The Rendering Equation

$$L_o(p, \omega_o) = \int_{\Omega} f(p, \omega_o, \omega_i) L_i(p, \omega_i) \cos \theta_i d\omega_i$$



Introduction. The Rendering Equation

- Every beam of light incident at p in direction ω_i
 - Multiplied by the BRDF of the material
 - And added up
 - Yields the reflected light at p in direction ω_o

$$L_o(p, \omega_o) = \int_{\Omega} f(p, \omega_o, \omega_i) L_i(p, \omega_i) \cos \theta_i d\omega_i$$

Introduction. The Rendering Equation

- Every beam of light incident at p in direction ω_i
 - May come from many different points in space
 - Light (which is not absorbed) continuously bounces off the surfaces in scene
- L_i is expressed as a function of L_o
 - Recursive formulation!!!

$$L_o(p, \omega_o) = \int_{\Omega} f(p, \omega_o, \omega_i) L_i(p, \omega_i) \cos \theta_i d\omega_i$$

Introduction. The Rendering Equation

- L_i is expressed as a function of L_o
 - Recursive formulation!!!

$$L_o(p, \omega_o) = \int_{\Omega} f(p, \omega_o, \omega_i) L_i(p, \omega_i) \cos \theta_i d\omega_i$$

- Means we need to apply recursivity for every ray, every bounce...

Outline


- *Introduction*
- **Derivation**
- Algorithms
- Practical issues
- Implementation in mobile

Derivation

- Many CG algorithms approximate the Rendering Equation
 - To accelerate its calculation
 - Otherwise it would be too costly


Derivation

- Direct lighting:
 - The simplest approach
 - Take into account a single bounce



Derivation

- 1 bounce vs 2 bounces vs 3 bounces
 - The larger the number of bounces the better the approximation




Derivation

- Ambient light
 - Replace indirect lighting with a constant
 - We assume that light reaches every point in every direction and with equal intensity
 - Ignoring the point's surrounding objects
- Traditional and computationally-efficient
- Rough approximation

Derivation

Ambient light



Derivation

- Ambient Occlusion
 - Indirect lighting approximation halfway between ambient light and true indirect illumination
 - Instead of using a constant ambient term, ambient occlusion produces an ambient term for every point in the scene

Derivation

- Ambient Occlusion makes two assumptions
 - The surface is a perfect diffuse surface
 - The BRDF becomes a constant

$$L_o(p, \omega_o) = k \int_{\Omega} L_i(p, \omega_i) \cos \theta_i d\omega_i$$

Derivation

- Ambient Occlusion makes two assumptions (ii)
 - Light potentially reaches a point p equally in all directions
 - But takes into account point's visibility

$$L_o(p, \omega_o) = k \int_{\Omega} V(p, \omega_i) \cos \theta_i d\omega_i$$

$$V(p, \omega_i) = \begin{cases} 0 & p \text{ occluded in direction } \omega_i \\ 1 & \text{otherwise} \end{cases}$$

Derivation

- Ambient Occlusion
 - Accounting for all visible directions
 - Upper part of the surface

$$L_o(p, \omega_o) = \frac{1}{\pi} \int_{\Omega} V(p, \omega_i) \cos \theta_i d\omega_i$$

Derivation

- Ambient Occlusion. Problem:
 - Indoor scenes would yield pitch black AO
 - Not the same occlusion comes from all the objects
- Solution:
 - Assume the closer objects occlude at a higher rate
 - Ambient Obscurance

Derivation

- Ambient Obscurance
 - Applies a fall-off function to the distance of intersection

$$L_o(p, \omega_o) = \frac{1}{\pi} \int_{\Omega} \rho(d(p, \omega_i)) \cos \theta_i d\omega_i$$

$$\rho(d) = \begin{cases} f(d) \in [0, 1] & d < \text{threshold} \\ 0 & \text{otherwise} \end{cases}$$

Derivation

- Ambient Obscurance
 - Fall-off function serves three purposes
 - Makes AO work for closed scenes
 - Cuts the occlusion to zero for far away points
 - AO relatively local computation
 - Gives nearby occluders greater occlusion value
- Ambient Occlusion and Ambient Obscurance commonly synonyms in literature

Outline

- *Introduction*
- *Derivation*
- **Algorithms**
- Practical issues
- Implementation in mobile

Algorithms


- Four main approaches
 - Ray-based
 - Geometry-based
 - Volume-based
 - Screen-space

Algorithms

- Four main approaches
 - **Ray-based**
 - Geometry-based
 - Volume-based
 - Screen-space

Algorithms. Ray-based

- Trace rays from the point to shade against the scene




Algorithms. Ray-based

- In practice
 - Use limited range for occlusion
 - Otherwise everything would be occluded in indoor scenes
 - Also faster to calculate because of finite radius
 - Use falloff function to smooth the transition
 - Do not solve analytically
 - Theoretically doable, but would be ridiculously expensive
 - Solution: Monte Carlo sampling
 - 256 – 1024 samples is usually enough


Algorithms. Ray-based

- Ray-based
 - Trace rays from the point to shade against the scene



Algorithms. Ray-based


- Cast a number of rays from the point to be shaded
- Determine occlusion distance, apply falloff, sum



- In reality use a fancy low-discrepancy sampling pattern

Algorithms. Ray-based


- Shadow rays
 - Usually faster than ordinary rays
 - Enough to detect any intersection, not necessarily closest one
 - Can do if we bake falloff function into the rays



Algorithms. Ray-based

- The amount of illumination a surface point is likely to receive in a scene, is calculated.
- A ray is sent from the camera to a surface point.
- This ray is split up into multiple rays that shoot off, into the scene, in a hemispherical fashion.
- These rays are biased to the normal of the surface.


Algorithms. Ray-based



Algorithms. Ray-based

- A calculation of the ratio of intersections and the number of cast rays is made, to produce the final illumination value.
- The less intersections we have, the brighter the point is.
- This is repeated again for the next surface point.

Algorithms. Ray-based

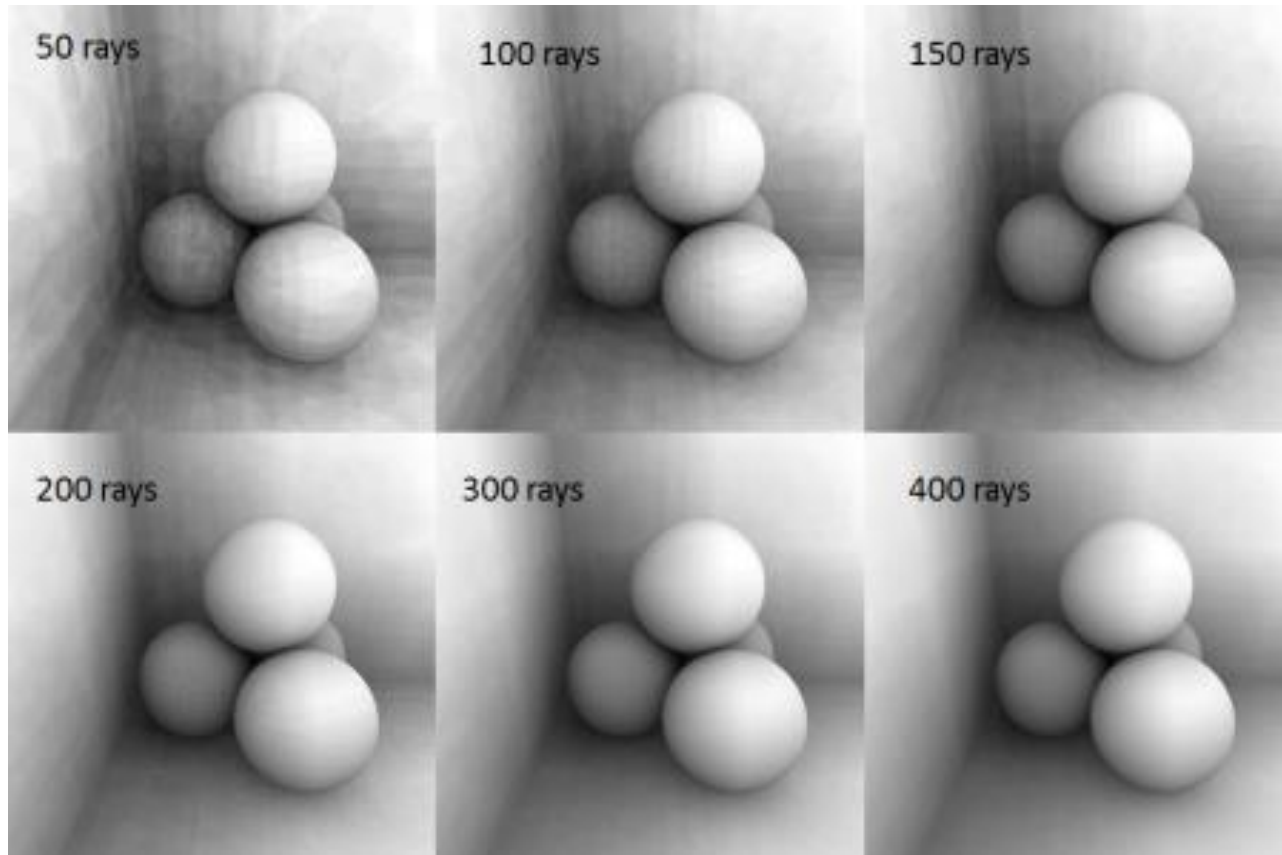


Algorithms. Ray-based

- Ray-based
 - High quality
 - Costly
 - Ray patterns
 - Sampling

Algorithms. Ray-based

- Ray-based



Algorithms


- Four main approaches
 - *Ray-based*
 - **Geometry-based**
 - Volume-based
 - Screen-space

Algorithms. Geometry-based

- RC is costly
 - Offline
- Simplify the query geometry
 - 2 layers: close geometry
 - Simplify objects


Algorithms. Geometry-based

- Simplify Geometry: Use polygon meshes as disk-shaped elements [Bunnell, 2005]
 - One element per vertex
 - Elements defined by position, normal, and area
 - Simplifies form factor calculation



Algorithms. Geometry-based

- Percentage of the hemisphere above a point occluded by an element (Solid Angle)
- Like radiosity form factor with 100% visibility



Emitter element E occludes receiver element R based on distance r and angles θ_E and θ_R

Algorithms. Geometry-based

- Calculate occlusion at a receiver by summing form factors:

Occlusion = 0

For each element E

Occlusion += form factor of E;

Algorithms. Geometry-based


- May simplify geometry further:




Algorithms. Geometry-based

- Must avoid double shadowing:

Ok



No Ok



Algorithms. Geometry-based

- Results



Environment Map




+ Ambient Occlusion



+ Indirect Lighting

Algorithms. Geometry-based




Algorithms. Geometry-based

- Other implementations:
 - Some proxy of geometry
 - Use multiple passes with shadow maps

Algorithms. Geometry-based

- Main idea: Approximating the geometry with spheres [Shanmugam and Arikan, 2007]:
 - Ray intersection is easy
 - Can be tricky to represent planar shapes




Algorithms. Geometry-based

- Two-levels:
 - High-frequency shadows – closer geometry
 - Querying the Normal-Depth (ND) map
 - Low-frequency shadows – distant objects
 - Sphere approximation with ray-length limitation


Algorithms. Geometry-based

- High-frequency shadows



Algorithms. Geometry-based

- Low-frequency shadows:
 - Approximating the geometry with spheres:



Algorithms. Geometry-based

- Geometry-based



Algorithms

- Four main approaches
 - *Ray-based*
 - *Geometry-based*
 - **Volume-based**
 - Screen-space

Algorithms. Volume-based

- Similar to geometry-based but
 - Scene is 3D
 - 3D texture encoding space occupancy
 - Voxels
 - Fragmented voxels

Algorithms. Volume-based

- Ambient Occlusion Volumes (McGuire 2010)
 - It computes the analog of a shadow volume for ambient light around each polygon
 - Applies a tunable occlusion function within the region it encloses

Algorithms. Volume-based


Ray Traced AO
5000 samples/pix

AOV (new)
1 sample/pix


Volumetric AO
64 samples/pix

Crytek SSAO
16 * 4x4 samples/pix


Sponza



City




House




Algorithms. Volume-based

- Ambient Occlusion Volumes (McGuire 2010)




Algorithms. Volume-based


- AO for large molecular models [Hermosilla et al., 2015]



Algorithms. Volume-based





Algorithms. Volume-based



Algorithms. Volume-based

- AO vs ground truth



Algorithms. Volume-based

- 1.36M atoms



Algorithms

- Four main approaches
 - *Ray-based*
 - *Geometry-based*
 - *Volume-based*
 - **Screen-space**

Algorithms. Screen-space

- Reduce calculation complexity
 - Use depth buffer to approximate scene geometry
 - Can also use Depth + Normal (aka ND buffer)
 - Calculation in screen-space → reduced complexity


Algorithms. Screen-space

- Uses depth buffer to approximate scene geometry



Algorithms. Screen-space

- Screen-space post-processing effect
 - Z-buffer data might be already available in a texture
 - Use it to estimate AO




Algorithms. Screen-space

- AO estimation
 - Sample neighbouring pixels
 - The calculated term is used to attenuate incoming light

Algorithms. Screen-space

- Use random samples inside a sphere centered at a surface point
 - Prevents banding
- Occlusion function used to relate sample depth delta and distance from the central point
 - Negative deltas give zero occlusion
 - Small positive depth delta produces high occlusion term
 - Large positive depth delta tend to zero
 - Because this calculation happens in screen space
 - Simple exponentials or look-ups used

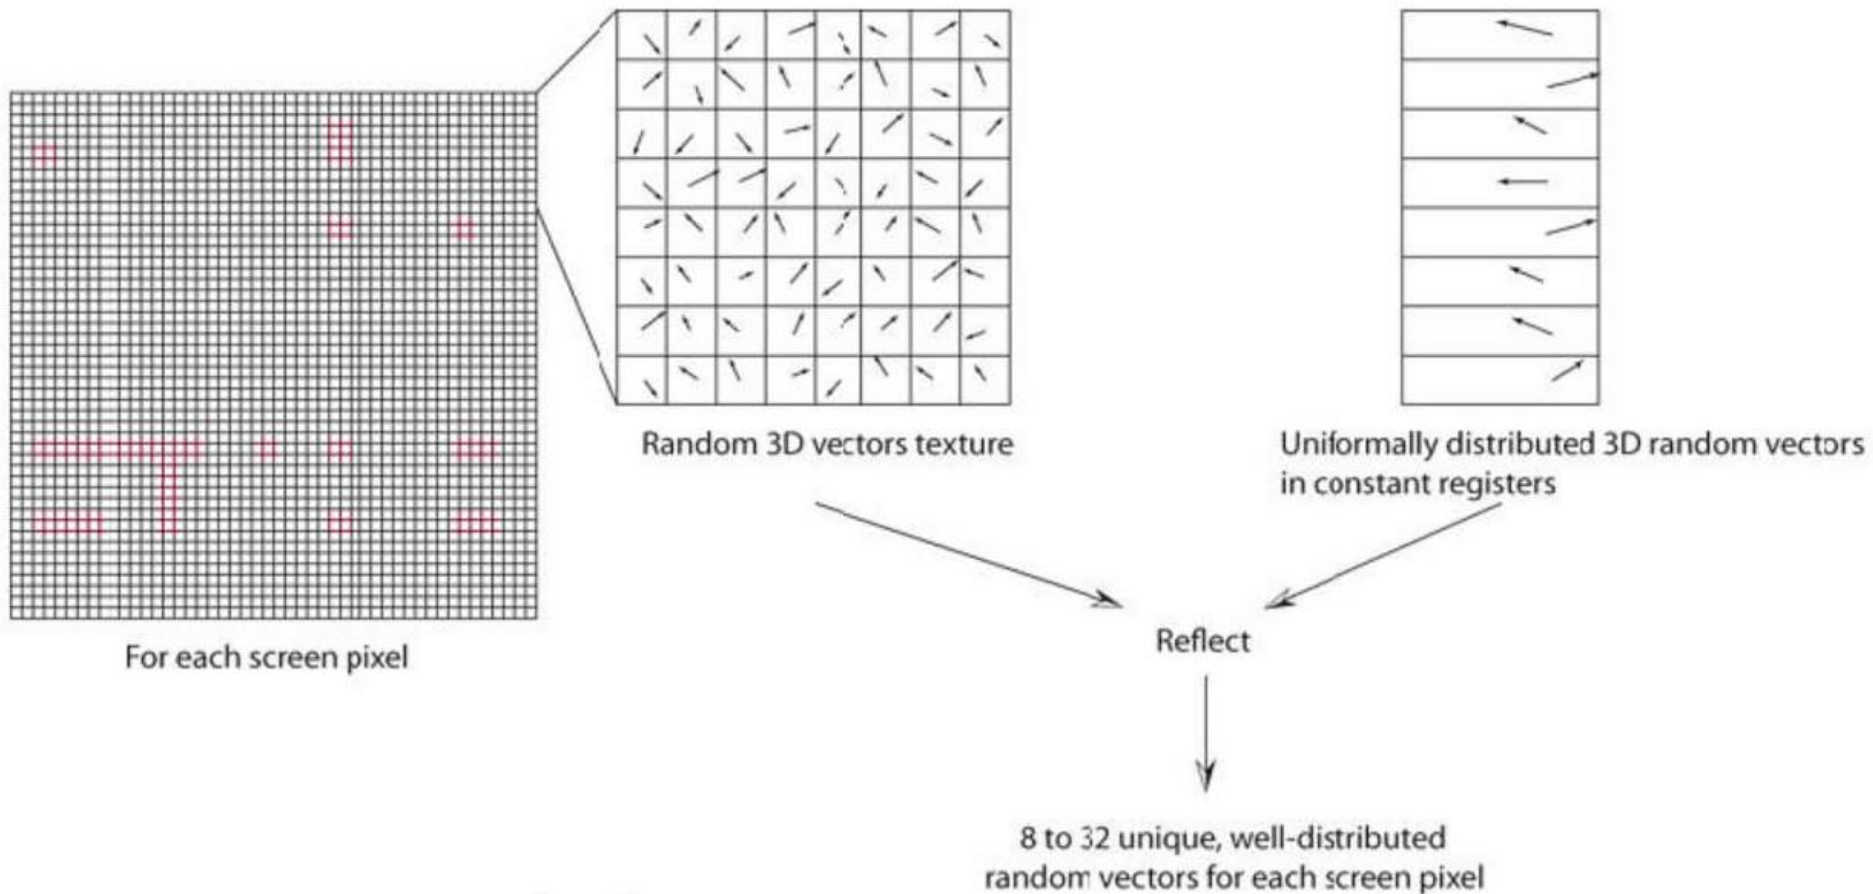
Algorithms. Screen-space



Algorithms. Screen-space


- Calculation:
 - Generate some number of random normal vectors per pixel:
 - 8-32: Starcraft II, 16: Cryengine II
 - Reflect random vectors off another set of varying length vectors with uniform distribution in solid sphere
 - Range of length is scaled by some artistic parameter
 - Samples then passed through the occlusion function

Algorithms. Screen-space




Algorithms. Screen-space

- The result is a noisy image




Algorithms. Screen-space

- The result is a noisy image
 - Solution: Blur the resulting image



Algorithms. Screen-space

- Self occlusion:
 - A sample may occlude itself
 - Flip offset vectors



Algorithms. Screen-space

- Issues:
 - No sample outside a certain target:
 - No shadow
 - Camera close to an object: Noise more noticeable
 - Increase the number of samples based on view proximity
 - Constrain the area in which the samples are taken

Algorithms. Screen-space

- Performance:
 - Stable, since it is a SS technique
 - Bottleneck at the random sampling
 - Tends to over illuminate solid object edges
 - Lower screen resolution for depth buffer often sufficient
 - $\frac{1}{4}$ size of the original depth render
 - Multiple SSAO functions can be used together (with different sampling constraints) to model different effects
 - Take the largest occlusion of all occlusion averages


Algorithms. Screen-space

- Horizon-Based AO [Bavoil et al., 2008]
 - For some radius around sample point:
 - Step through depth buffer in some number of randomized directions for a fixed number of samples
 - Find highest altitudes from center within radius (horizon)
 - Average the weighted samples over the directions
 - Intergrade radiance over through visible angle

Algorithms. Screen-space

- HBAO:

- Step through depth buffer in some number of randomized directions for a fixed number of samples
- Find highest altitudes from center within radius (horizon)




Algorithms. Screen-space

- AO vs HBAO

Half-Resolution AO
6x6 samples / AO pixel
No Blur

AO = 3.5 ms @ 800x600
On GeForce GTX 280



Algorithms. Screen-space

- AO vs HBAO

Half-Resolution AO
6x6 samples / AO pixel
15x15 Blur

AO = 3.5 ms @ 800x600
Blur = 2.5 ms @ 1600x1200
On GeForce GTX 280



Algorithms. Screen-space

- HBAO
 - Interactive frame rates (2008)
 - Also relies on blurring
 - Might require some biasing in horizon angle
 - Per sample falloff (radial falloff function)
 - Jitter samples
 - More accurately simulates ray casting AO

Algorithms. Screen-space

- Battlefield 3.
 - No SSAO



Algorithms. Screen-space

- Battlefield 3.
 - SSAO



Algorithms. Screen-space

- Battlefield 3.
 - HBAO



Algorithms. Screen-space

- Battlefield 3.
 - SSAO



Algorithms. Screen-space

- Battlefield 3. HBAO

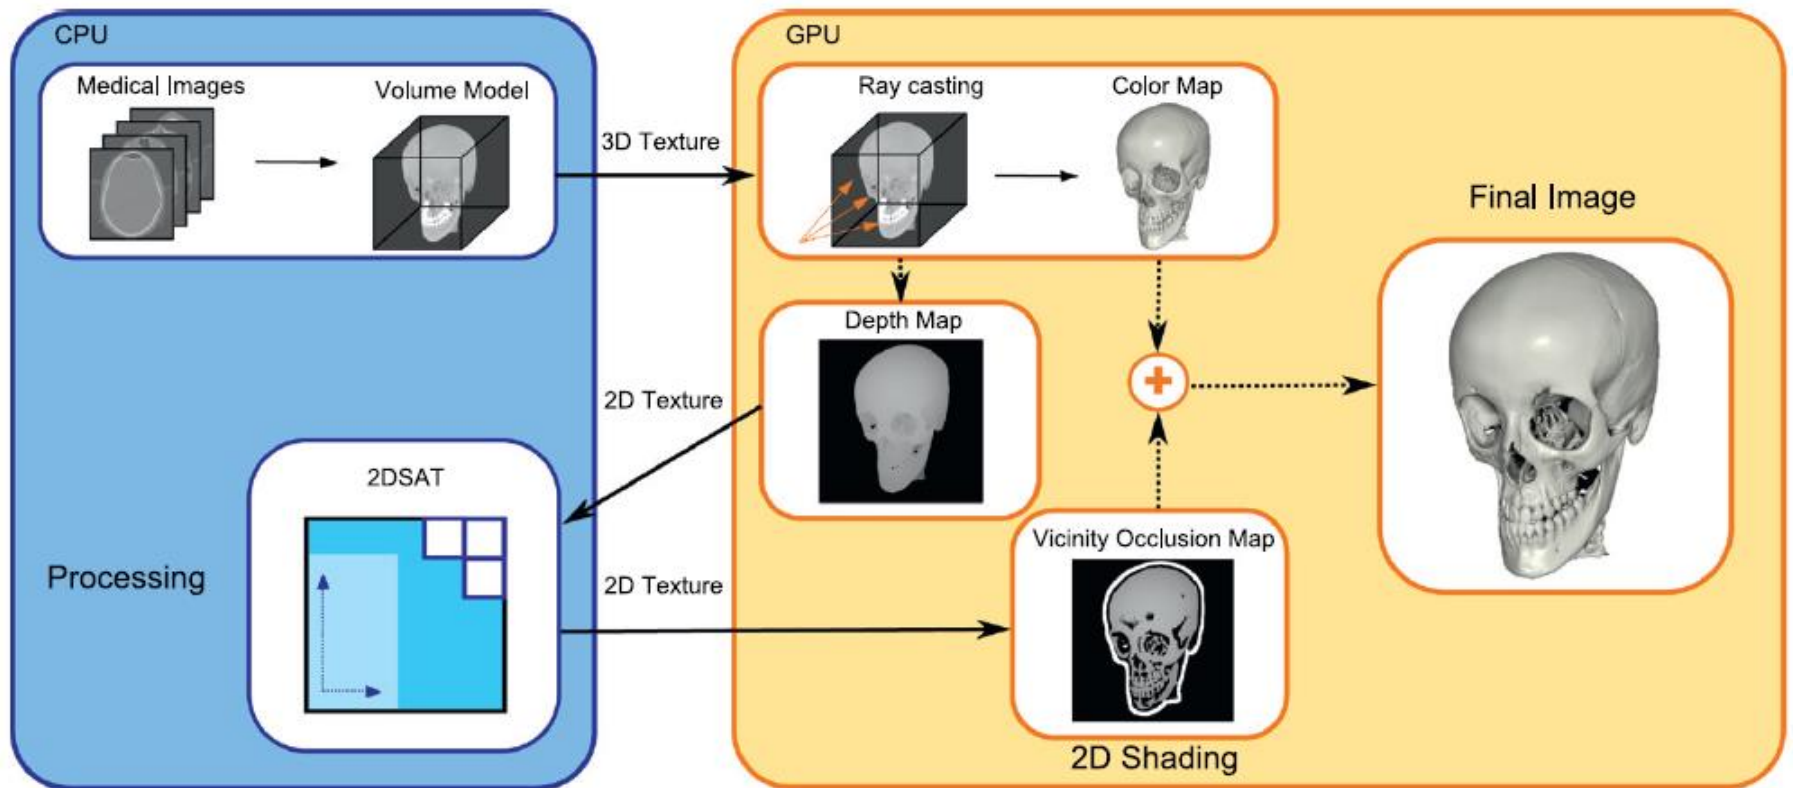


Algorithms. Screen-space

- Vicinity Occlusion Maps: AO for volumetric models [Díaz et al., 2008]
 - Approximates ambient occlusion screen-based
 - Uses a data structure based on Summed Area Tables:
 - Reduces sampling cost
 - Provides the blur
 - Halos for free

Algorithms. Screen-space

- Vicinity Occlusion Maps



Algorithms. Screen-space

- Vicinity Occlusion Maps



Algorithms. Screen-space

- Vicinity Occlusion Maps

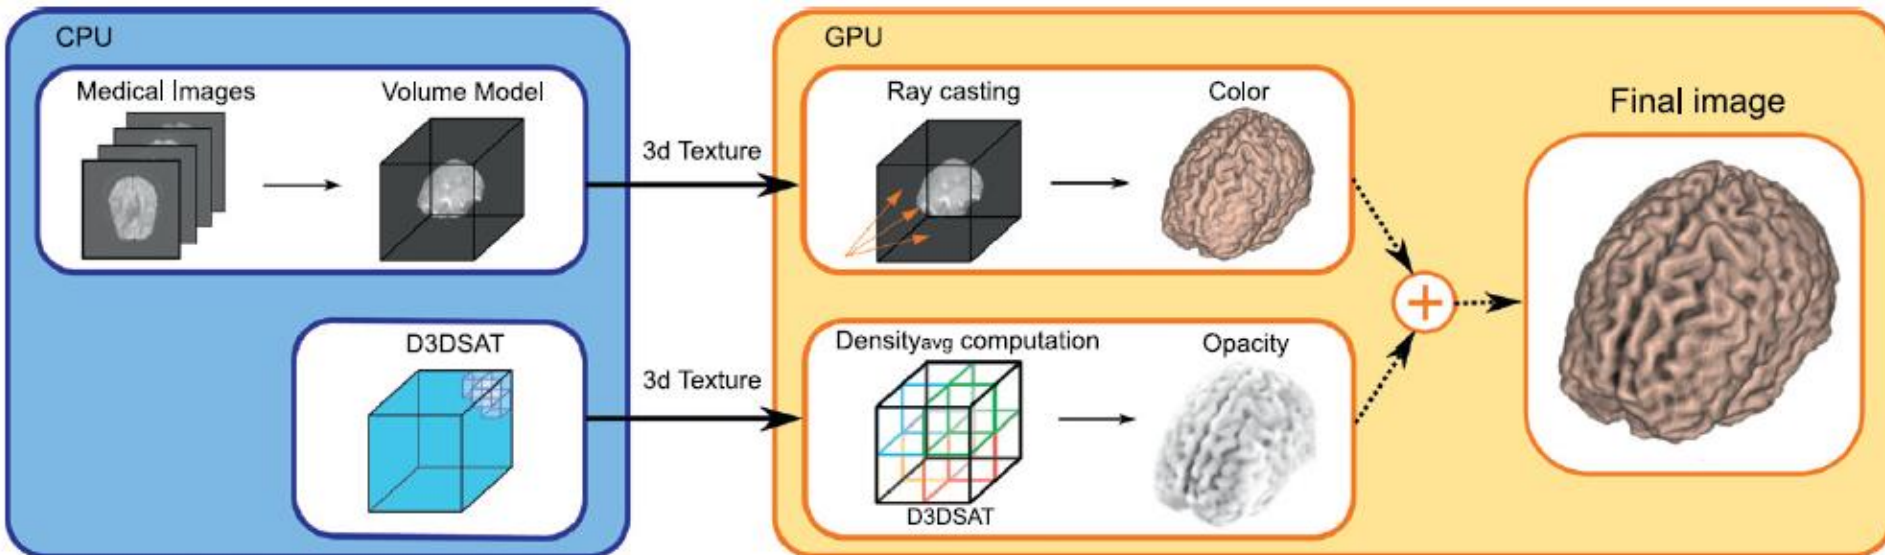


Algorithms. Screen-space




Algorithms. Screen-space

- DSAT3D: View independent AO for volumes [Díaz et al., 2010]
 - Actually more a Volume-based method
 - Uses a 3D SAT for view independent AO calculation



Algorithms. Screen-space

- DSAT3D




Ray casting



DSAT3D ambient occlusion

Algorithms. Screen-space

- DSAT3D
 - Combine with 2D VOM



Algorithms. Screen-space

- Conclusions:
 - SSAO methods avoid pre-computation
 - Can be used in dynamic scenes
 - Different techniques concentrate on different aspects of the rendering equation
 - Some techniques may cause (or generate) color bleeding

Algorithms. Screen-space

- Other approaches:
 - Reflective shadow maps
 - Directional occlusion
 - ...

Outline

- *Introduction*
- *Derivation*
- *Algorithms*
- **Practical issues**
- Implementation in mobile


Practical issues

- Usually sampling-limited
 - May reduce resolution for depth map
- Distant geometry is tricky
 - Smart fall-off function may help
- Trade noise vs banding

Practical issues




- Banding & Noise

- A fixed sampling pattern produces banding (left).
- Random sampling removes banding but introduces noise (middle).
- SSAO output is typically blurred to remove noise (right).



Practical issues

- Sampling
 - Desktop implementations may take 16-32 samples per pixel




Practical issues

- Reducing resolution
 - AO is usually low frequency, so SSAO/HBAO is performed at half resolution for significant speedup
 - Smart blur is done using full resolution z-buffer to avoid edge bleeding
 - Small high frequency geometry like grass can cause flickering / shimmering due to the half resolution

Practical issues

- Blurring: Typical Gaussian blurring smoothens out edges
 - Use bilateral filtering



Practical issues

- Blurring. Bilateral filtering

$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_q - I_p|) I_q$$

$$G_{\sigma}(x) = \frac{1}{\sigma \sqrt{2\pi}} \exp \left(-\frac{x^2}{2\sigma^2} \right)$$

Practical issues


- Solution in Battlefield 3: Temporal filtering
 - AO is dependent only on scene geometry, not camera position
 - Use AO from last frame
 - Reproject it to the current view and interpolate between new and old AO

Outline

- *Introduction*
- *Derivation*
- *Algorithms*
- *Practical issues*
- **Implementation in mobile**


Implementation in mobile

- SSAO: Alchemy-based in OpenGL ES



Implementation in mobile

- Overview of the approach:
 - Sample in a disk (texture space)
 - Unproject to get sample
 - Apply falloff function to map angle (n , sample- p) to AO in $[0,1]$
 - Use p and s normal to prevent self-occlusion
 - Penalise large depth discontinuities



Implementation in mobile

- Implementation

```
#define NSAMPLES 4
#define NSAMPLESf 4.0

precision mediump float;

uniform sampler2D depth;
uniform sampler2D normal;
uniform sampler2D rotation;
uniform mat4 iProjection;

uniform vec2 samples[NSAMPLES];
uniform float radius;
uniform vec2 ssaoSize;
uniform float rotationWidth;
uniform vec2 texelSize;

varying vec2 texCoord;
```


Implementation in mobile

- Implementation

```
varying vec2 texCoord;

// Unproject the given point in texture coordinates to view space
// coordinates.
// (s,t, linear depth) -> (x,y,z)
vec3 unproject (vec2 st, float d)
{
    vec2 xy = st*2.0 - 1.0; // x/w and y/w in [-1,1]
    vec4 pfar = (iProjection * vec4(xy, 1.0, 1.0));
    pfar.xyz /= pfar.w;
    return (pfar * d).xyz;
}
```

Implementation in mobile

- Implementation

```
void main ()
{
    vec2 frag      = texCoord; // frag in [0,1]
    float pdepth   = texture2D(depth, frag).r;
    float pradius   = radius * (1.0 - pdepth);
    vec3 pnormal    = texture2D(normal, frag).rgb*2.0 - 1.0;
    vec3 pview      = unproject(frag, pdepth);

    vec2 rotst     = frag * ssaoSize / rotationWidth;
    vec2 rotvec     = texture2D(rotation, rotst).xy*2.0 - 1.0;

    // Alchemy AO
    float AO = 0.0;
    for (int i = 0; i < NSAMPLES; ++i)
    {
```

```

float AO = 0.0;
for (int i = 0; i < NSAMPLES; ++i)
{
    vec2 svec = pradius * reflect(samples[i].xy, rotvec);
    vec2 qfrag = frag + svec;
    float qdepth = texture2D(depth, qfrag).r;
    vec3 qview = unproject(qfrag, qdepth);
    vec3 qnormal = texture2D(normal, qfrag).rgb*2.0 - 1.0;

    // Compute occlusion based on angle
    vec3 v = normalize(qview - pview);
    float a = smoothstep(0.0, 1.0, dot(v, pnormal));

    // Avoid self-shadowing
    float b = 1.0 - dot(pnormal, qnormal);

    float diff = qdepth - pdepth;

    // Penalise large depth discontinuities
    float dbias = 0.05;
    float c = step(abs(diff), dbias);

    AO += a * b * c;
}
AO = 1.0 - AO/NSAMPLESf;

gl_FragColor.r = AO;

```

Ambient Occlusion

Pere-Pau Vázquez

ViRVIG – UPC