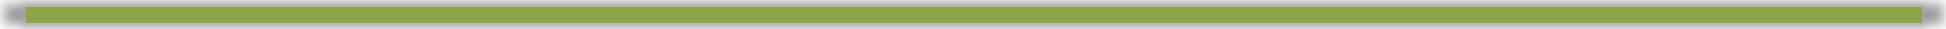


# Hard shadows

---

Pere-Pau Vázquez  
ViRVIG – UPC

# Outline



- Motivation
- Shadow Maps
- Shadow Volumes

# Motivation

---

- Why do we need shadows?
  - How High is the tennis player?



Without Shadow



3 inches?



1 foot?



3 feet?

# Motivation

---

- Why do we need shadows?
  - Shadows are popular in games!



# Motivation

---

- Shadows do significantly contribute to the realism of rendered images
- Global effect which is expensive to compute
- Lot of work done in this area
- Which algorithms are suitable for interactive applications ?

# Motivation

---

- Shadows play an important role in our understanding of 3D geometry:  
Help to understand relative object position and size in a scene



# Motivation

---

- Shadows can also help us understanding the geometry of a complex receiver



# Motivation

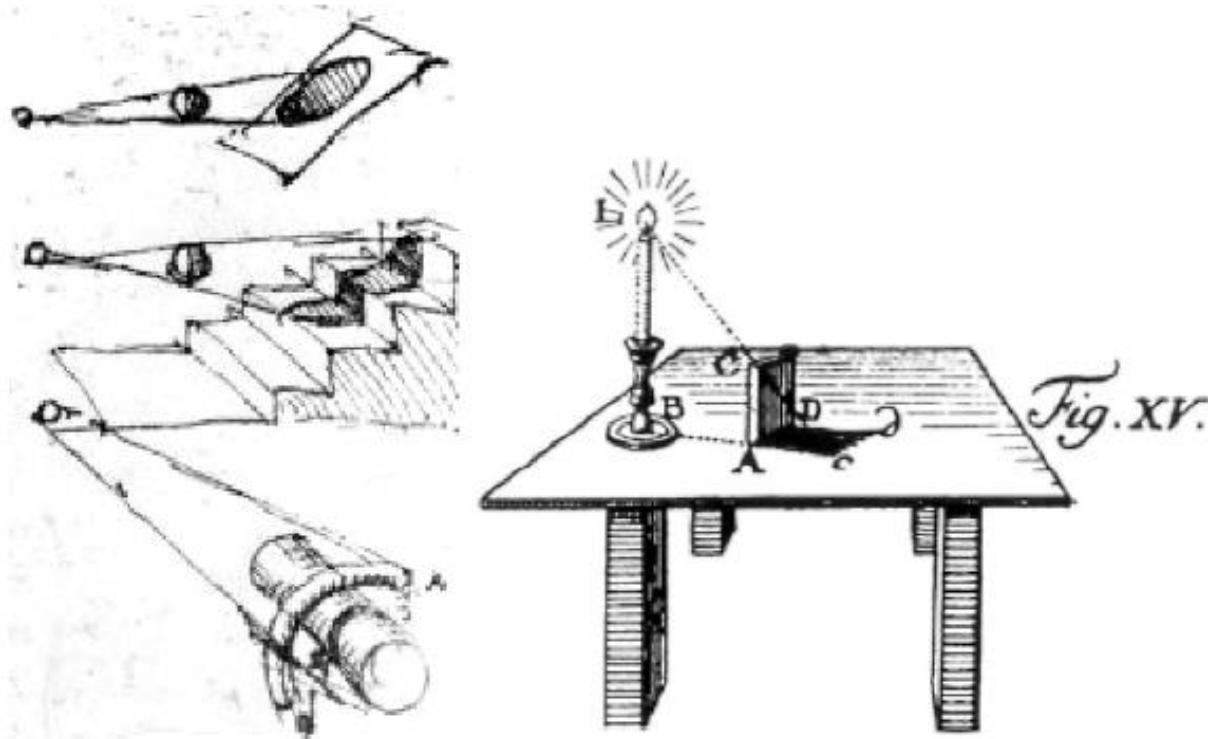
---

- Finally, shadows provide useful visual cues that help in understanding the geometry of a complex occluder



# Motivation

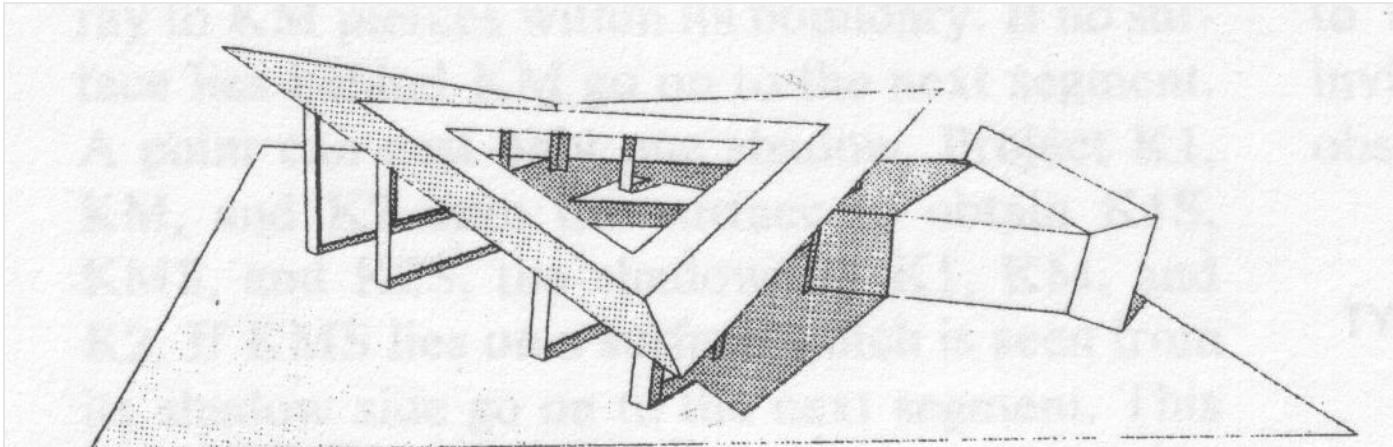
---



**Figure 1:** Left: Study of shadows by Leonardo da Vinci<sup>48</sup> — Right: Shadow construction by Lambert<sup>35</sup>.

# Motivation

- Efficient methods for computing shadows are still challenging task in computer graphics



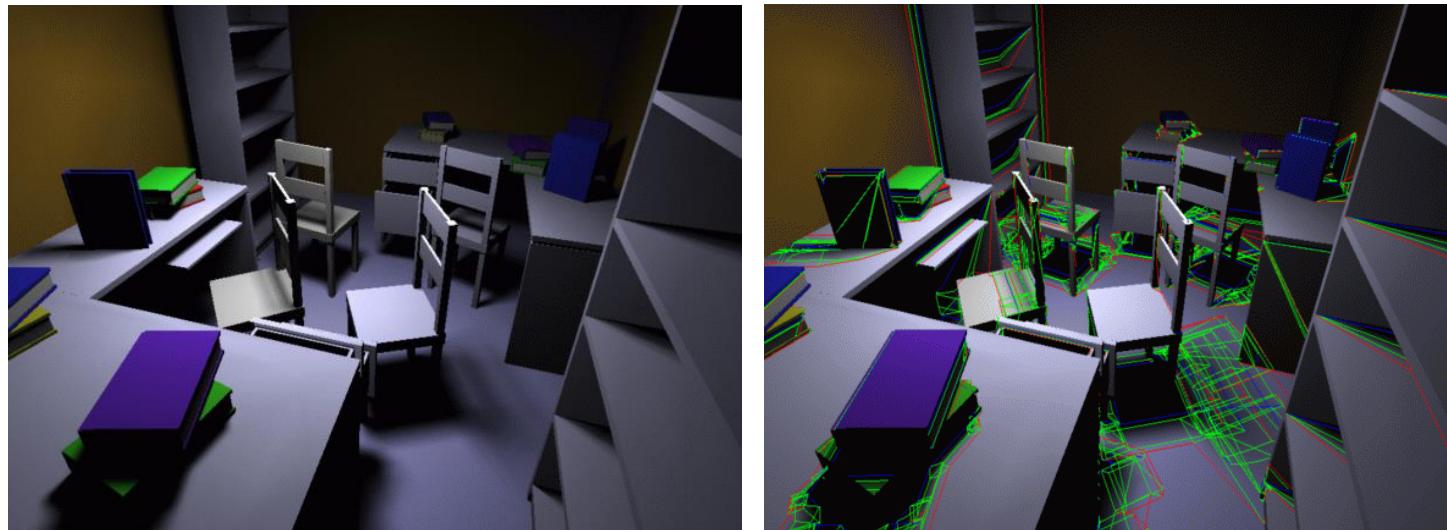
**Figure 5 – A higher angle view of the building. 7094 calculation time for this picture was about 30 minutes.**

- Early work on shadow techniques  
[Arthur Appel, IBM Research, 1968]

# Motivation

---

- Many algorithms proposed to compute shadows



- But only a few techniques are suitable for real-time or interactive applications

# Motivation

---

- Refresh rate ~ 30 to 60 fps
  - Games
  - VR Simulations
  - Virtual TV Set

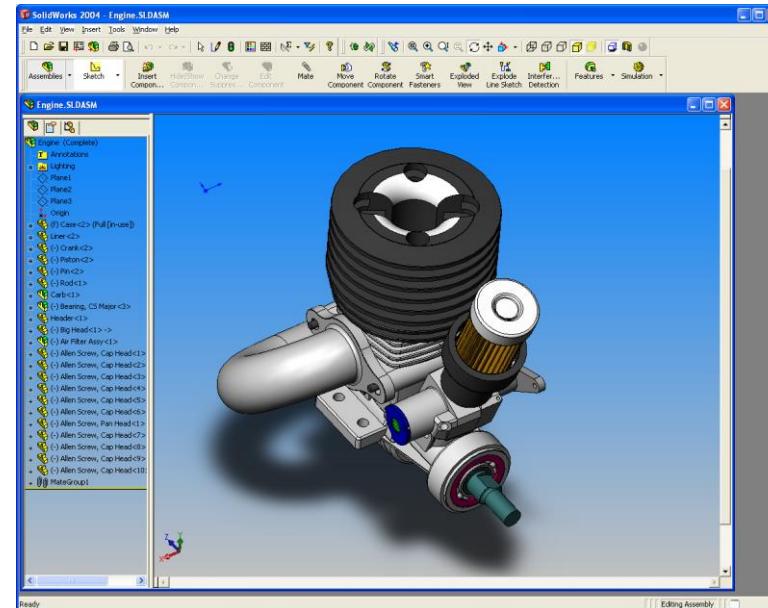


*Doom3*

- Only hardware-accelerated methods
  - CPU resources dedicated for other tasks
    - Sound, physics, AI, input

# Motivation

- Refresh rate ~ 10 to 30 fps
  - CAD/CAM
  - Digital Content Creation (DCC)
  - VR Interaction
- Hybrid approaches possible
  - Combined CPU & GPU shadow technique

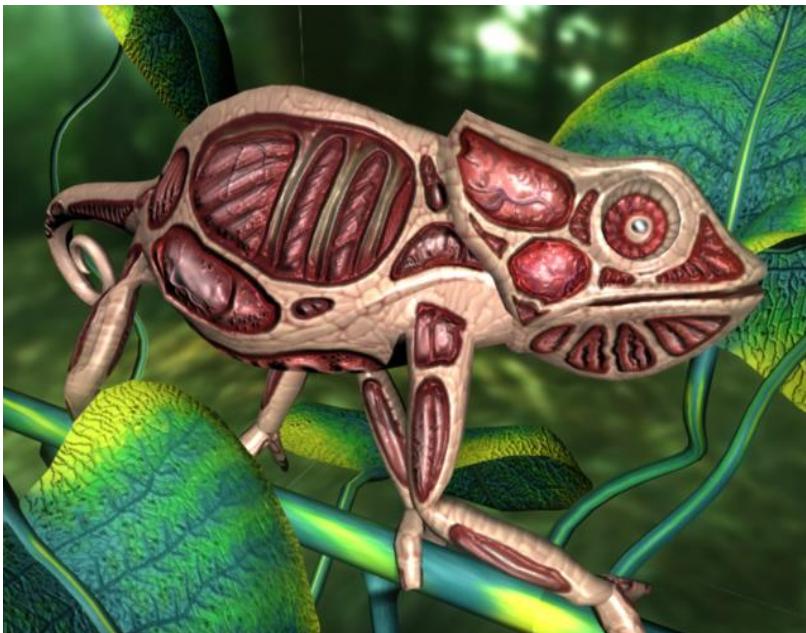


*SolidWorks*

# Motivation

---

- New group of applications
  - Hardware-accelerated rendering methods for offline rendering (preview, production)
    - Example: Alias Maya6 hardware-renderer



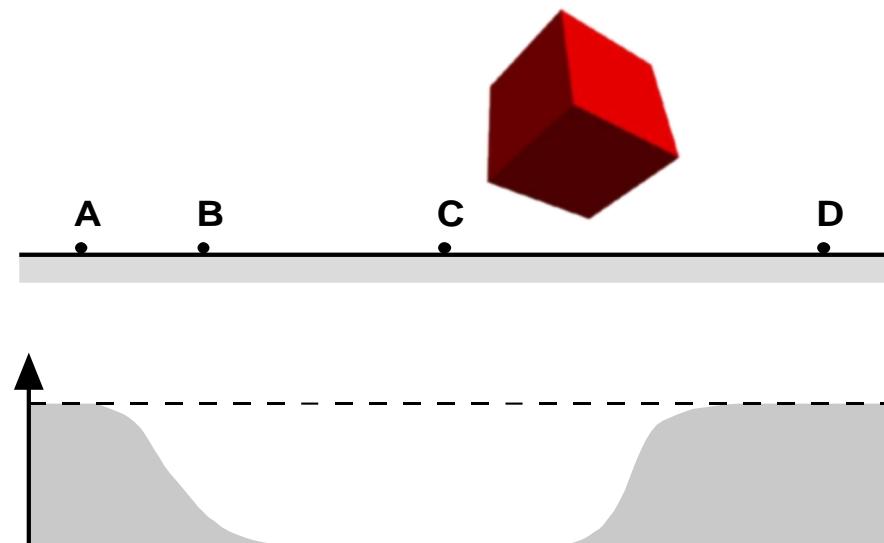
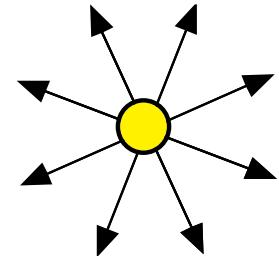
*Maya Software*



*Maya OpenGL*

# Motivation

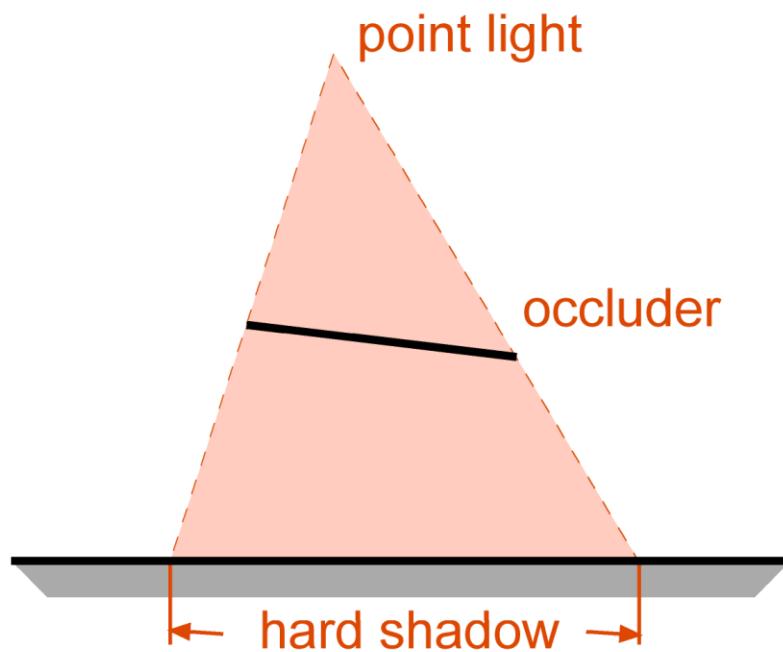
- To compute shadows we need to analyze the spatial relationship of
  - Receiver
  - Light source
  - Occluders
- Result: intensity distribution on receiver



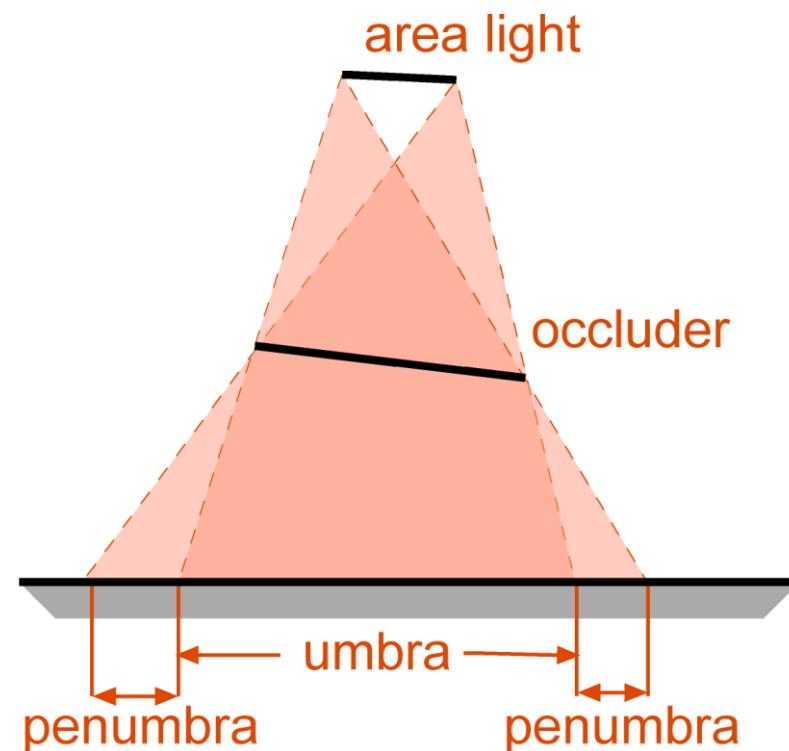
# Motivation

---

- Hard shadows vs. Soft shadows



Hard shadow edges  
- unreal (like geometry)  
+ easy computation



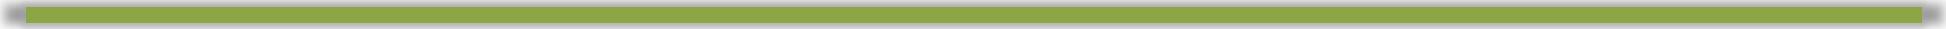
Soft shadow boundaries  
+ realistic  
- expensive

# Motivation

---

- Ways of thinking about shadows
  - As volumes of space that are dark
    - Shadow Volumes [Franklin Crow 77]
  - As places not seen by a light source looking at the scene
    - Shadow Maps [Lance Williams 78]

# Outline

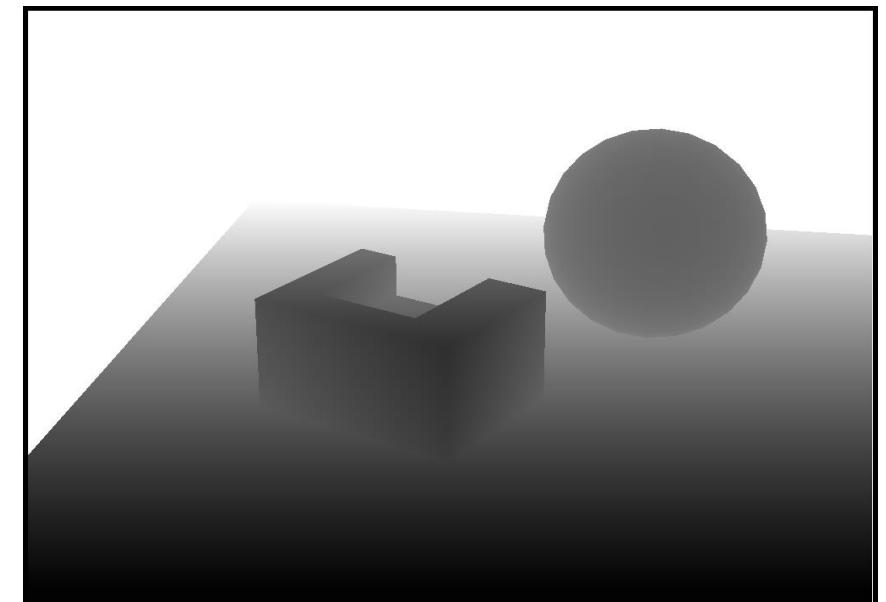
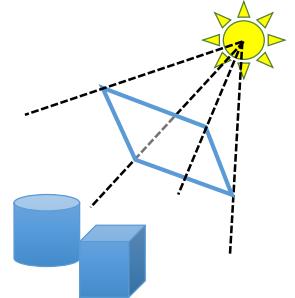


- *Motivation*
- Shadow Maps
- Shadow Volumes

# Shadow Maps

Shadow Map (light's view)

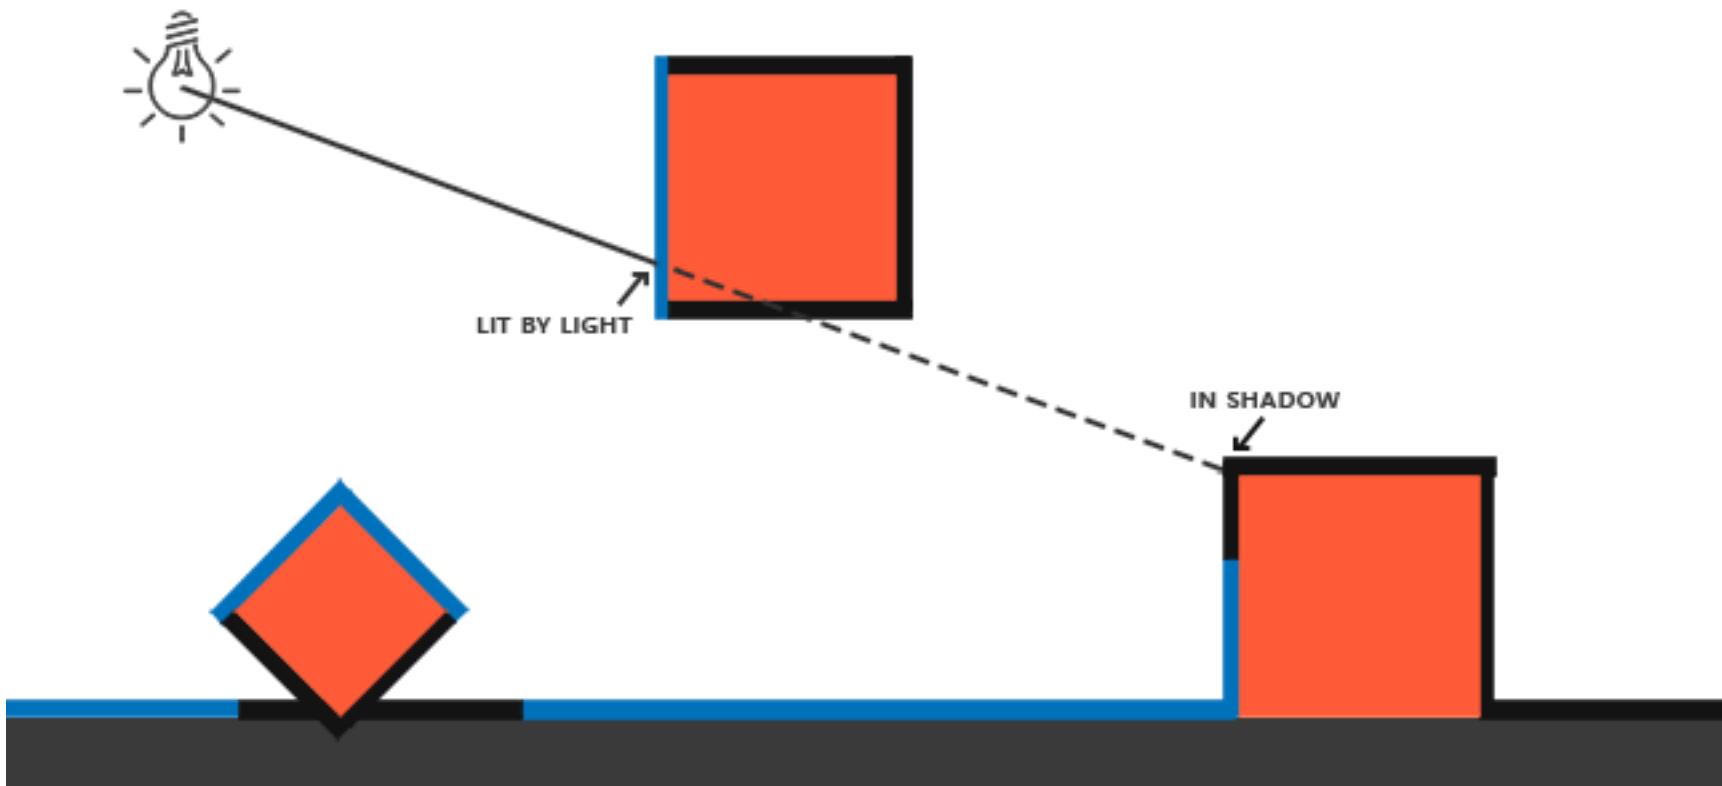
- Basic Algorithm – the simple explanation:
  - Render image from light source
    - Represents geometry in light
  - Render from camera
    - Test if rendered point is
      - visible in the light's view
        - If so -> point in light
        - Else -> point in shadow



# Shadow Maps

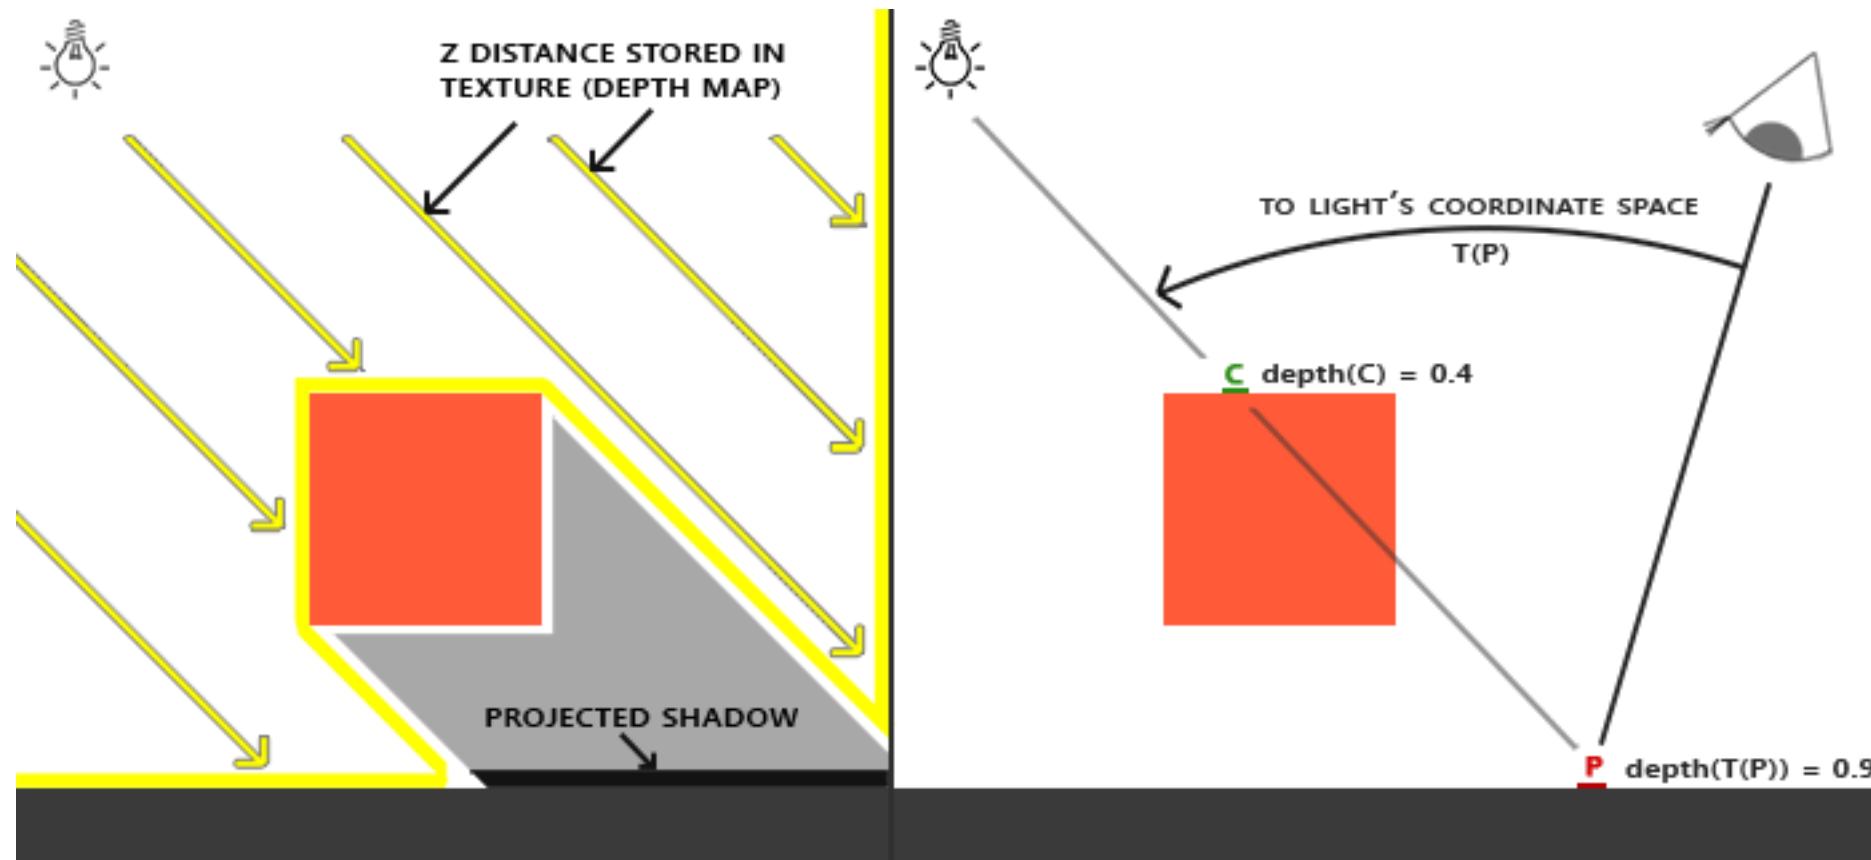
---

- Basic Algorithm – illustrated:

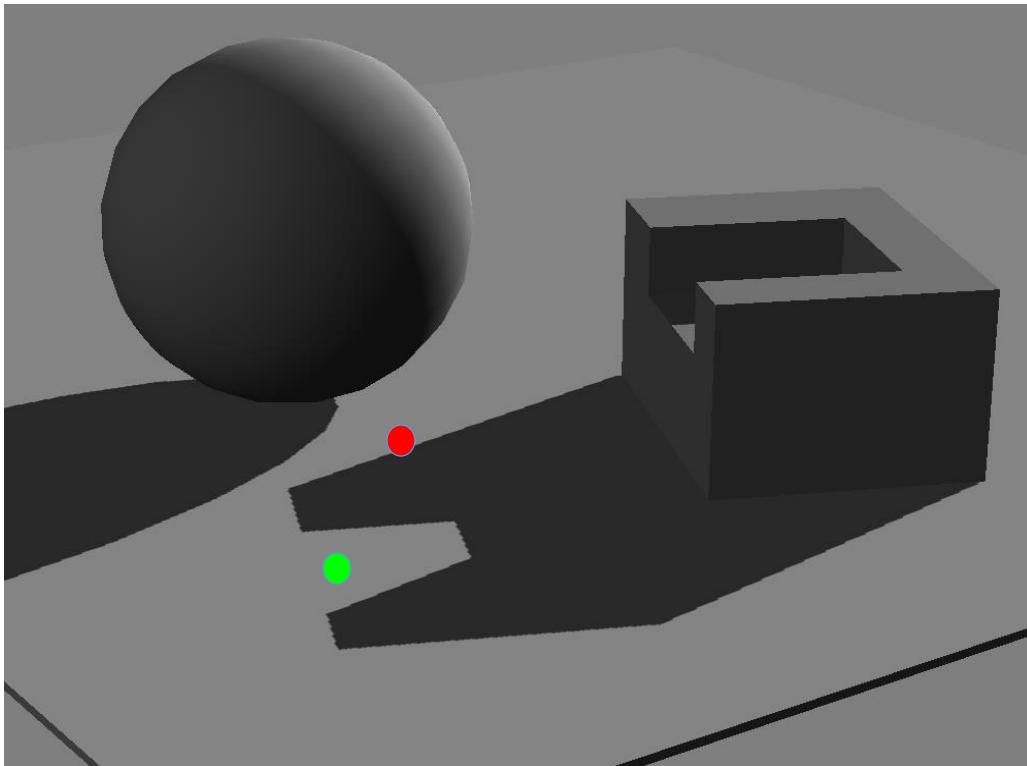


# Shadow Maps

- Basic Algorithm – illustrated:

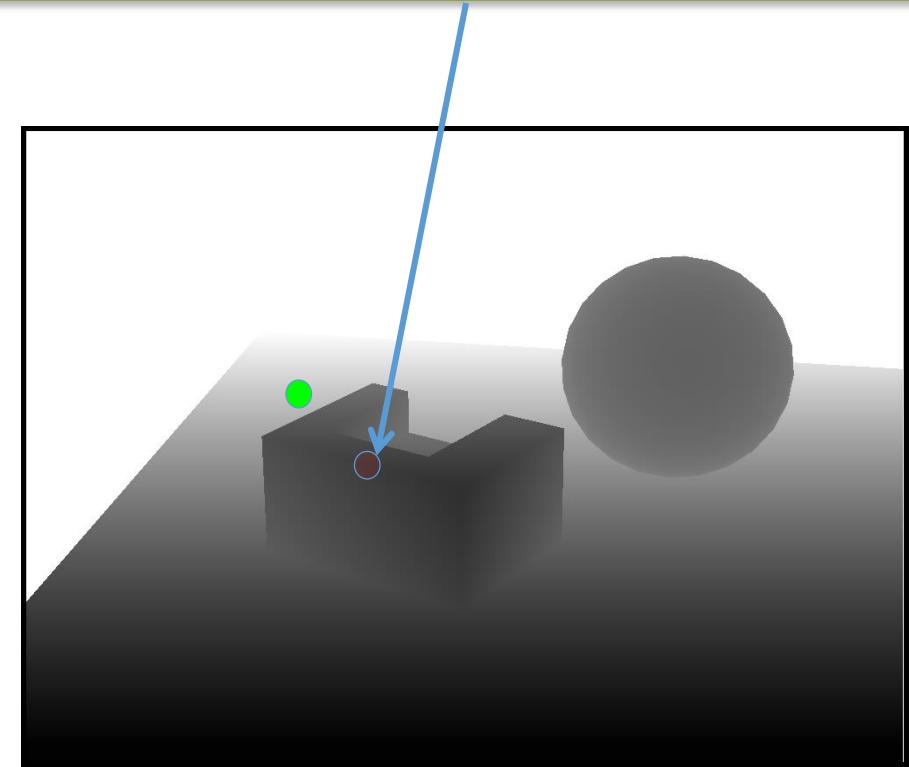


# Shadow Maps



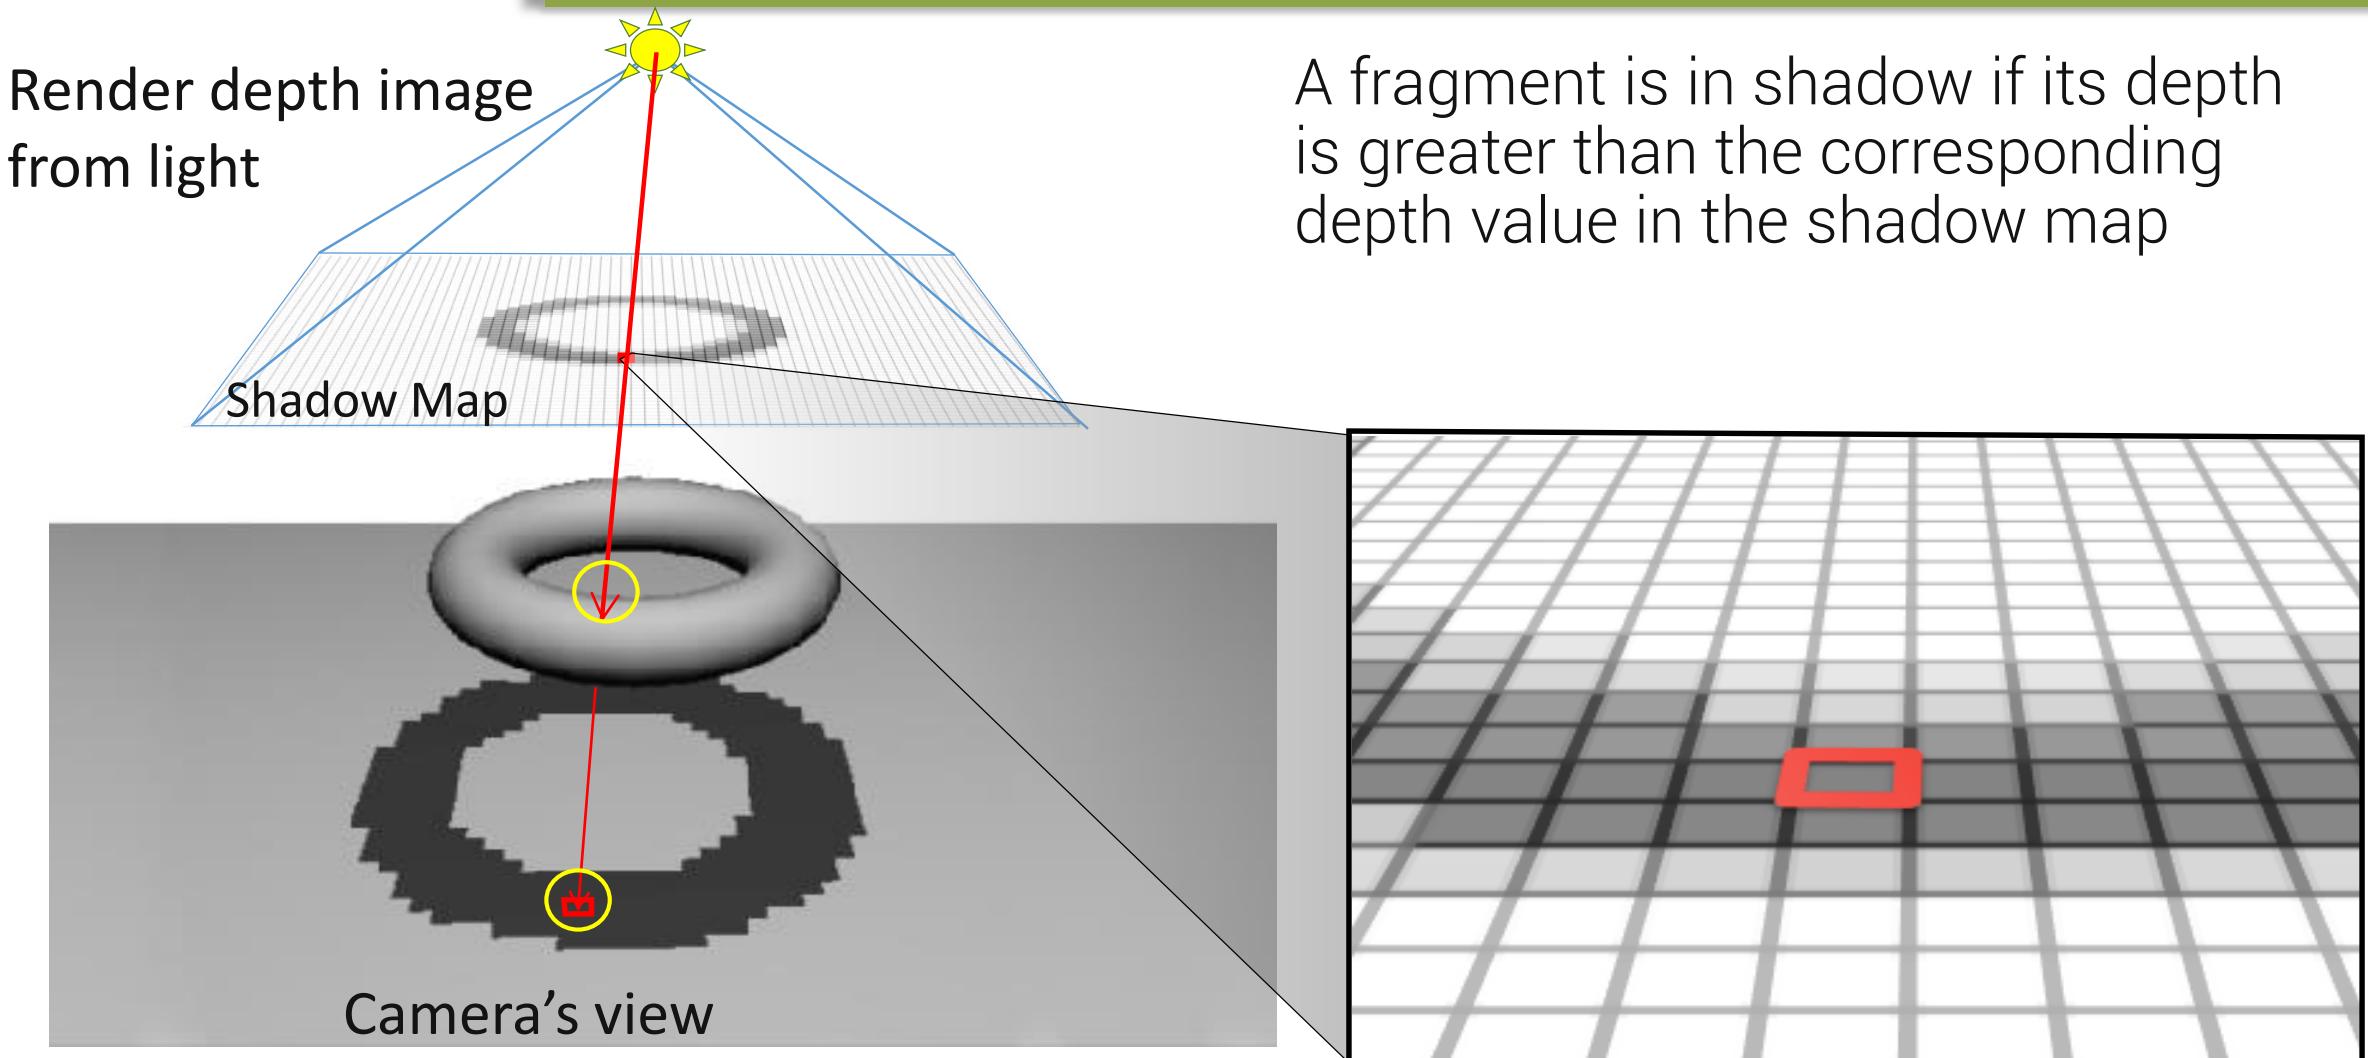
Camera's view

Point not represented in shadow map (point is behind box)



Light's view  
(Shadow Map)

# Depth Comparison



# Shadow Maps. Implementation

- Initial configuration parameters (background color, depth test...)

```
// Dark blue background  
glClearColor(0.0f, 0.0f, 0.4f, 0.0f);  
  
// Enable depth test  
 glEnable(GL_DEPTH_TEST);  
  
 // Accept fragment if it closer to the camera than the former one  
 glDepthFunc(GL_LESS);  
  
 // Cull triangles which normal is not towards the camera  
 glEnable(GL_CULL_FACE);
```

# Shadow Maps. Implementation

- Initial configuration parameters (background color, depth test...)

```
// Dark blue background  
glClearColor(0.0f, 0.0f, 0.4f, 0.0f);  
  
// Enable depth test  
 glEnable(GL_DEPTH_TEST);  
  
// Accept fragment if it closer to the camera than the former one  
 glDepthFunc(GL_LESS);  
  
// Cull triangles which normal is not towards the camera  
 glEnable(GL_CULL_FACE);
```

# Shadow Maps. Implementation

- Initial configuration parameters (background color, depth test...)

```
// Dark blue background  
glClearColor(0.0f, 0.0f, 0.4f, 0.0f);  
  
// Enable depth test  
 glEnable(GL_DEPTH_TEST);  
  
 // Accept fragment if it closer to the camera than the former one  
 glDepthFunc(GL_LESS);  
  
 // Cull triangles which normal is not towards the camera  
 glEnable(GL_CULL_FACE);
```

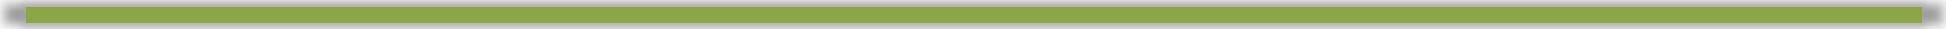
# Shadow Maps. Implementation

---

- Generate the framebuffer object (optionally multiple textures)

```
GLuint FramebufferName = 0;  
 glGenFramebuffers(1, &FramebufferName);  
 glBindFramebuffer(GL_FRAMEBUFFER, FramebufferName);
```

# Shadow Maps. Implementation



- Generate the depth texture

```
GLuint depthTexture;  
glGenTextures(1, &depthTexture);  
 glBindTexture(GL_TEXTURE_2D, depthTexture);
```

# Shadow Maps. Implementation

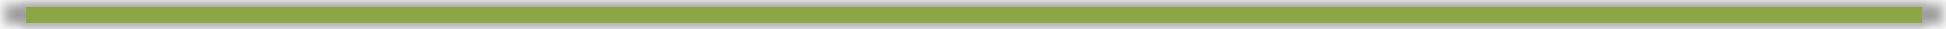
- Configure the depth texture

```
// Creating the internal memory and data structures with a void texture
glTexImage2D(GL_TEXTURE_2D, 0, GL_DEPTH_COMPONENT16, 1024, 1024, 0, GL_DEPTH_COMPONENT, GL_FLOAT, 0);

// Configuration parameters (for posterior access)
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, GL_LINEAR);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, GL_LINEAR);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_S, GL_CLAMP_TO_EDGE);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_T, GL_CLAMP_TO_EDGE);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_COMPARE_FUNC, GL_LEQUAL);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_COMPARE_MODE, GL_COMPARE_R_TO_TEXTURE);

// Attach the depth
glFramebufferTexture(GL_FRAMEBUFFER, GL_DEPTH_ATTACHMENT, depthTexture, 0);
```

# Shadow Maps. Implementation



- Ensure we render with no color
  - Typically 2x-4x faster on modern GPUs
  - Can render with stencil too
    - Depth+stencil are generated “at the same time”
  - Usually called “depth-only pass”

```
glDrawBuffer(GL_NONE);
```

# Shadow Maps. Implementation

- Create shaders & pass parameters

```
// Create and compile our GLSL program from the shaders
GLuint programID = LoadShaders( "ShadowMapping_SimpleVersion.vertexshader", "ShadowMapping_SimpleVersion.fragmentshader" );

// Get a handle for our "myTextureSampler" uniform
GLuint TextureID  = glGetUniformLocation(programID, "myTextureSampler");

// Get a handle for our "MVP" uniform
GLuint MatrixID = glGetUniformLocation(programID, "MVP");
GLuint DepthBiasID = glGetUniformLocation(programID, "DepthBiasMVP");
GLuint ShadowMapID = glGetUniformLocation(programID, "shadowMap");
```

# Shadow Maps. Implementation

- Create shaders & pass parameters

```
// Create and compile our GLSL program from the shaders
GLuint programID = LoadShaders( "ShadowMapping_SimpleVersion.vertexshader", "ShadowMapping_SimpleVersion.fragmentshader" );

// Get a handle for our "myTextureSampler" uniform
GLuint TextureID  = glGetUniformLocation(programID, "myTextureSampler");

// Get a handle for our "MVP" uniform
GLuint MatrixID = glGetUniformLocation(programID, "MVP");
GLuint DepthBiasID = glGetUniformLocation(programID, "DepthBiasMVP");
GLuint ShadowMapID = glGetUniformLocation(programID, "shadowMap");
```

# Shadow Maps. Implementation

- Render from the light source

```
// Compute the MVP matrix from the Light's point of view
glm::vec3 lightPos(5, 20, 20);
glm::mat4 depthProjectionMatrix = glm::perspective<float>(45.0f, 1.0f, 2.0f, 50.0f);
glm::mat4 depthViewMatrix = glm::lookAt(lightPos, lightPos-lightInvDir, glm::vec3(0,1,0));

glm::mat4 depthModelMatrix = glm::mat4(1.0);
glm::mat4 depthMVP = depthProjectionMatrix * depthViewMatrix * depthModelMatrix;

// Send our transformation to the currently bound shader, in the "MVP" uniform
glUniformMatrix4fv(depthMatrixID, 1, GL_FALSE, &depthMVP[0][0]);

drawGeometryDepth(indices, vertexbuffer, elementbuffer);
```

# Shadow Maps. Implementation

- Render from the light source

```
// Compute the MVP matrix from the Light's point of view
glm::vec3 lightPos(5, 20, 20);
glm::mat4 depthProjectionMatrix = glm::perspective<float>(45.0f, 1.0f, 2.0f, 50.0f);
glm::mat4 depthViewMatrix = glm::lookAt(lightPos, lightPos-lightInvDir, glm::vec3(0,1,0));

glm::mat4 depthModelMatrix = glm::mat4(1.0);
glm::mat4 depthMVP = depthProjectionMatrix * depthViewMatrix * depthModelMatrix;

// Send our transformation to the currently bound shader, in the "MVP" uniform
glUniformMatrix4fv(depthMatrixID, 1, GL_FALSE, &depthMVP[0][0]);

drawGeometryDepth(indices, vertexbuffer, elementbuffer);
```

# Shadow Maps. Implementation

- Render from the light source. Shaders:

```
1 #version 330 core
2
3 // Input vertex data, different for all executions of this shader.
4 layout(location = 0) in vec3 vertexPosition_modelspace;
5
6 // Values that stay constant for the whole mesh.
7 uniform mat4 depthMVP;
8
9 void main(){
10     gl_Position = depthMVP * vec4(vertexPosition_modelspace,1);
11 }
```

# Shadow Maps. Implementation

- Render from the light source. Shaders:

```
#version 330 core

// Output data
layout(location = 0) out float fragmentdepth;

void main(){
    fragmentdepth = gl_FragCoord.z;
}
```

# Shadow Maps. Implementation

---

- Render from the observer's view

```
// Render to the screen
glBindFramebuffer(GL_FRAMEBUFFER, 0);
glViewport(0,0,windowWidth,windowHeight); // Render on the whole framebuffer, complete from the Lower Left corner

glEnable(GL_CULL_FACE);
glCullFace(GL_BACK); // Cull back-facing triangles -> draw only front-facing triangles

// Clear the screen
glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);
```

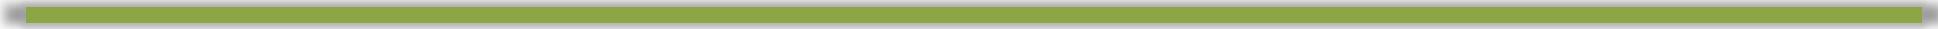
# Shadow Maps. Implementation

- Render from the observer's view

```
// Use our shader
glUseProgram(programID);

// Compute the MVP matrix from keyboard and mouse input
computeMatricesFromInputs();
glm::mat4 ProjectionMatrix = getProjectionMatrix();
glm::mat4 ViewMatrix = getViewMatrix();
//ViewMatrix = glm::LookAt(glm::vec3(14,6,4), glm::vec3(0,1,0), glm::vec3(0,1,0));
glm::mat4 ModelMatrix = glm::mat4(1.0);
glm::mat4 MVP = ProjectionMatrix * ViewMatrix * ModelMatrix;
```

# Shadow Maps. Implementation



- Render from the observer's view

```
glm::mat4 biasMatrix(  
    0.5, 0.0, 0.0, 0.0,  
    0.0, 0.5, 0.0, 0.0,  
    0.0, 0.0, 0.5, 0.0,  
    0.5, 0.5, 0.5, 1.0  
);  
  
glm::mat4 depthBiasMVP = biasMatrix*depthMVP;
```

# Shadow Maps. Implementation

- Render from the observer's view

```
// Send our transformation to the currently bound shader,  
// in the "MVP" uniform  
glUniformMatrix4fv(MatrixID, 1, GL_FALSE, &MVP[0][0]);  
glUniformMatrix4fv(DepthBiasID, 1, GL_FALSE, &depthBiasMVP[0][0]);  
  
// Bind our texture in Texture Unit 0  
glActiveTexture(GL_TEXTURE0);  
glBindTexture(GL_TEXTURE_2D, Texture);  
// Set our "myTextureSampler" sampler to use Texture Unit 0  
glUniform1i(TextureID, 0);  
  
glActiveTexture(GL_TEXTURE1);  
glBindTexture(GL_TEXTURE_2D, depthTexture);  
glUniform1i(ShadowMapID, 1);  
  
drawGeometry(indices, vertexbuffer, elementbuffer, uvbuffer, normalbuffer, Texture);
```

# Shadow Maps. Implementation

- Render from the observer's view

```
// Send our transformation to the currently bound shader,  
// in the "MVP" uniform  
glUniformMatrix4fv(MatrixID, 1, GL_FALSE, &MVP[0][0]);  
glUniformMatrix4fv(DepthBiasID, 1, GL_FALSE, &depthBiasMVP[0][0]);  
  
// Bind our texture in Texture Unit 0  
glActiveTexture(GL_TEXTURE0);  
glBindTexture(GL_TEXTURE_2D, Texture);  
// Set our "myTextureSampler" sampler to use Texture Unit 0  
glUniform1i(TextureID, 0);  
  
glActiveTexture(GL_TEXTURE1);  
glBindTexture(GL_TEXTURE_2D, depthTexture);  
glUniform1i(ShadowMapID, 1);  
  
drawGeometry(indices, vertexbuffer, elementbuffer, uvbuffer, normalbuffer, Texture);
```

# Shadow Maps. Implementation

- Render from the observer's view

```
// Send our transformation to the currently bound shader,  
// in the "MVP" uniform  
glUniformMatrix4fv(MatrixID, 1, GL_FALSE, &MVP[0][0]);  
glUniformMatrix4fv(DepthBiasID, 1, GL_FALSE, &depthBiasMVP[0][0]);  
  
// Bind our texture in Texture Unit 0  
glActiveTexture(GL_TEXTURE0);  
glBindTexture(GL_TEXTURE_2D, Texture);  
// Set our "myTextureSampler" sampler to use Texture Unit 0  
glUniform1i(TextureID, 0);  
  
glActiveTexture(GL_TEXTURE1);  
glBindTexture(GL_TEXTURE_2D, depthTexture);  
glUniform1i(ShadowMapID, 1);  
  
drawGeometry(indices, vertexbuffer, elementbuffer, uvbuffer, normalbuffer, Texture);
```

# Shadow Maps. Implementation

- Render from the observer's view

```
// Send our transformation to the currently bound shader,  
// in the "MVP" uniform  
glUniformMatrix4fv(MatrixID, 1, GL_FALSE, &MVP[0][0]);  
glUniformMatrix4fv(DepthBiasID, 1, GL_FALSE, &depthBiasMVP[0][0]);  
  
// Bind our texture in Texture Unit 0  
glActiveTexture(GL_TEXTURE0);  
glBindTexture(GL_TEXTURE_2D, Texture);  
// Set our "myTextureSampler" sampler to use Texture Unit 0  
glUniform1i(TextureID, 0);  
  
glActiveTexture(GL_TEXTURE1);  
glBindTexture(GL_TEXTURE_2D, depthTexture);  
glUniform1i(ShadowMapID, 1);  
  
drawGeometry(indices, vertexbuffer, elementbuffer, uvbuffer, normalbuffer, Texture);
```

# Shadow Maps. Implementation

- Render from the observer's view

```
glActiveTexture(GL_TEXTURE1);
 glBindTexture(GL_TEXTURE_2D, depthTexture);
 glUniform1i(ShadowMapID, 1);

 drawGeometry(indices, vertexbuffer, elementbuffer, uvbuffer, normalbuffer, Texture);

 // Swap buffers
 glfwSwapBuffers(window);
 glfwPollEvents();
```

# Shadow Maps. Implementation

- Render from the observer's view. Vertex shader

```
1 #version 330 core
2
3 // Input vertex data, different for all executions of this shader.
4 layout(location = 0) in vec3 vertexPosition_modelspace;
5 layout(location = 1) in vec2 vertexUV;
6 layout(location = 2) in vec3 vertexNormal_modelspace;
7
8 // Output data ; will be interpolated for each fragment.
9 out vec2 UV;
10 out vec4 ShadowCoord;
11
12 // Values that stay constant for the whole mesh.
13 uniform mat4 MVP;
14 uniform mat4 DepthBiasMVP;
15
```

# Shadow Maps. Implementation

- Render from the observer's view . Vertex shader

```
1 #version 330 core
2
3 // Input vertex data, different for all executions of this shader.
4 layout(location = 0) in vec3 vertexPosition_modelspace;
5 layout(location = 1) in vec2 vertexUV;
6 layout(location = 2) in vec3 vertexNormal_modelspace;
7
8 // Output data ; will be interpolated for each fragment.
9 out vec2 UV;
10 out vec4 ShadowCoord;
11
12 // Values that stay constant for the whole mesh.
13 uniform mat4 MVP;
14 uniform mat4 DepthBiasMVP;
15
```

# Shadow Maps. Implementation

- Render from the observer's view . Vertex shader

```
1 #version 330 core
2
3 // Input vertex data, different for all executions of this shader.
4 layout(location = 0) in vec3 vertexPosition_modelspace;
5 layout(location = 1) in vec2 vertexUV;
6 layout(location = 2) in vec3 vertexNormal_modelspace;
7
8 // Output data ; will be interpolated for each fragment.
9 out vec2 UV;
10 out vec4 ShadowCoord;
11
12 // Values that stay constant for the whole mesh.
13 uniform mat4 MVP;
14 uniform mat4 DepthBiasMVP;
```

# Shadow Maps. Implementation

- Render from the observer's view . Vertex shader

```
16  
17 void main(){  
18  
19     // Output position of the vertex, in clip space : MVP * position  
20     gl_Position = MVP * vec4(vertexPosition_modelspace,1);  
21  
22     ShadowCoord = DepthBiasMVP * vec4(vertexPosition_modelspace,1);  
23  
24     // UV of the vertex. No special space for this one.  
25     UV = vertexUV;  
26 }  
27
```

# Shadow Maps. Implementation

- Render from the observer's view . Vertex shader

```
16  
17 void main(){  
18  
19     // Output position of the vertex, in clip space : MVP * position  
20     gl_Position = MVP * vec4(vertexPosition_modelspace,1);  
21  
22     ShadowCoord = DepthBiasMVP * vec4(vertexPosition_modelspace,1);  
23  
24     // UV of the vertex. No special space for this one.  
25     UV = vertexUV;  
26 }  
27
```

# Shadow Maps. Implementation

- Render from the observer's view. Vertex shader

```
16  
17 void main(){  
18  
19     // Output position of the vertex, in clip space : MVP * position  
20     gl_Position = MVP * vec4(vertexPosition_modelspace,1);  
21  
22     ShadowCoord = DepthBiasMVP * vec4(vertexPosition_modelspace,1);  
23  
24     // UV of the vertex. No special space for this one.  
25     UV = vertexUV;  
26 }  
27
```

# Shadow Maps. Implementation

- Render from the observer's view. Fragment shader

```
1 #version 330 core
2
3 // Interpolated values from the vertex shaders
4 in vec2 UV;
5 in vec4 ShadowCoord;
6
7 // Output data
8 layout(location = 0) out vec3 color;
9
10 // Values that stay constant for the whole mesh.
11 uniform sampler2D myTextureSampler;
12 uniform sampler2DShadow shadowMap;
```

# Shadow Maps. Implementation

- Render from the observer's view. Fragment shader

```
1 #version 330 core
2
3 // Interpolated values from the vertex shaders
4 in vec2 UV;
5 in vec4 ShadowCoord;
6
7 // Output data
8 layout(location = 0) out vec3 color;
9
10 // Values that stay constant for the whole mesh.
11 uniform sampler2D myTextureSampler;
12 uniform sampler2DShadow shadowMap;
```

# Shadow Maps. Implementation

- Render from the observer's view. Fragment shader

```
1 #version 330 core
2
3 // Interpolated values from the vertex shaders
4 in vec2 UV;
5 in vec4 ShadowCoord;
6
7 // Output data
8 layout(location = 0) out vec3 color;
9
10 // Values that stay constant for the whole mesh.
11 uniform sampler2D myTextureSampler;
12 uniform sampler2DShadow shadowMap;
```

# Shadow Maps. Implementation

- Render from the observer's view. Fragment shader

```
13  
14 void main(){  
15  
16     // Light emission properties  
17     vec3 LightColor = vec3(1,1,1);  
18  
19     // Material properties  
20     vec3 MaterialDiffuseColor = texture( myTextureSampler, UV ).rgb;  
21  
22     float visibility = texture( shadowMap, vec3(ShadowCoord.xy, (ShadowCoord.z)/ShadowCoord.w) );  
23  
24     color = visibility * MaterialDiffuseColor * LightColor;  
25  
26 }
```

# Shadow Maps. Implementation

- Render from the observer's view. Fragment shader

```
13  
14 void main(){  
15  
16     // Light emission properties  
17     vec3 LightColor = vec3(1,1,1);  
18  
19     // Material properties  
20     vec3 MaterialDiffuseColor = texture( myTextureSampler, UV ).rgb;  
21  
22     float visibility = texture( shadowMap, vec3(ShadowCoord.xy, (ShadowCoord.z)/ShadowCoord.w) );  
23  
24     color = visibility * MaterialDiffuseColor * LightColor;  
25  
26 }
```

# Shadow Maps. Implementation

- Render from the observer's view. Fragment shader

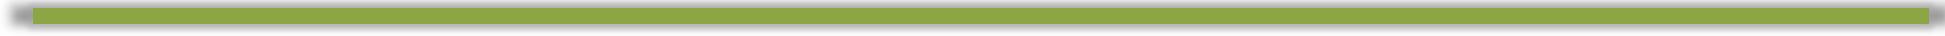
```
13  
14 void main(){  
15  
16     // Light emission properties  
17     vec3 LightColor = vec3(1,1,1);  
18  
19     // Material properties  
20     vec3 MaterialDiffuseColor = texture( myTextureSampler, UV ).rgb;  
21  
22     float visibility = texture( shadowMap, vec3(ShadowCoord.xy, (ShadowCoord.z)/ShadowCoord.w) );  
23  
24     color = visibility * MaterialDiffuseColor * LightColor;  
25  
26 }
```

# Shadow Maps. Implementation

- Render from the observer's view. Fragment shader

```
13  
14 void main(){  
15  
16     // Light emission properties  
17     vec3 LightColor = vec3(1,1,1);  
18  
19     // Material properties  
20     vec3 MaterialDiffuseColor = texture( myTextureSampler, UV ).rgb;  
21  
22     float visibility = texture( shadowMap, vec3(ShadowCoord.xy, (ShadowCoord.z)/ShadowCoord.w) );  
23  
24     color = visibility * MaterialDiffuseColor * LightColor;  
25  
26 }
```

# Shadow Maps



- Pros
  - Very efficient: "This is as fast as it gets"
- Cons...

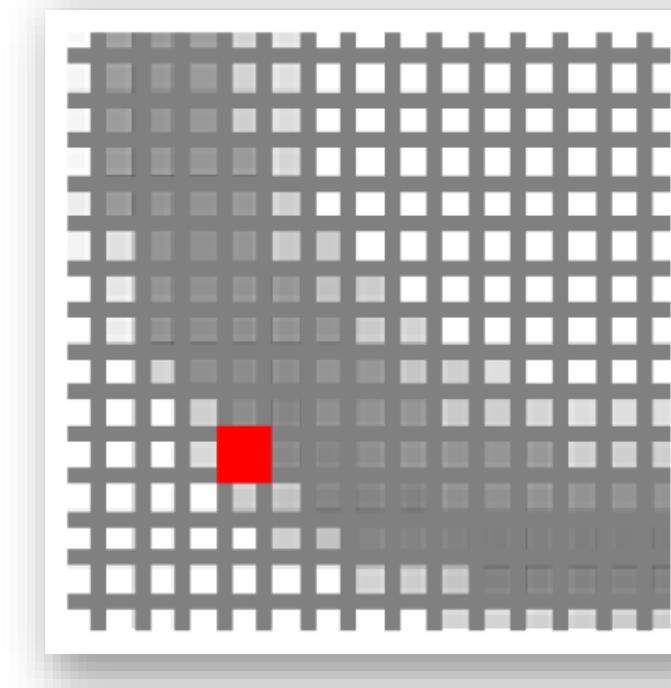
# Shadow Maps - Problems

---

- Low Shadow Map resolution results in jagged shadows

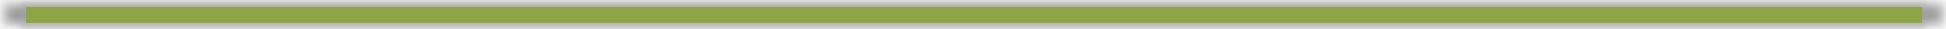


from viewpoint



from light

# Shadow Maps - Problems



In addition:

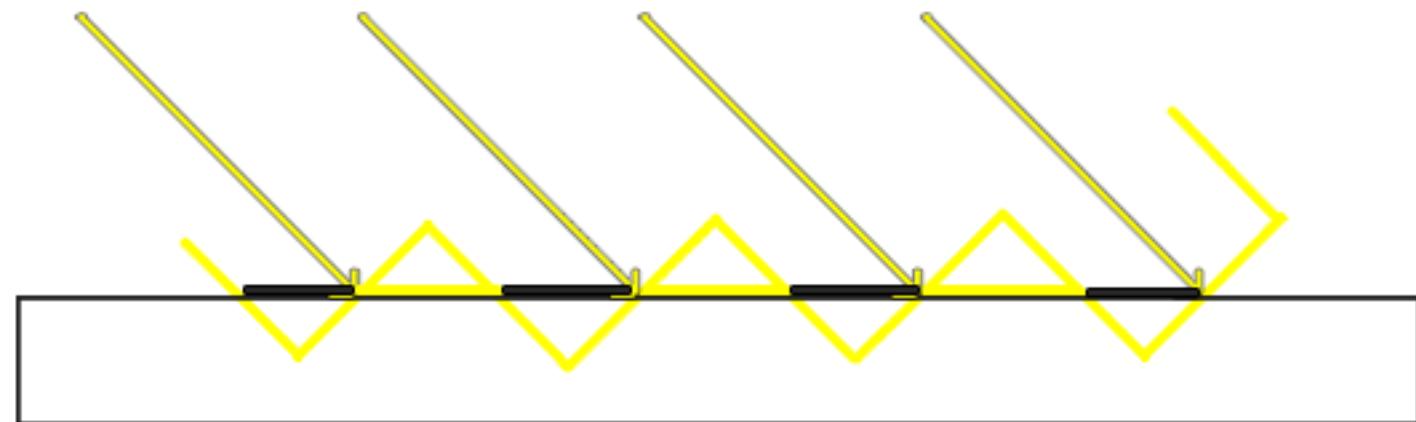
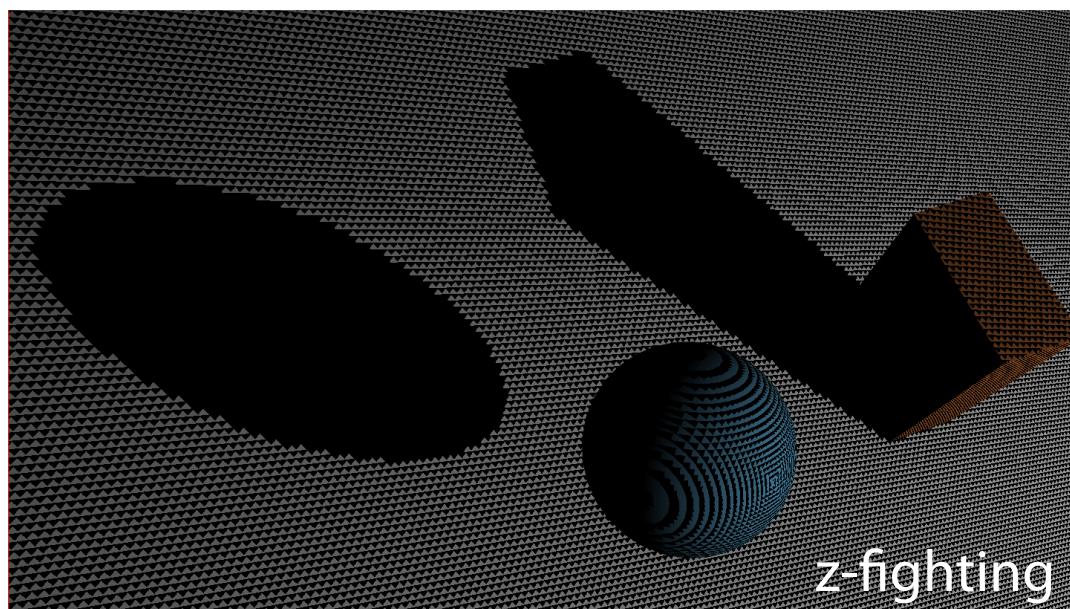
- A tolerance threshold (bias) needs to be tuned **for each scene** for the depth comparison

# Shadow Maps - Problems

---

Typical problems:

- Shadow acne, due to z-fighting
  - Light position and view position are very similar, floating point operations lead to z-fighting

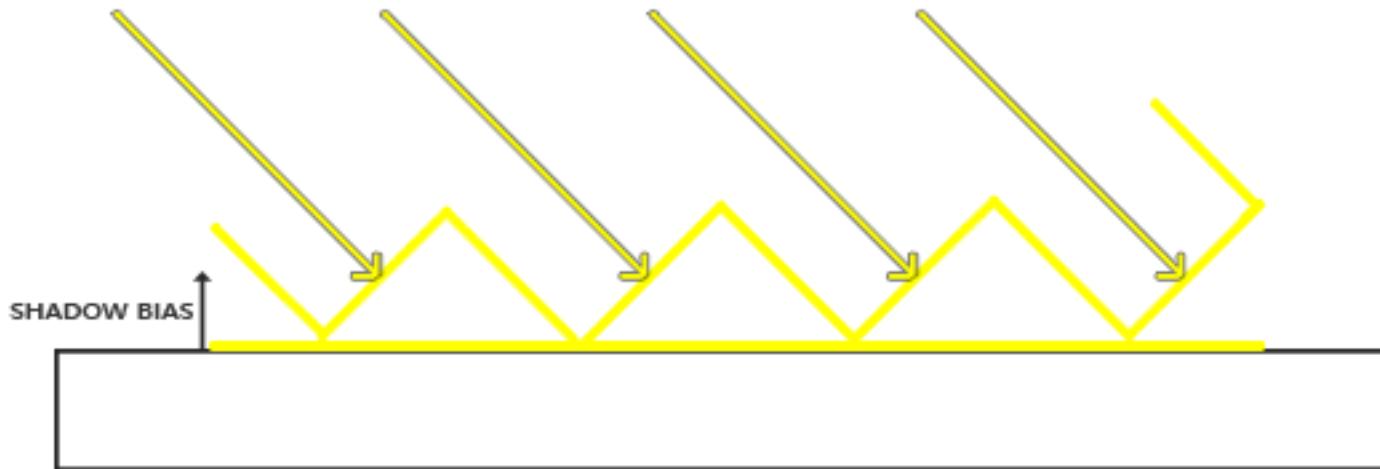


# Shadow Maps - Problems

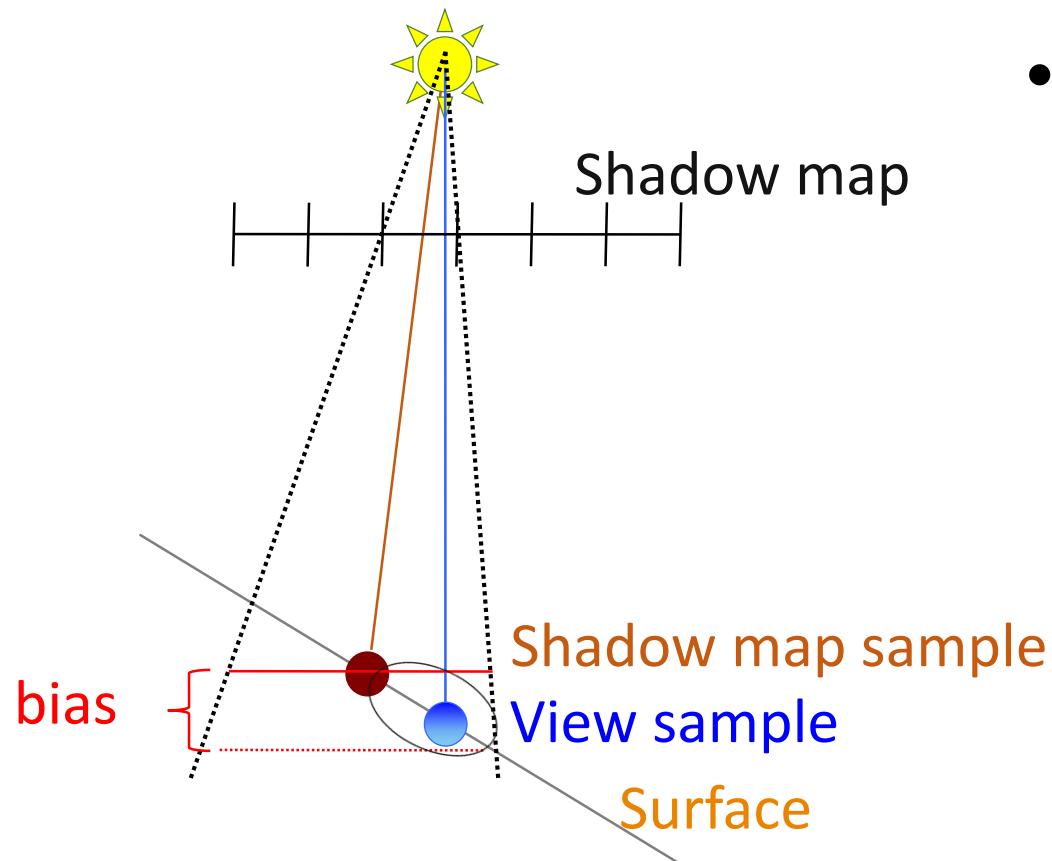
---

Typical problems:

- Shadow acne, can be solved using shadow bias (offset the depth of the surface – or the shadow map – a small bias to avoid this

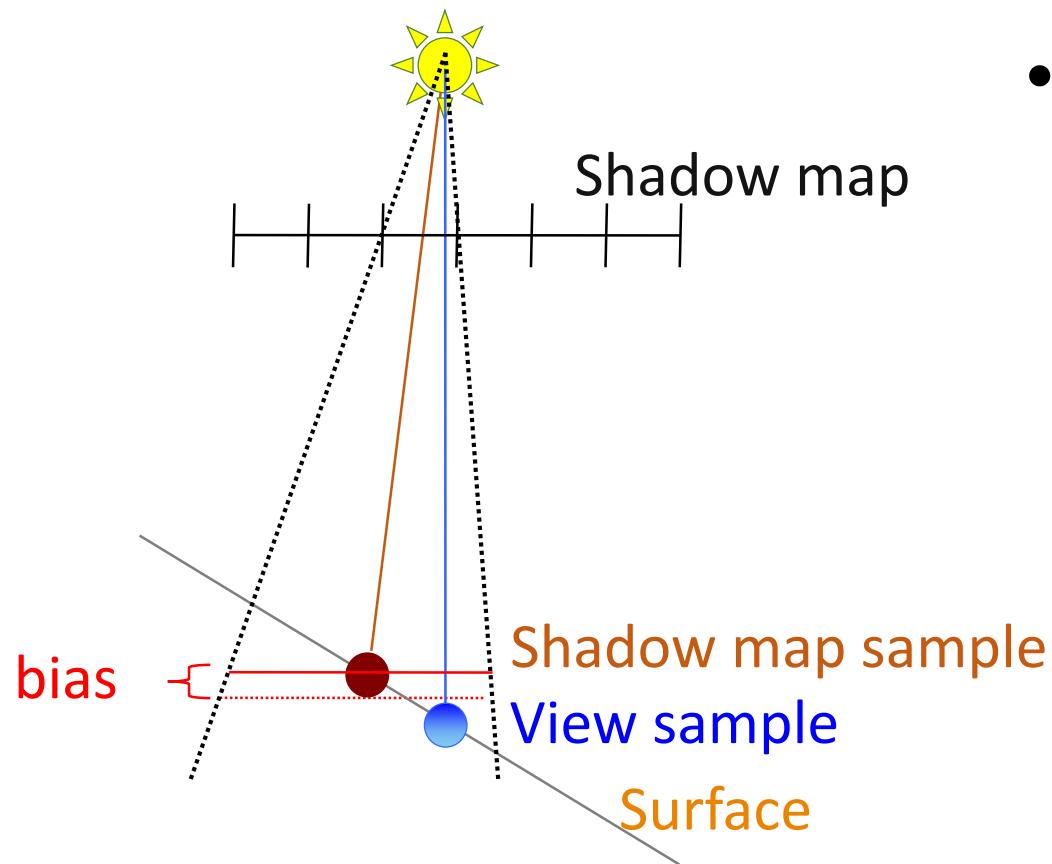


# Bias

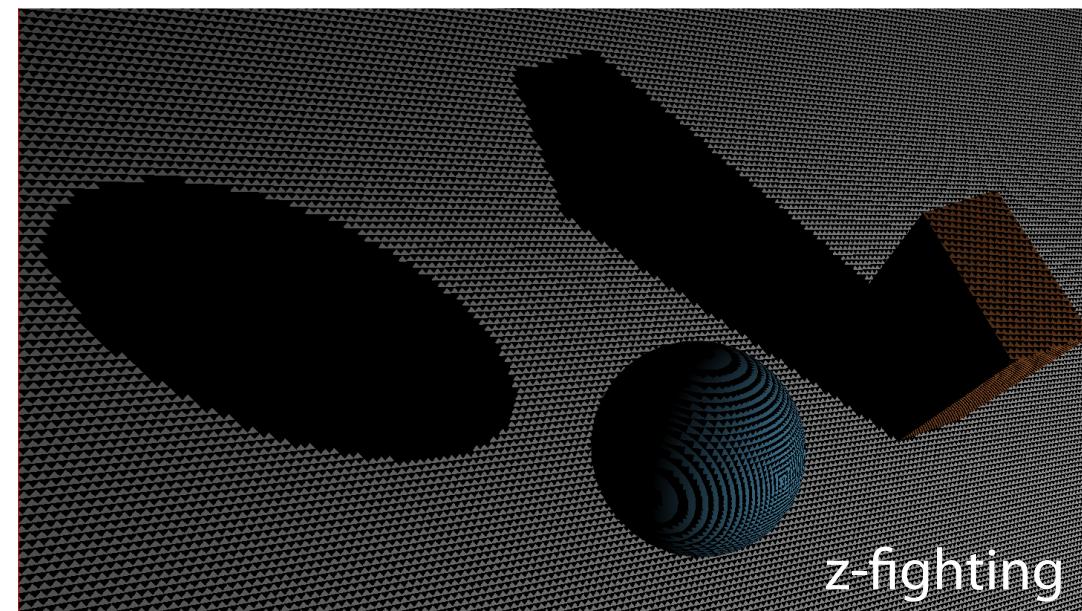


- Need a tolerance threshold (depth bias) when comparing depths to avoid surface self shadowing

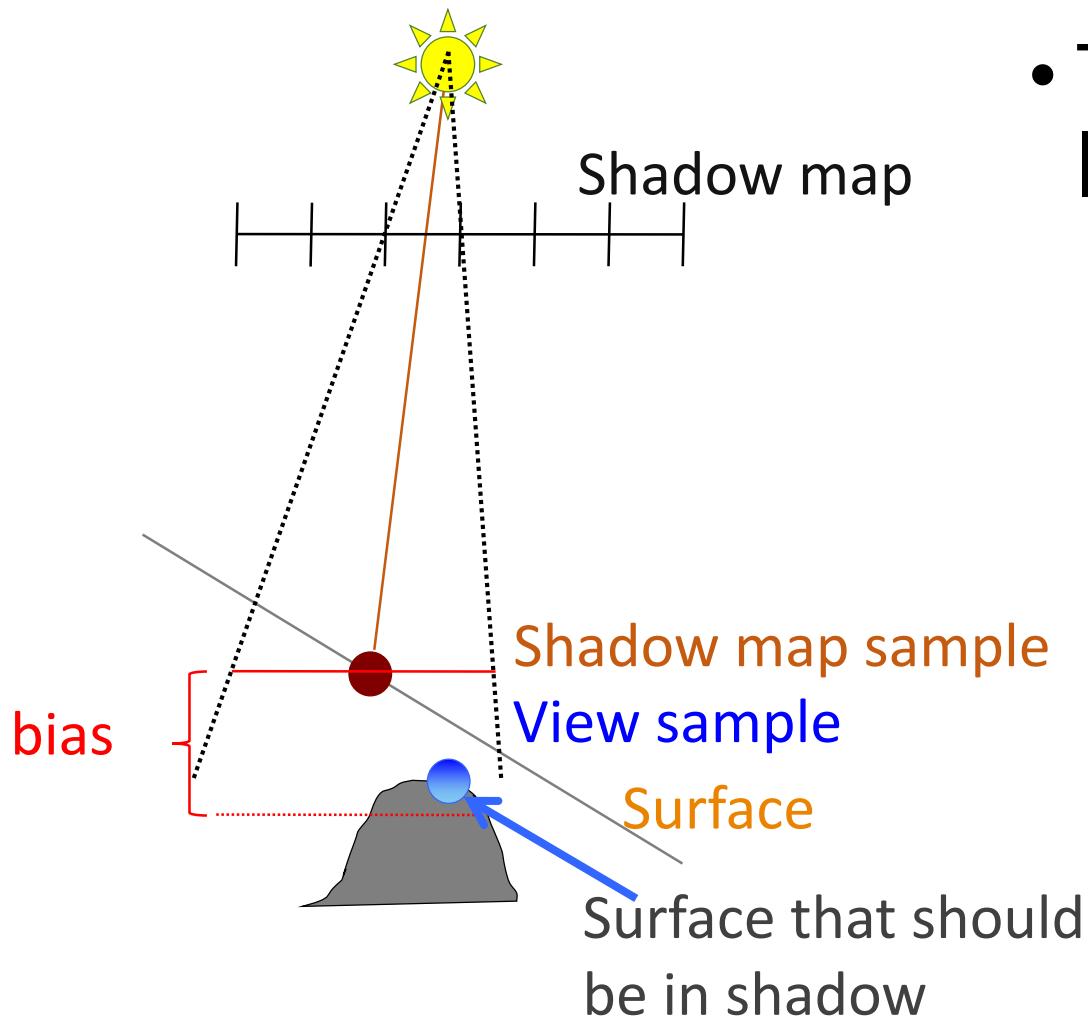
# Bias



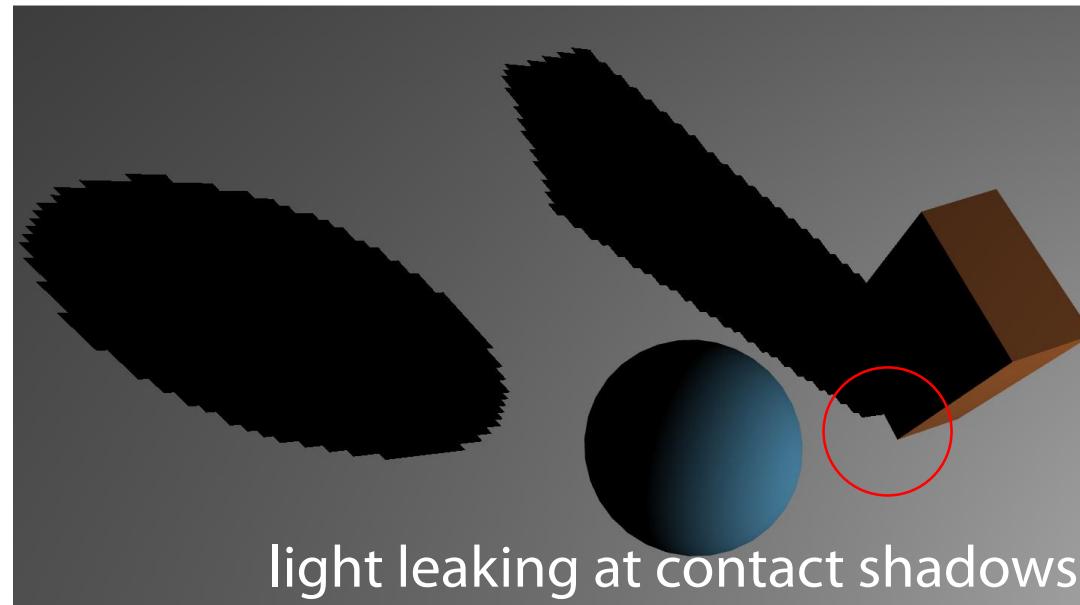
- Need a tolerance threshold (depth bias) when comparing depths to avoid surface self shadowing



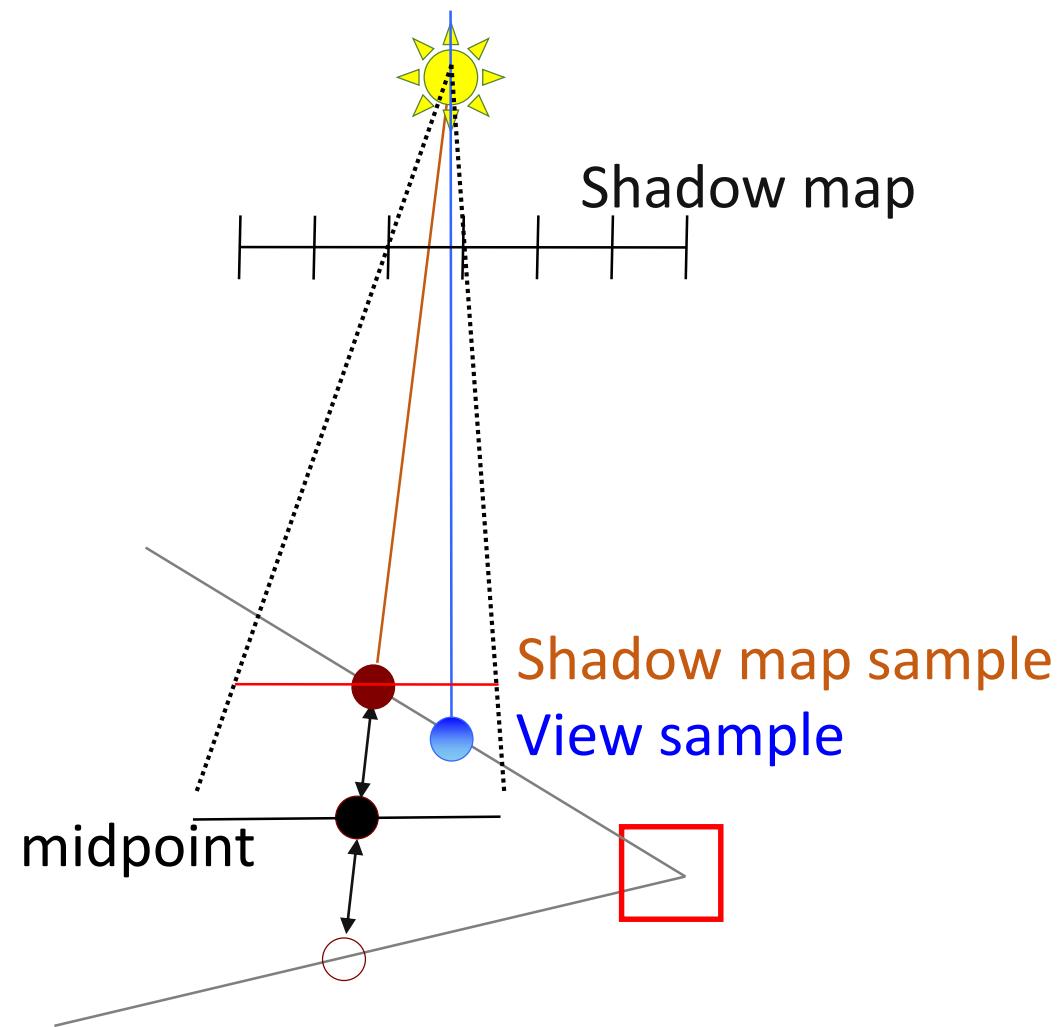
# Bias



- Too much bias leads to light leaking (aka peter panning)

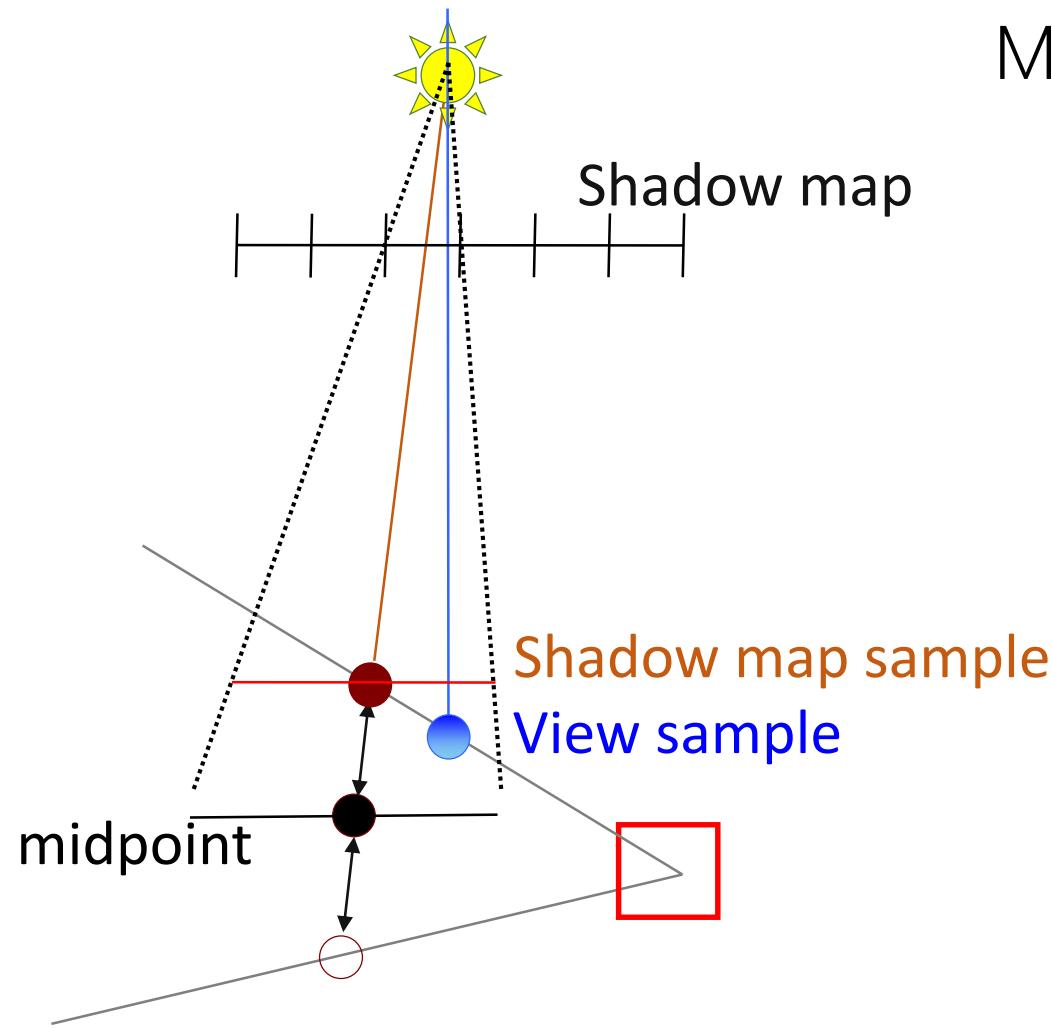


# Ameliorating the Bias

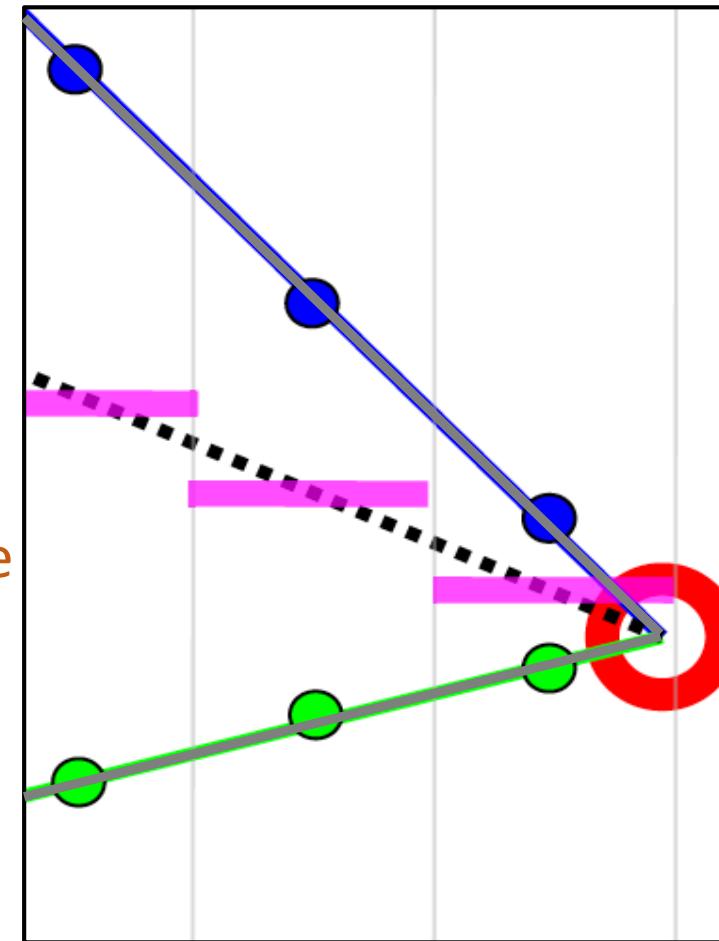


- Midpoint Shadow Maps [Woo 92]
- Second Depth Shadow Mapping [Wang and Molnar94]
- Dual Depth Layer [Weiskopf and Ertl 04]

# Ameliorating the Bias



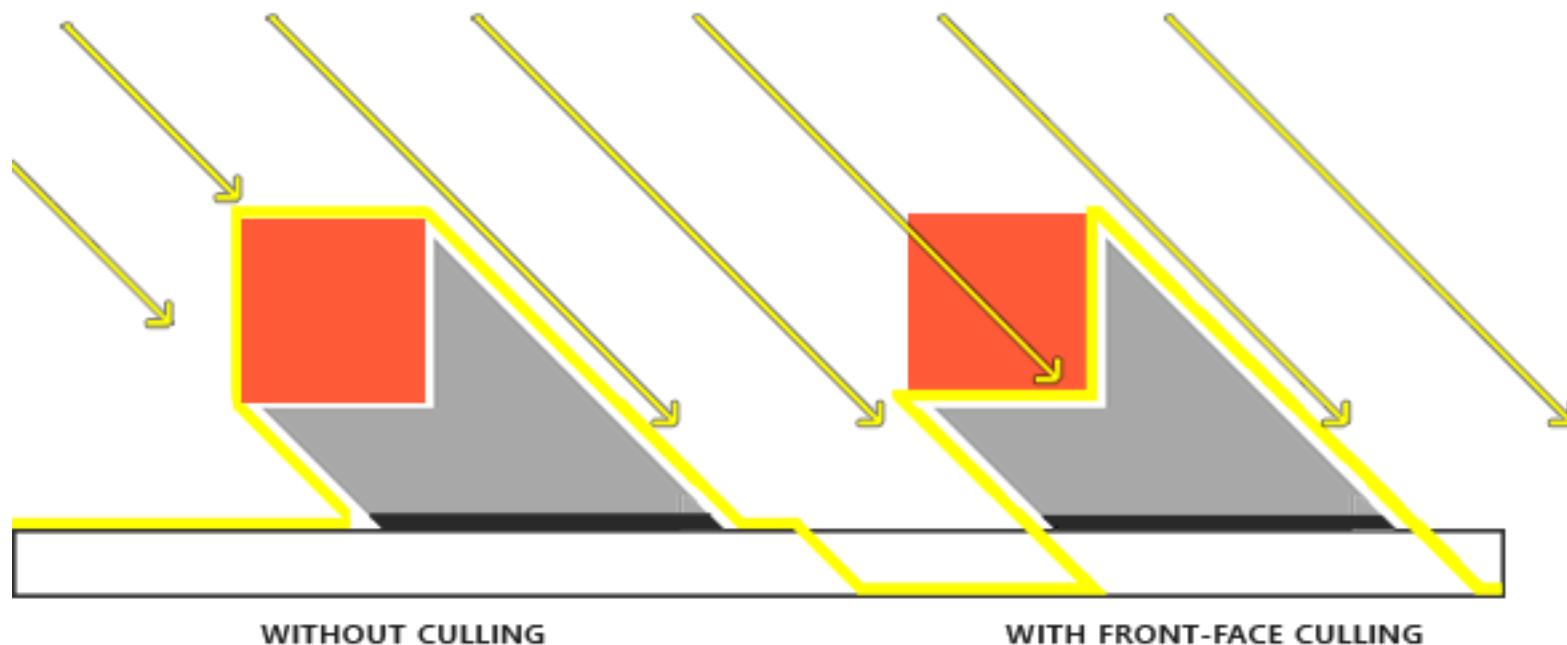
Midpoint Shadow Maps [Woo 92]



# Ameliorating the Bias

---

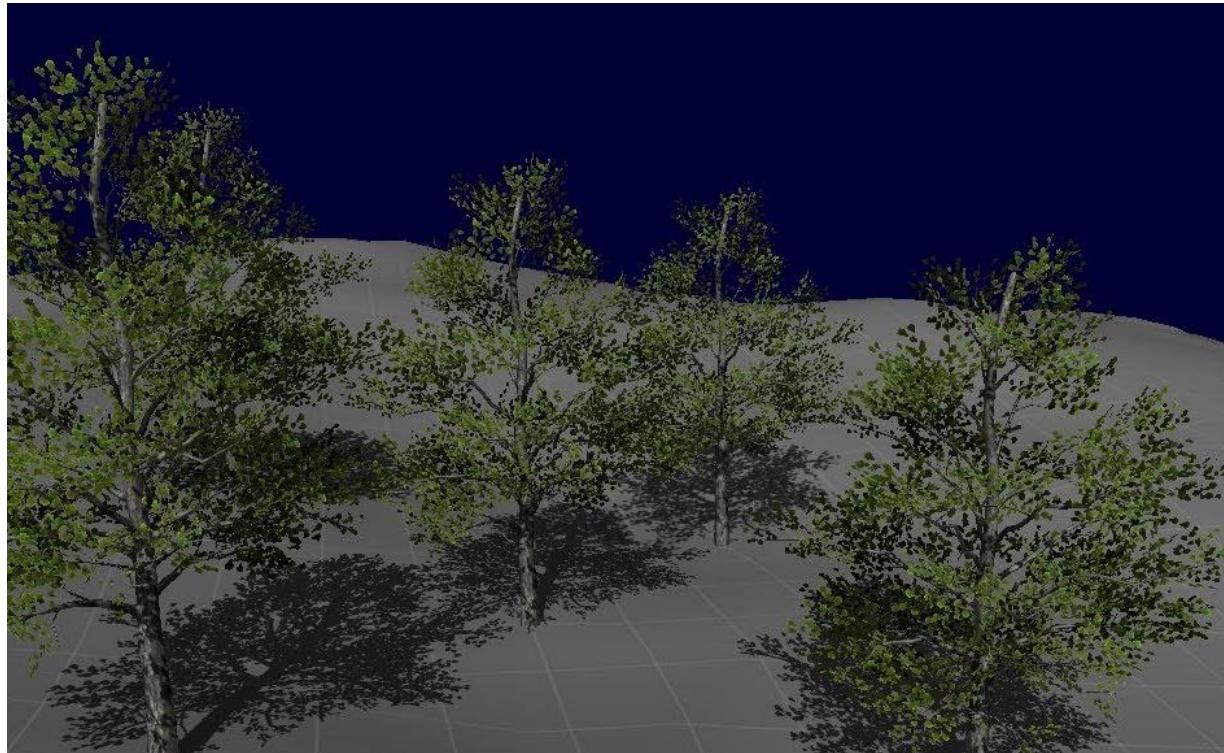
- We can also use the back faces for the depth map creation
  - Using front-face culling



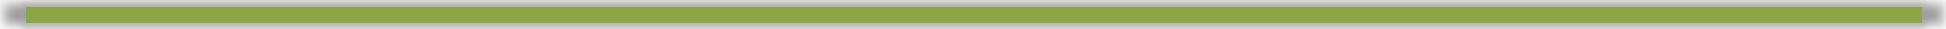
# Shadow Maps

---

- Shadow maps are the basis of most modern shadowing techniques
  - E.g. cascaded shadow maps



# Outline

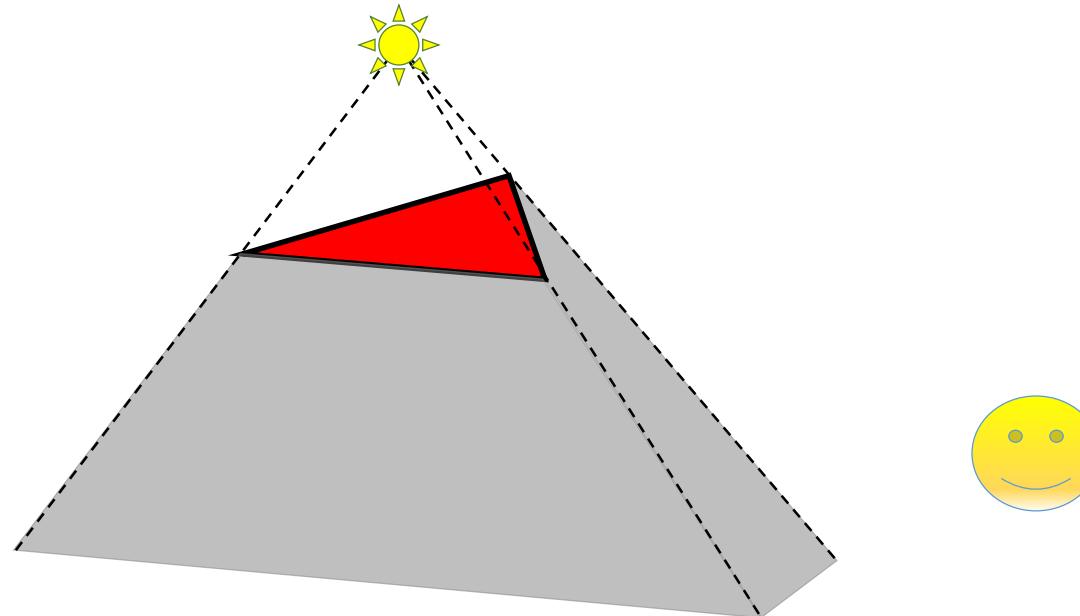


- *Motivation*
- *Shadow Maps*
- *Shadow Volumes*

# Shadow Volumes. Concept

---

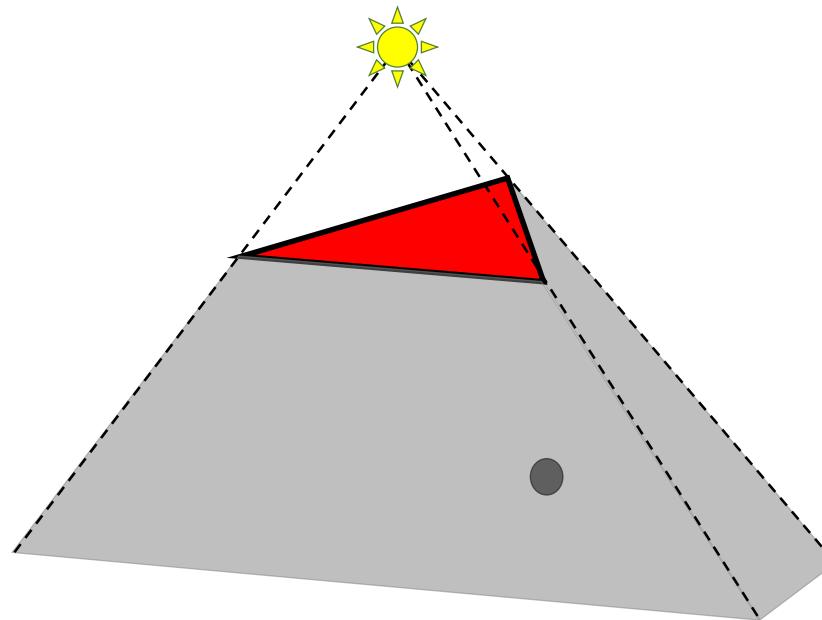
- Concept
  - Create volumes of “space in shadow” from each triangle
    - Each triangle creates 3 quads that extends to infinity



# Shadow Volumes. Concept

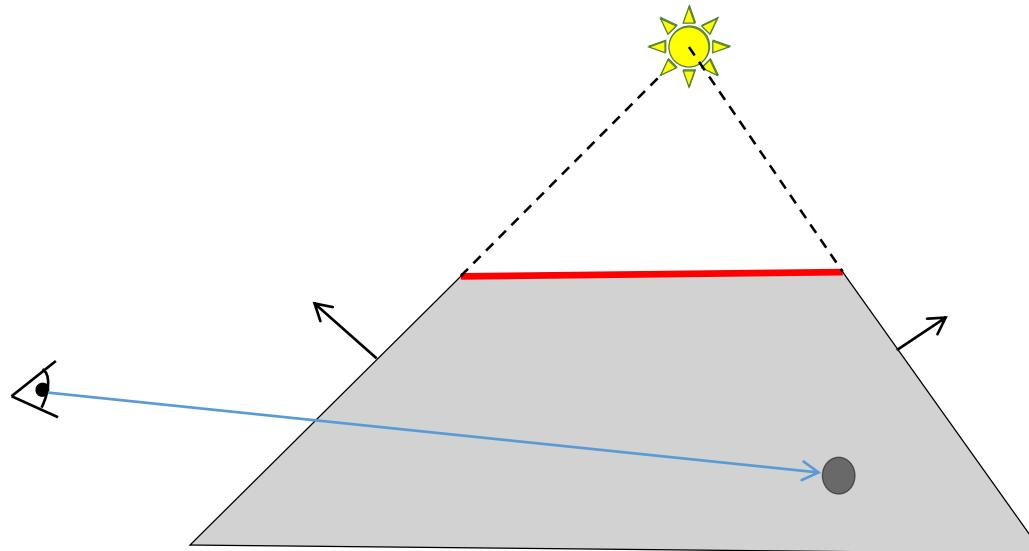
---

- To test a point, count how many shadow volumes it is located within. One or more means point is in shadow



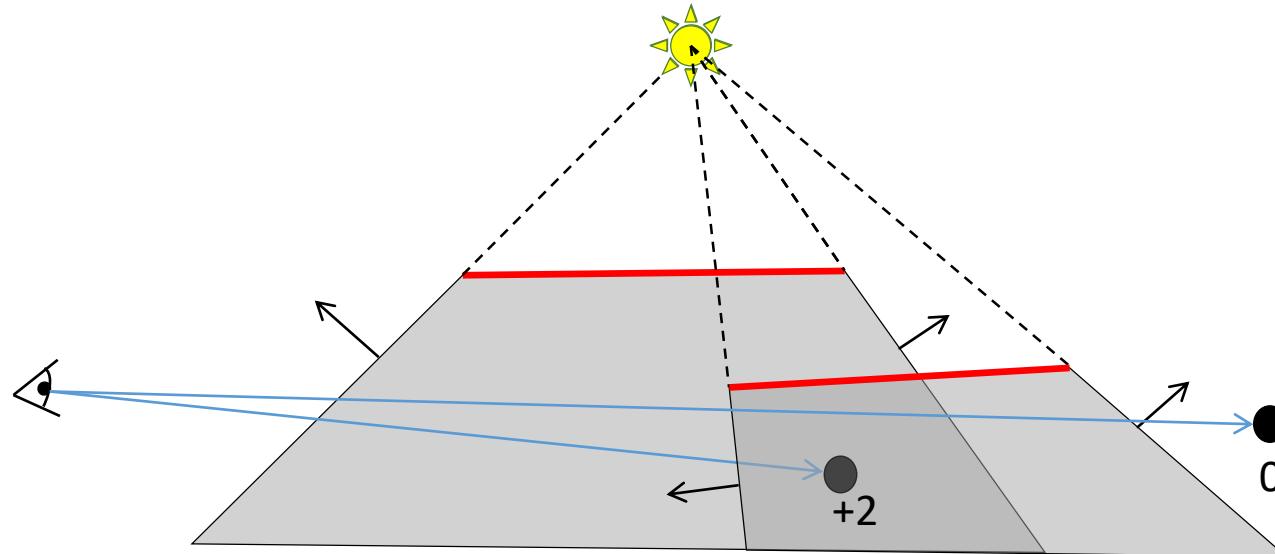
# Shadow Volumes. Concept

- To test a point, count how many shadow volumes it is located within. One or more means point is in shadow



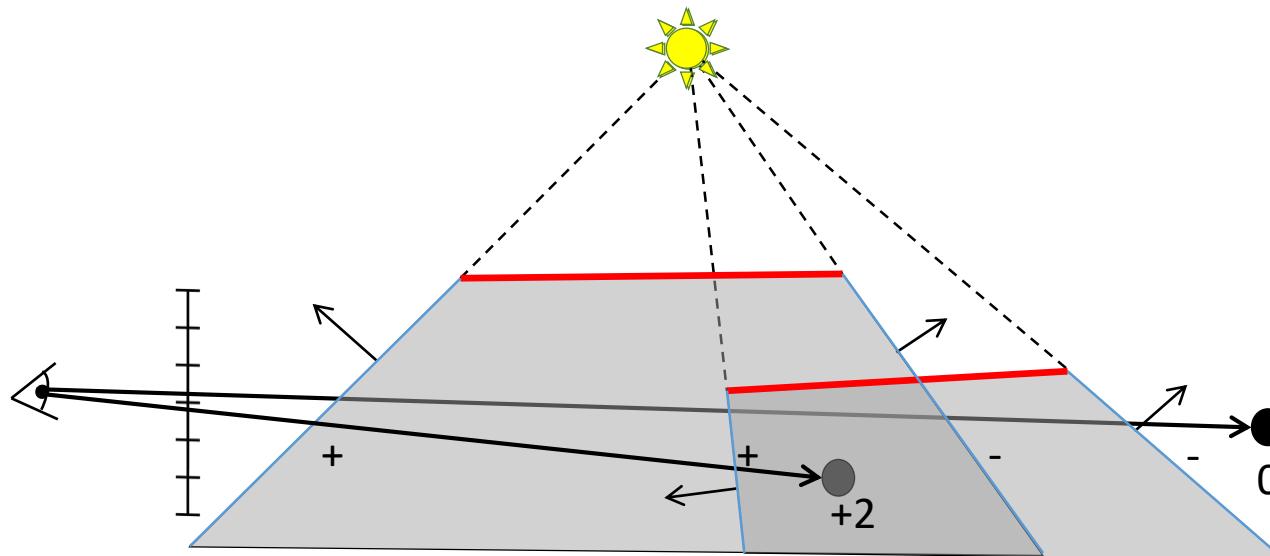
# Shadow Volumes. Concept

- To test a point, count how many shadow volumes it is located within. One or more means point is in shadow



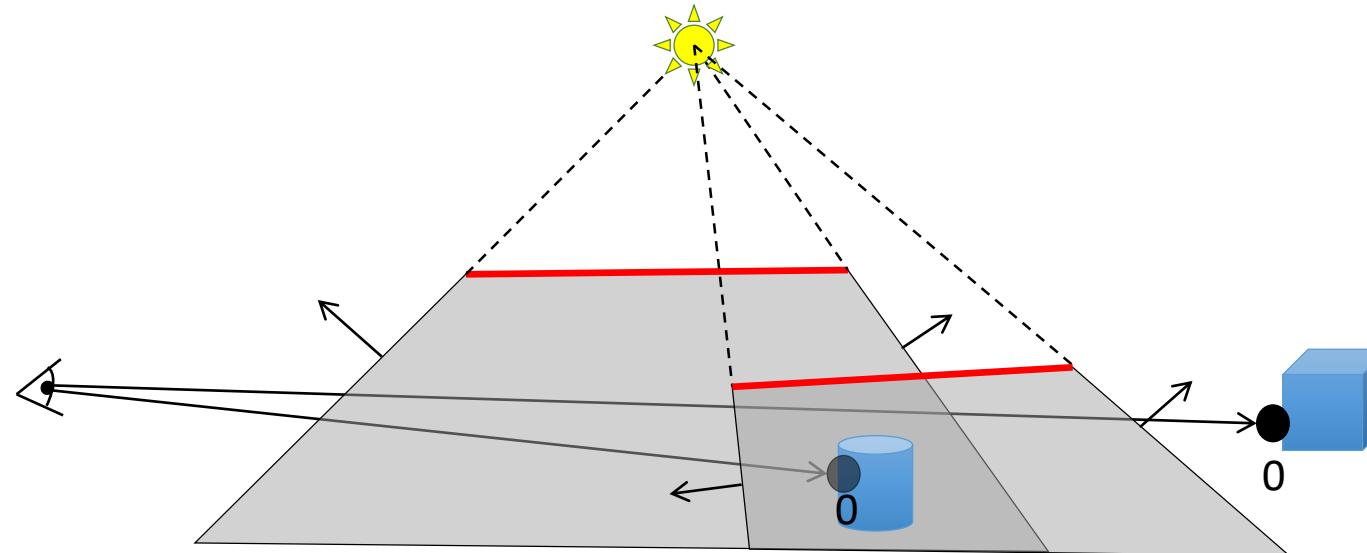
# Shadow Volumes. Concept

- A counter per pixel
- If we go through more frontfacing than backfacing polygons, then the point is in shadow



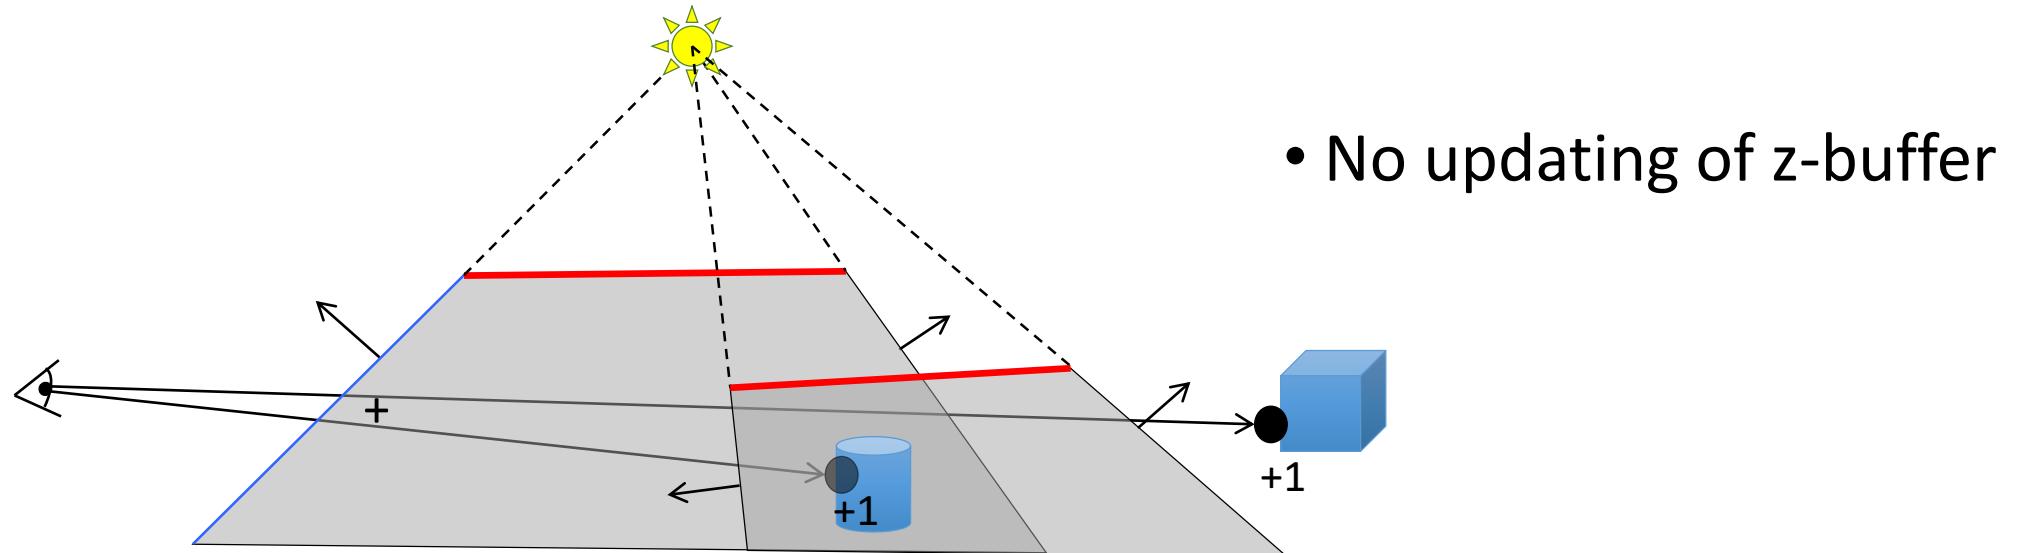
# Shadow Volumes. Stencil buffer

- Perform counting with the stencil buffer
  - Render front facing shadow quads to the stencil buffer
    - Inc stencil value, since those represents entering shadow volume
  - Render back facing shadow quads to the stencil buffer
    - Dec stencil value, since those represents exiting shadow volume



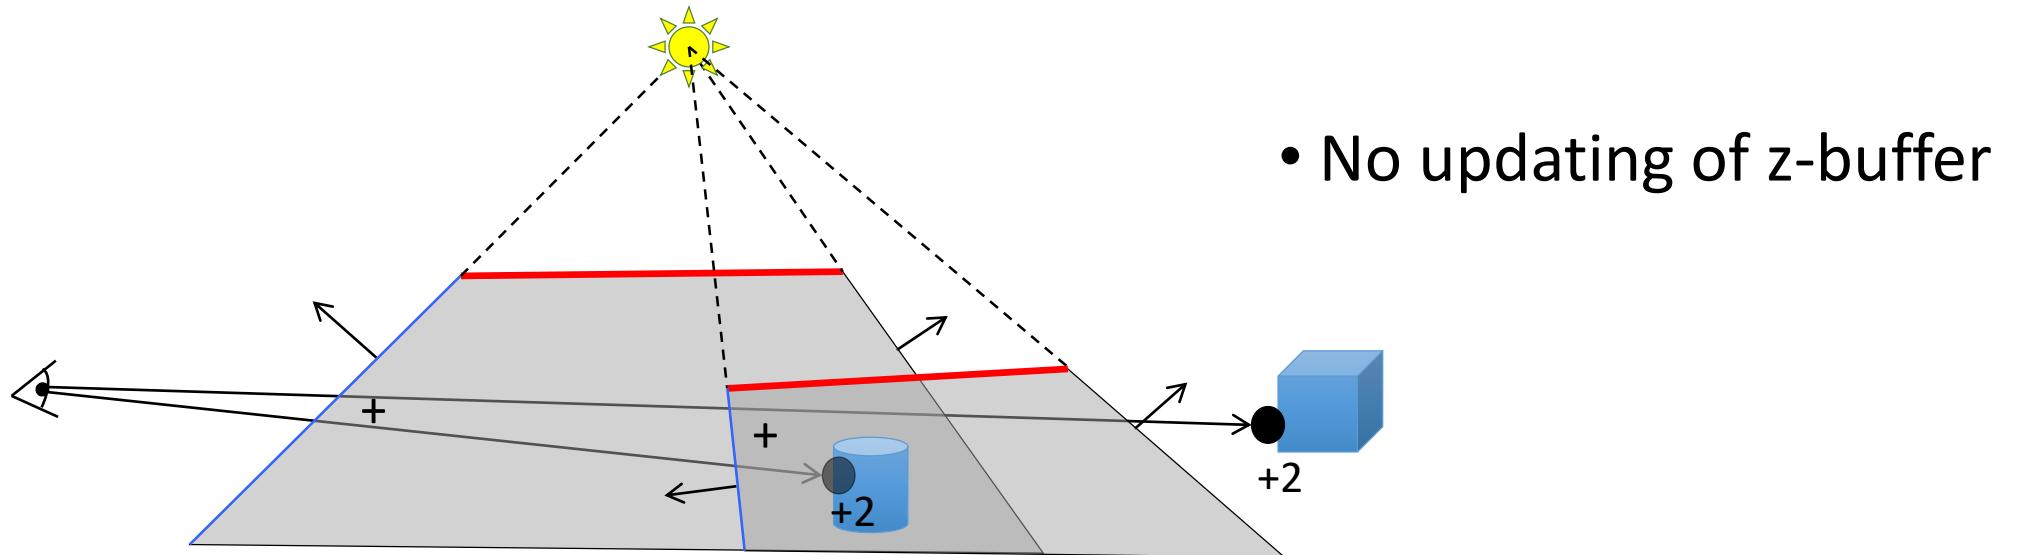
# Shadow Volumes. Stencil buffer

- Perform counting with the stencil buffer
  - Render front facing shadow quads to the stencil buffer
    - Inc stencil value, since those represents entering shadow volume
  - Render back facing shadow quads to the stencil buffer
    - Dec stencil value, since those represents exiting shadow volume



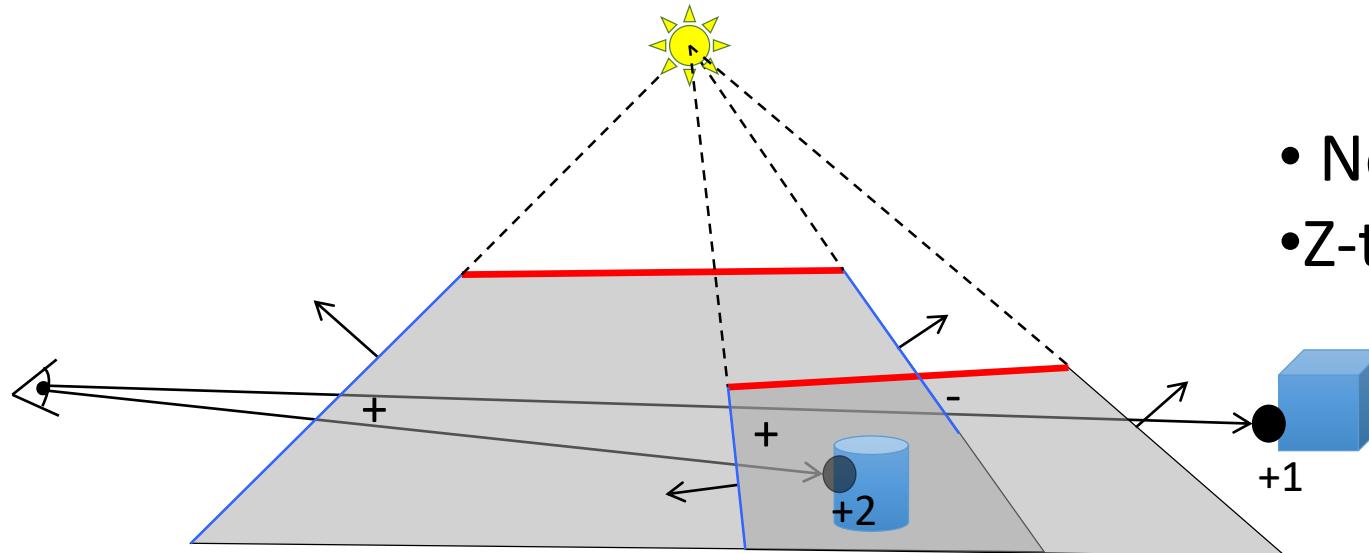
# Shadow Volumes. Stencil buffer

- Perform counting with the stencil buffer
  - Render front facing shadow quads to the stencil buffer
    - Inc stencil value, since those represents entering shadow volume
  - Render back facing shadow quads to the stencil buffer
    - Dec stencil value, since those represents exiting shadow volume



# Shadow Volumes. Stencil buffer

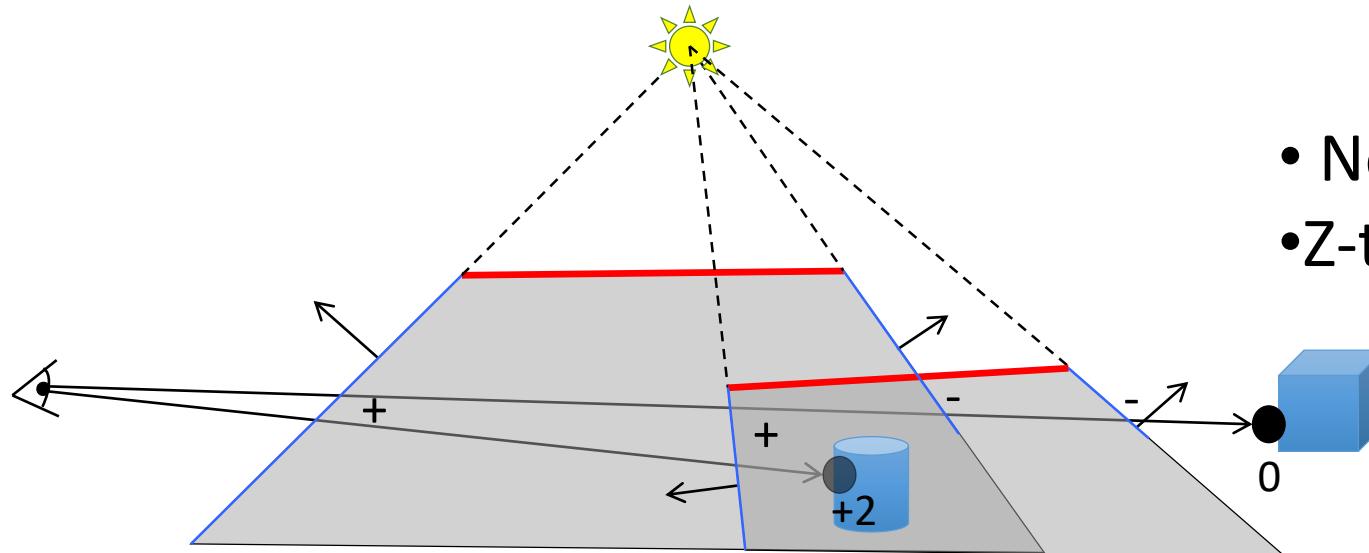
- Perform counting with the stencil buffer
  - Render front facing shadow quads to the stencil buffer
    - Inc stencil value, since those represents entering shadow volume
  - Render back facing shadow quads to the stencil buffer
    - Dec stencil value, since those represents exiting shadow volume



- No updating of z-buffer
- Z-test is enabled as usual

# Shadow Volumes. Stencil buffer

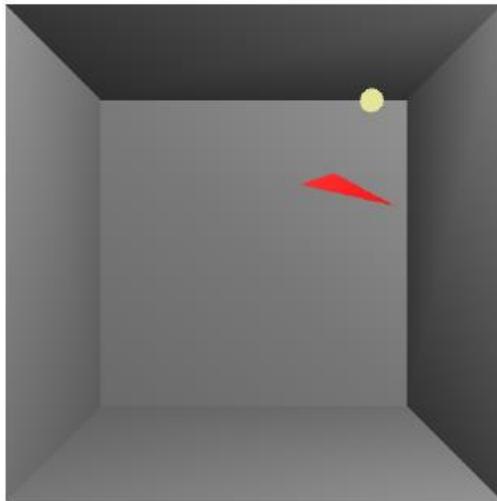
- Perform counting with the stencil buffer
  - Render front facing shadow quads to the stencil buffer
    - Inc stencil value, since those represents entering shadow volume
  - Render back facing shadow quads to the stencil buffer
    - Dec stencil value, since those represents exiting shadow volume



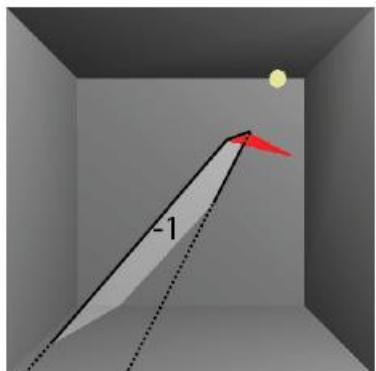
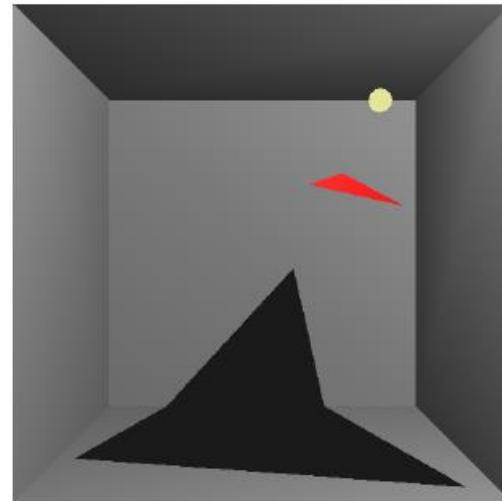
- No updating of z-buffer
- Z-test is enabled as usual

# Z-pass: How the stencil buffer is used

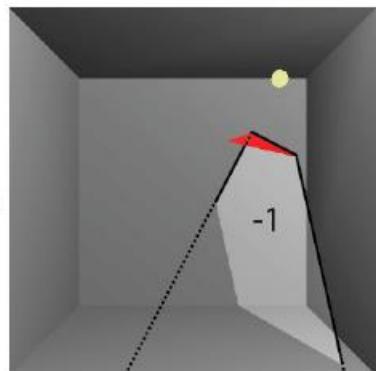
What we have...



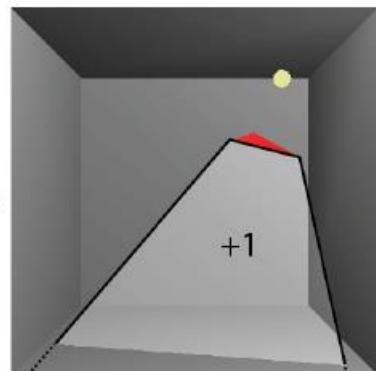
What we want...



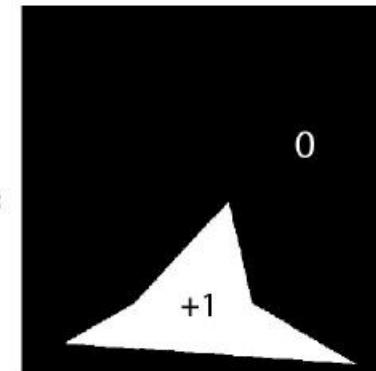
+



+



=



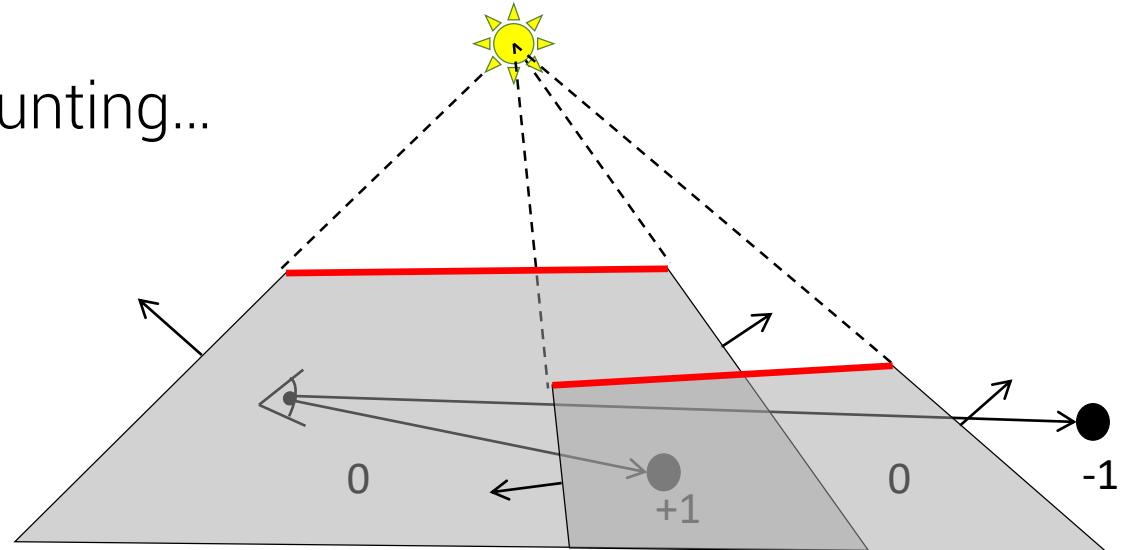
# Shadow Volumes. Stencil buffer

---

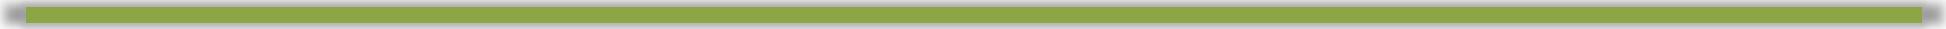
- A three pass process [Heidmann91]
  - 1st pass: Render ambient lighting
  - 2nd pass:
    - Draw to stencil buffer only
      - Turn off updating of z-buffer and writing to color buffer but still use standard depth test
      - Set stencil operation to
        - incrementing stencil buffer count for frontfacing shadow volume quads, and
        - decrementing stencil buffer count for backfacing shadow volume quads
      - use `glStencilOpSeparate(...)`
    - 3rd pass: Render diffuse and specular where stencil buffer is 0.

# Shadow Volumes. Eye Location Problem

- If the eye is located inside one or more shadow volumes, then the count will be wrong
- Solution:
  - Offset stencil buffer with the #shadow volumes that the eye is located within
  - Or modify the way we do the counting...



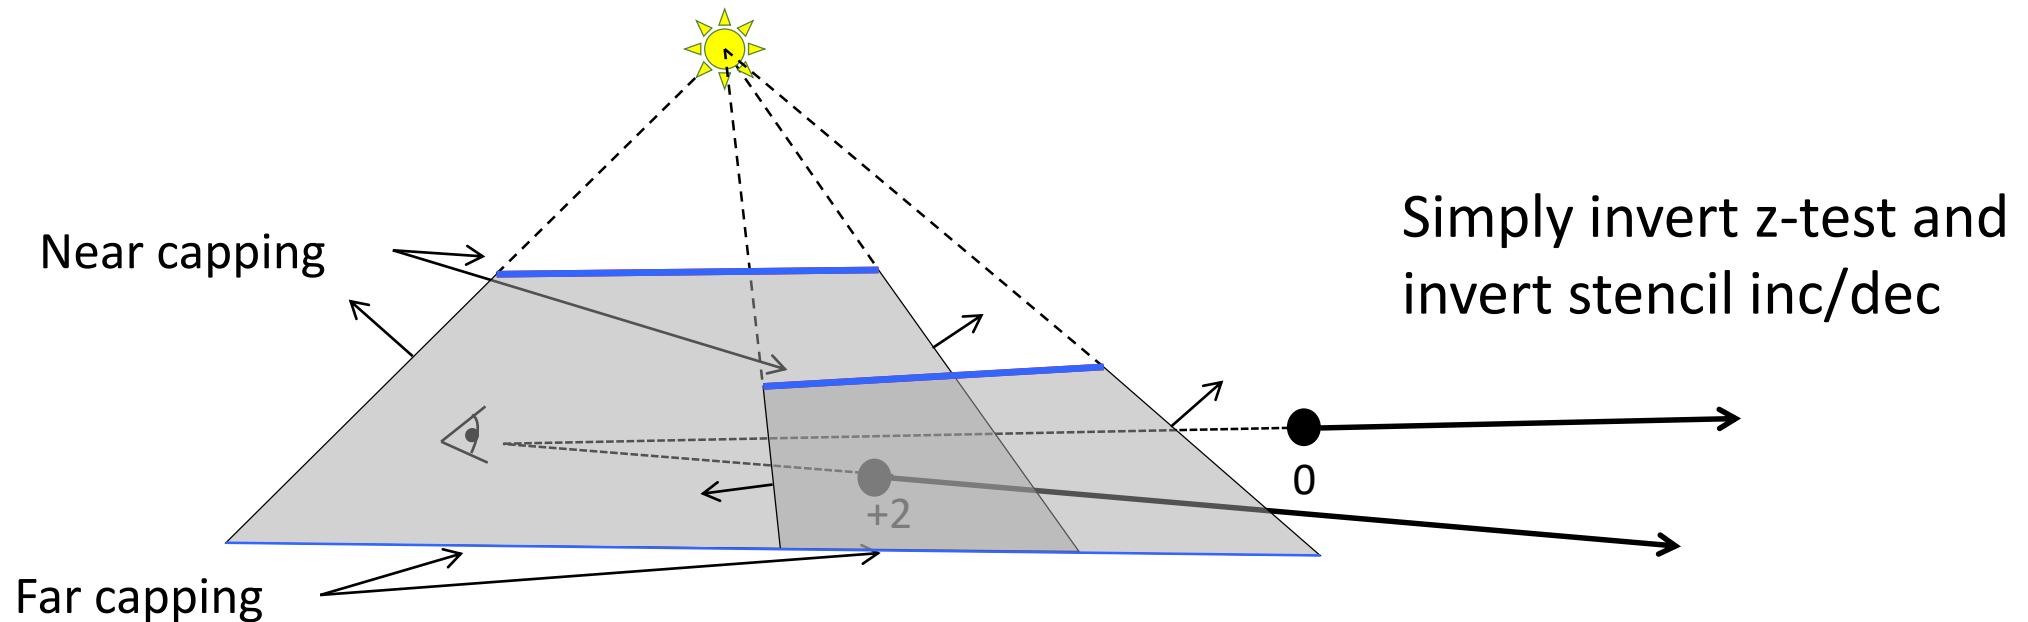
# Shadow Volumes. The Z-fail Algorithm



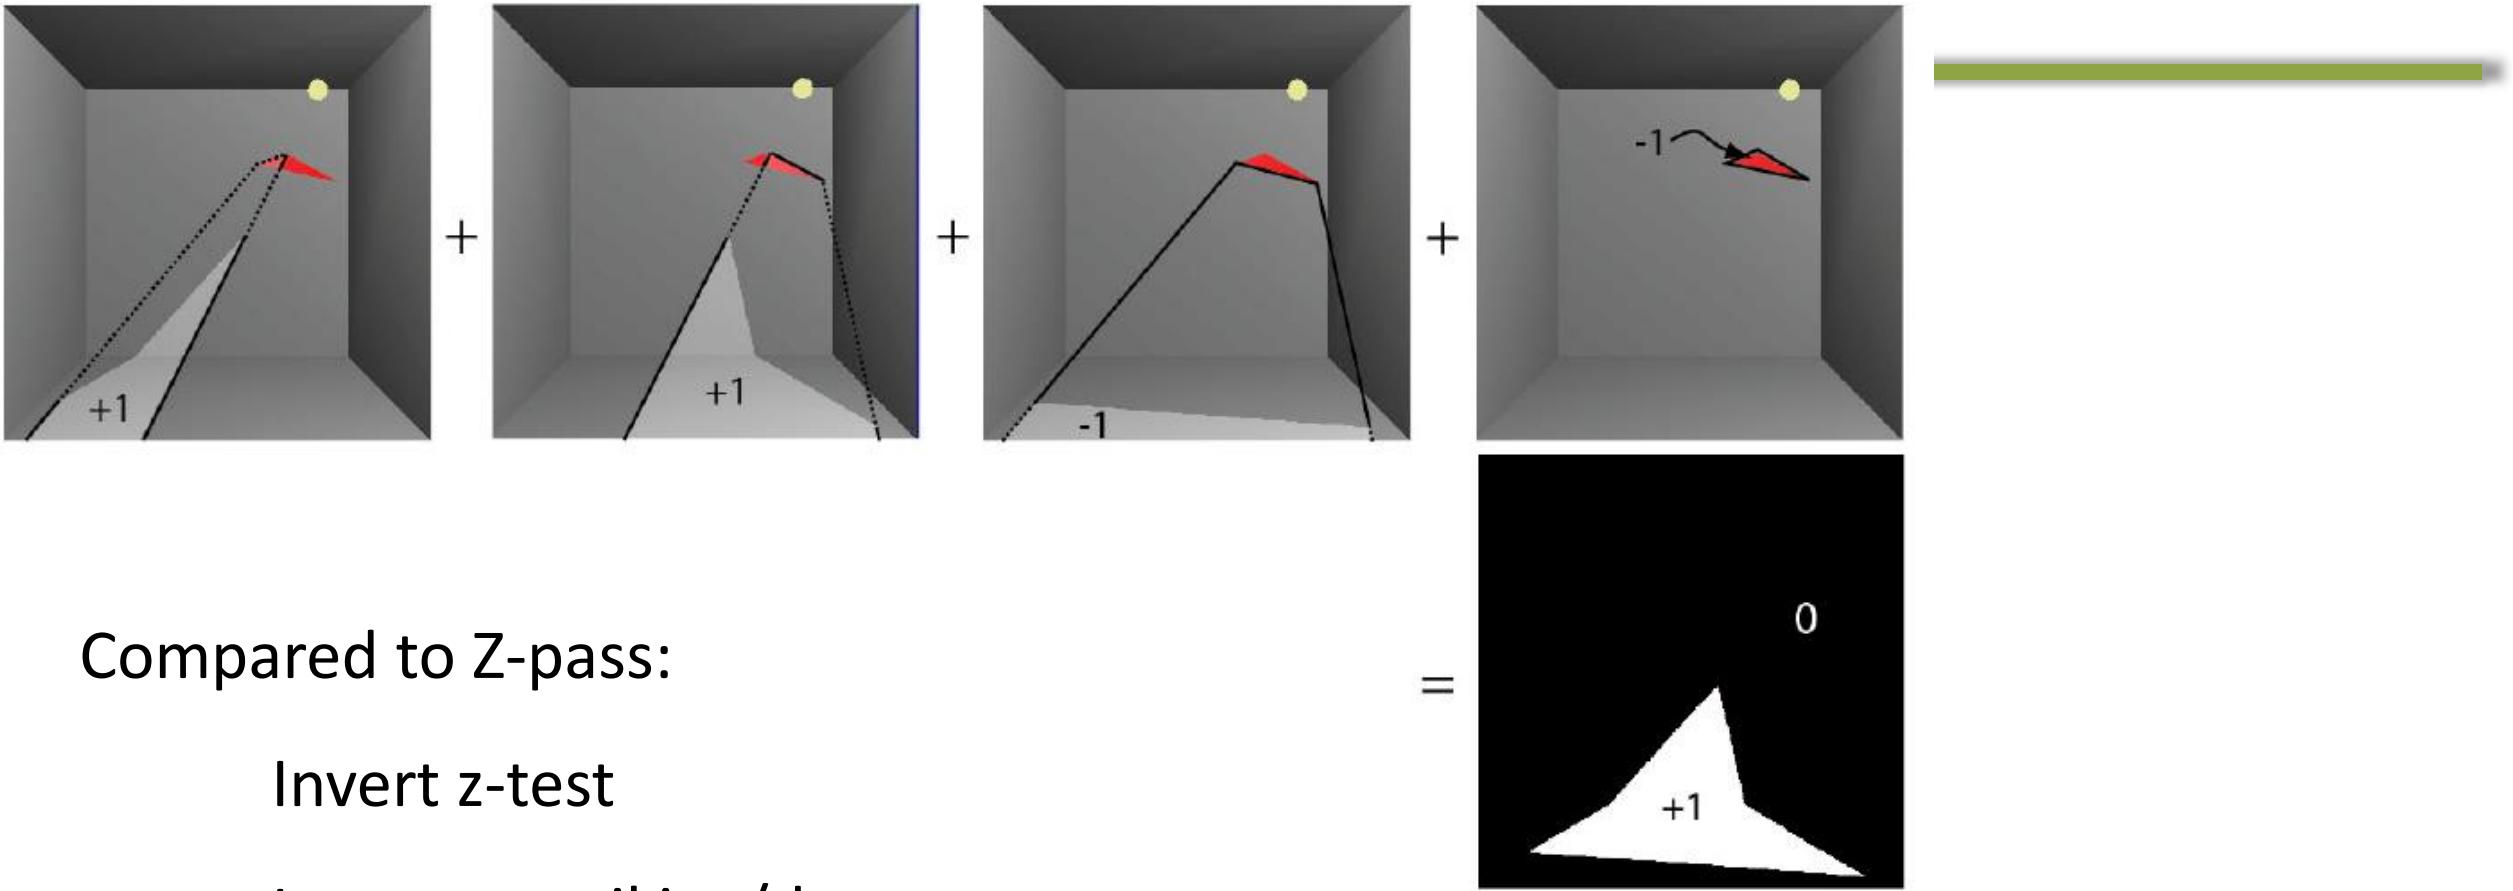
- By [Carmack00] and [Bilodeau and Songy 99]
  - “Carmacks Reverse”
- Count to infinity instead of to the eye
  - Add caps to the volume frustums
  - We can choose any reference location for the counting
  - A point in light avoids any offset
  - Infinity is always in light – if we cap the shadow volumes at infinity

# Shadow Volumes. The Z-fail Algorithm

- By [Carmack00] and [Bilodeau and Songy 99]
  - Count to infinity instead of to the eye



# Z-fail



Compared to Z-pass:

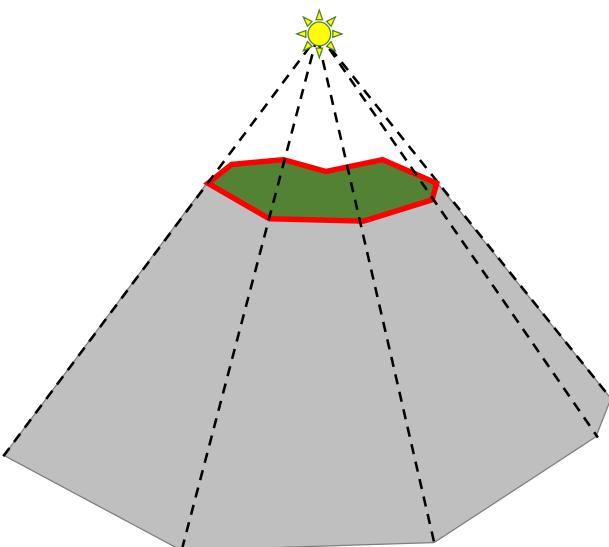
Invert z-test

Invert stencil inc/dec

I.e., count to infinity instead of from eye.

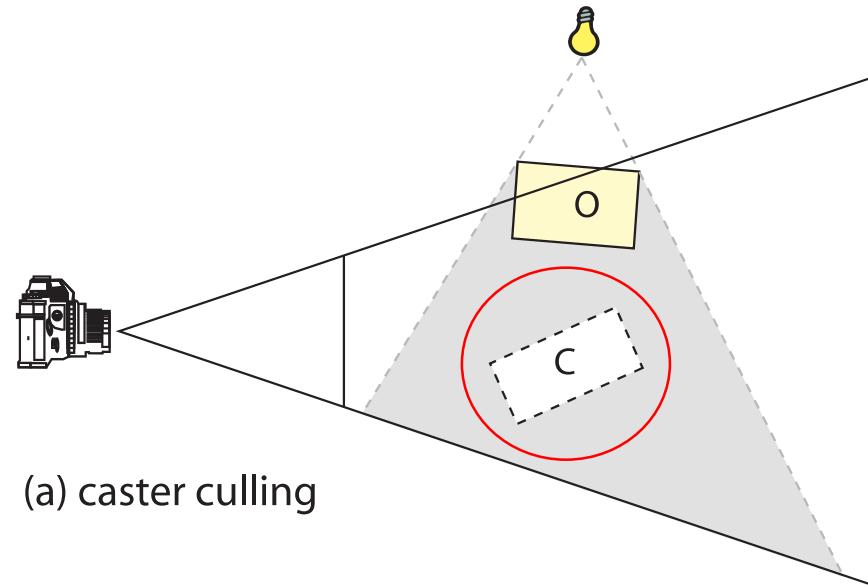
# Shadow Volumes - Summary

- Pros:
  - High quality
- Cons:
  - OVERDRAW



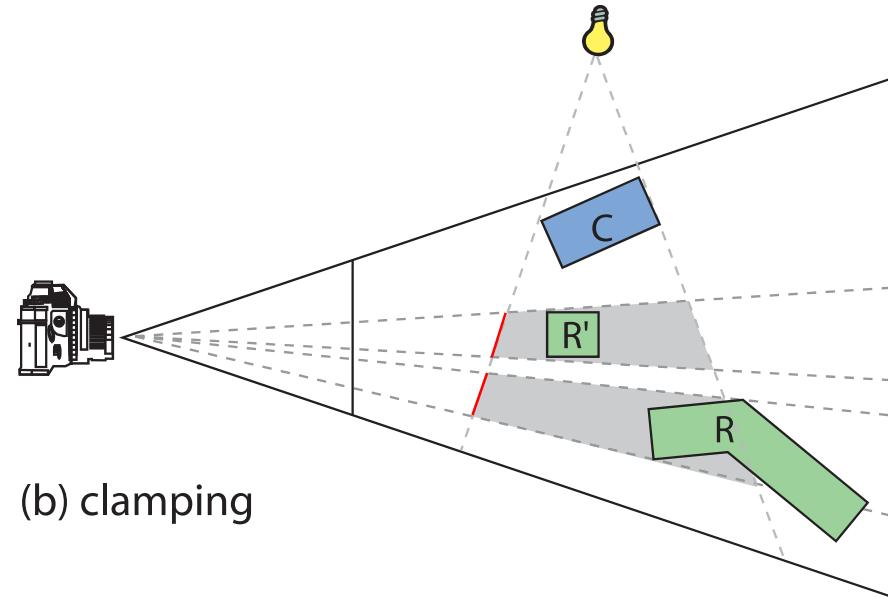
# Shadow Volumes. Culling and Clamping

- Culling of Shadow Volumes [Lloyd et al. 2004][Stich et al. 2007]
  - Culling of Shadow Casters if it is located totally within shadow
    - Tested against a shadow depth map



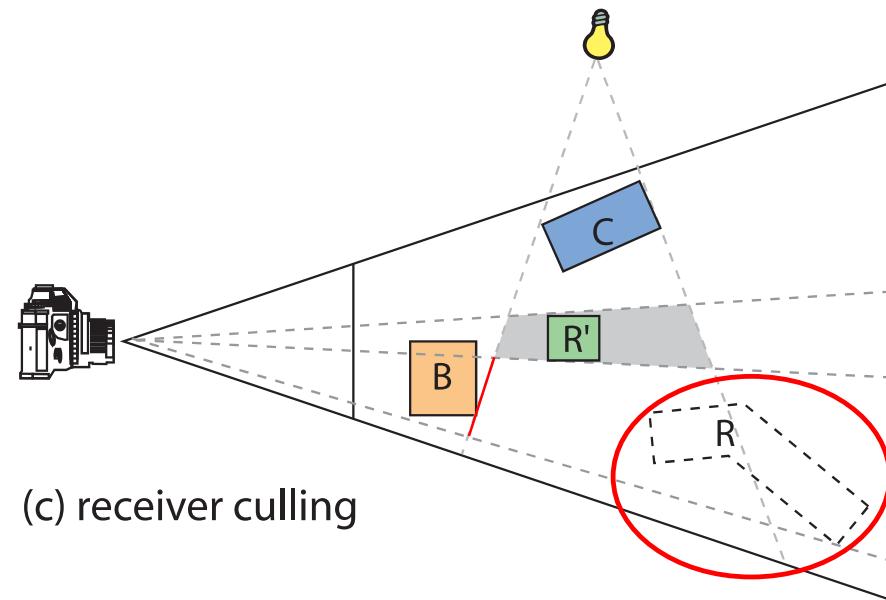
# Shadow Volumes. Culling and Clamping

- Clamping of Shadow Volumes  
[Lloyd et al. 2004] [Eisemann and Decoret 2006]
  - Idea: Only render parts of shadow quads that affects a shadow receiver
    - Tested against AABB around shadow receivers

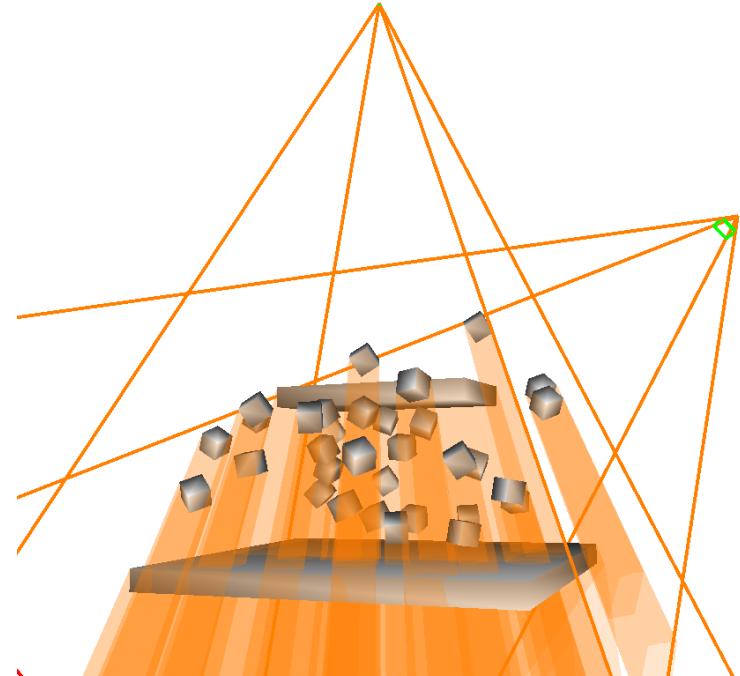


# Shadow Volumes. Culling and Clamping

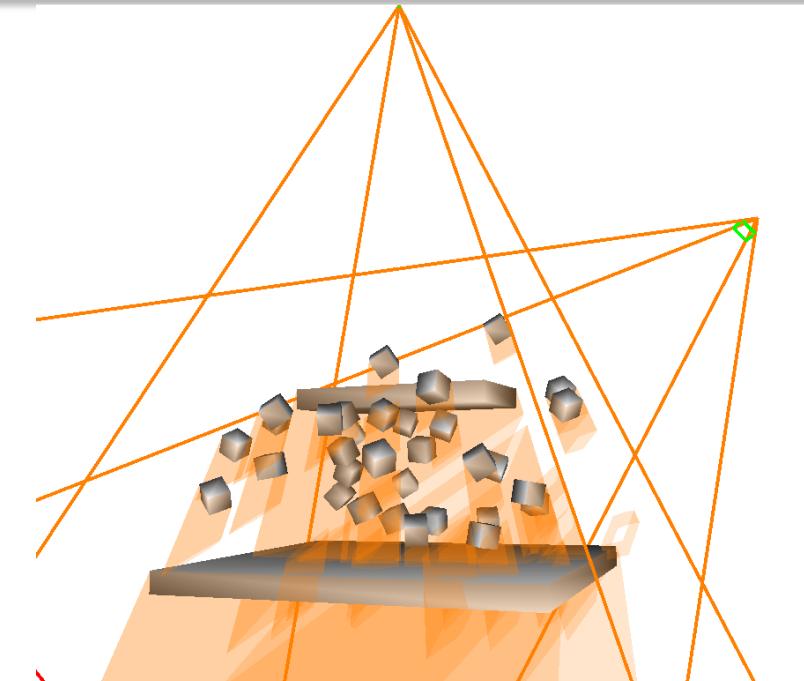
- Culling of Shadow Volumes  
[Lloyd et al. 2004] [Eisemann and Decoret 2006]
  - Receiver Culling
    - Idea: Cull part of shadow volumes where shadow receivers are not visible from the eye



# Shadow Volumes. Culling and Clamping



Shadow volumes

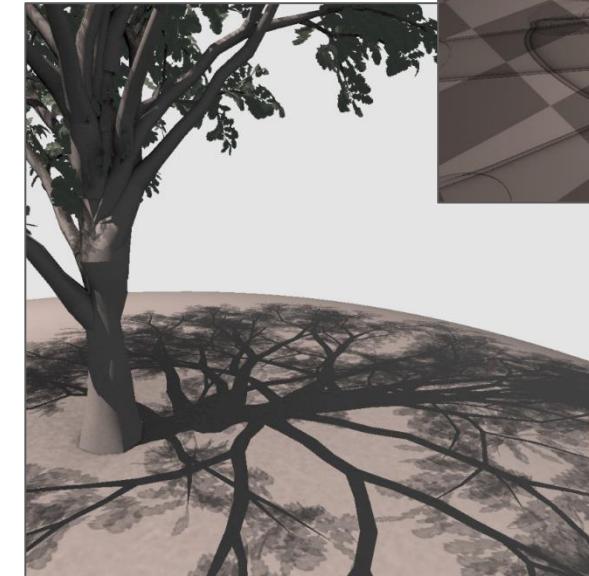
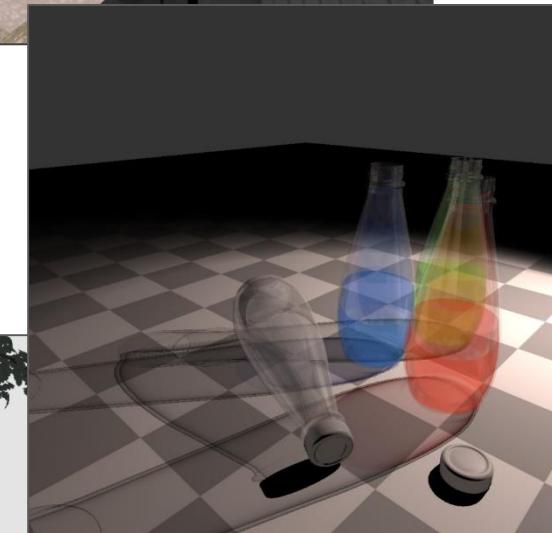


CC Shadow volumes

Illustrates reduced depth complexity when  
using Culling and Clamping

# Per-Triangle Shadow Volumes

- Per-triangle Shadow Volumes  
[Sintorn et al., 2011]
  - Idea:
    - Return to using one shadow volume per triangle
      - 100% robust + easiest.
    - Achieve efficiency by changing rasterization primitive from shadow quad to triangle shadow volume against hierarchical shadow buffer
    - Using CUDA rasterization
  - Allows for a much more general and robust algorithm



# Per-Triangle Shadow Volumes

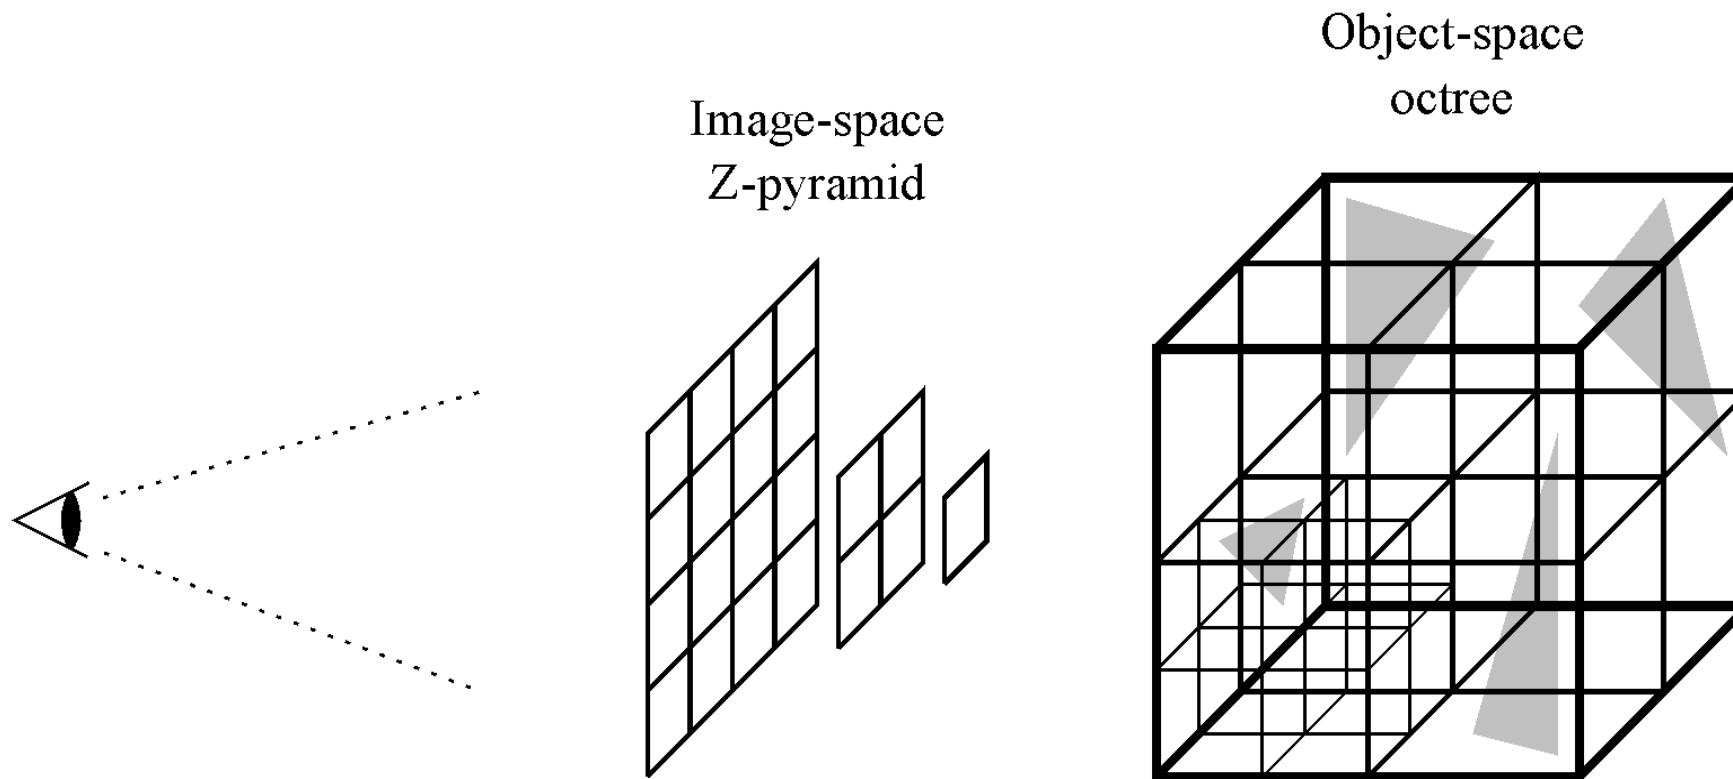
---

- Render the scene off-screen with ambient color, direct lighting
  - Store samples (color & depth)
    - Will evaluate visibility of each sample
      - Using a shadow buffer (one bit per sample)
    - Need to classify visibility wrt the faces of the frustum generated by each triangle
  - Use a two-level hierarchical depth buffer and a two-level hierarchical shadow buffer

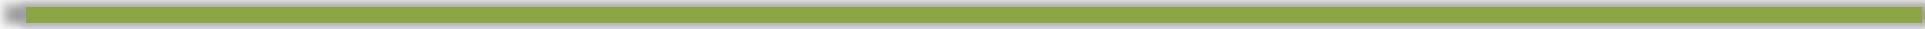
# Per-Triangle Shadow Volumes

---

- Hierarchical depth buffer stores minimum depth per tile
  - Accelerates depth tests by using higher levels
  - Can be cached (less memory required)

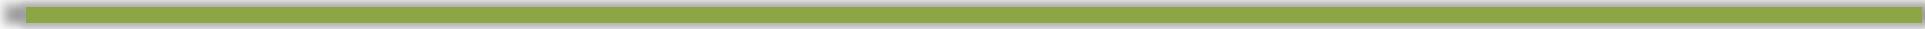


# Per-Triangle Shadow Volumes



- Upper level of hierarchical depth buffer contains min and max values
  - Define a bounding box (bounding frustum) named tile
    - Need to determine whether samples lie *inside* the box

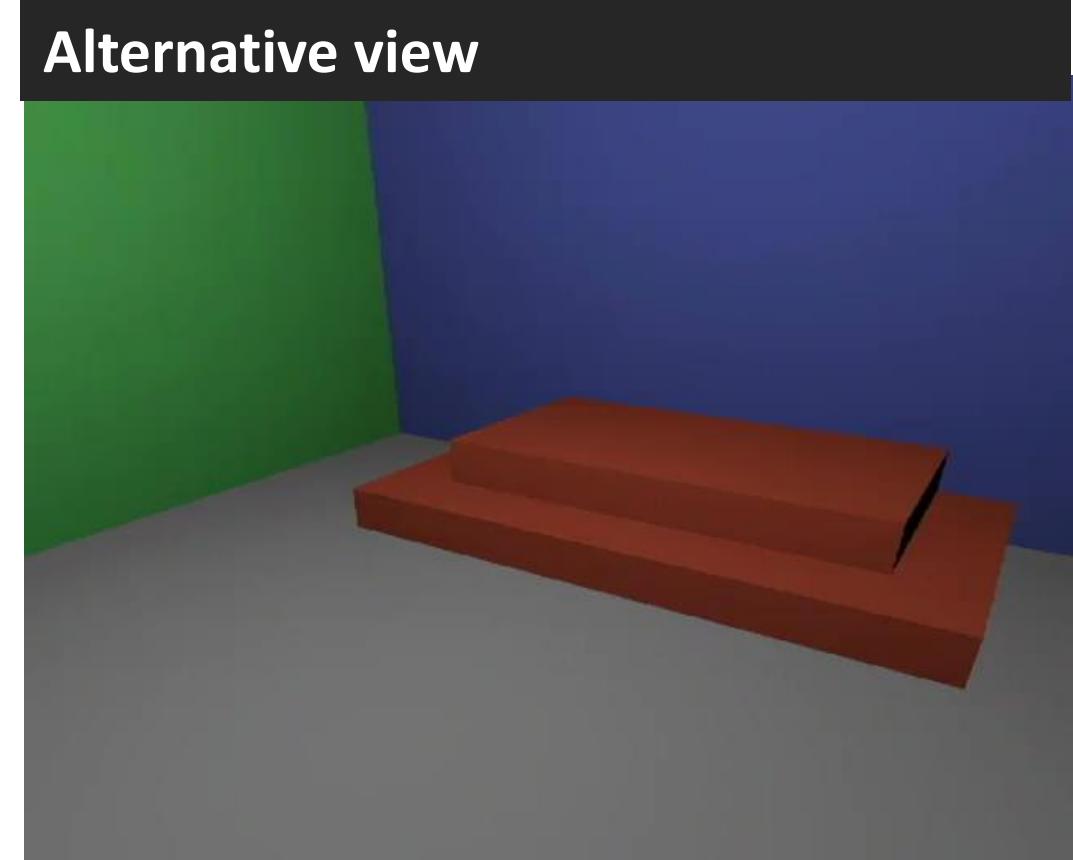
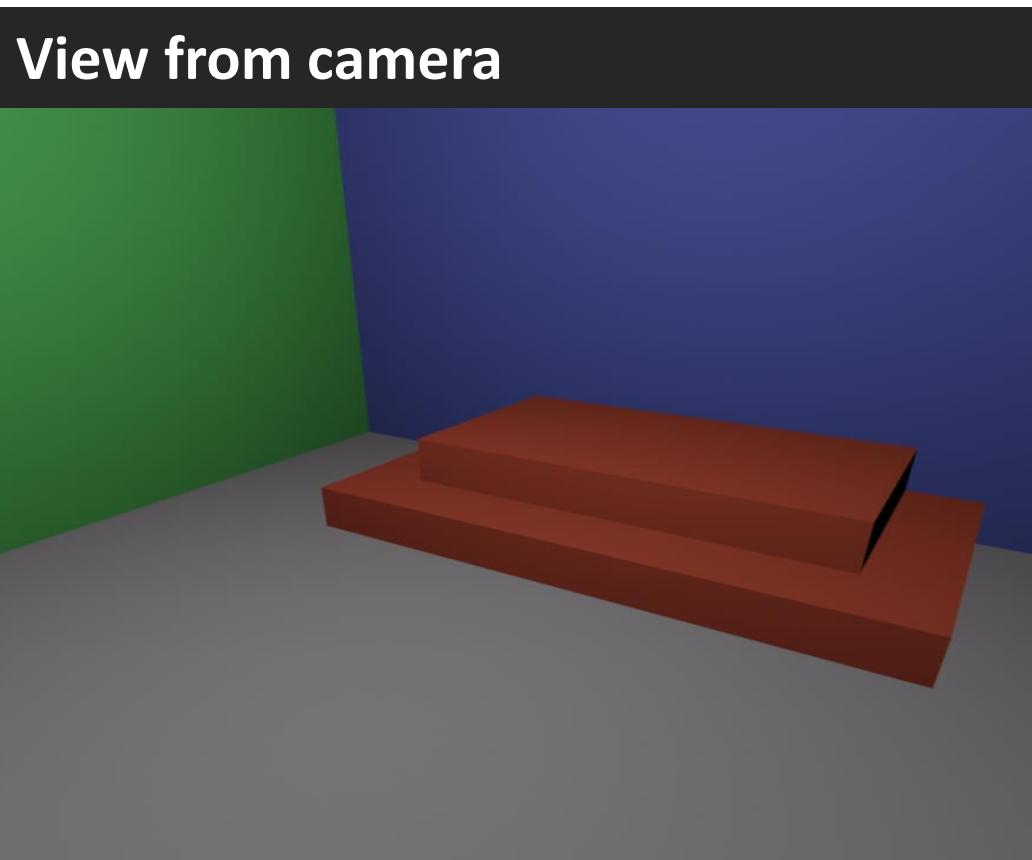
# Per-Triangle Shadow Volumes



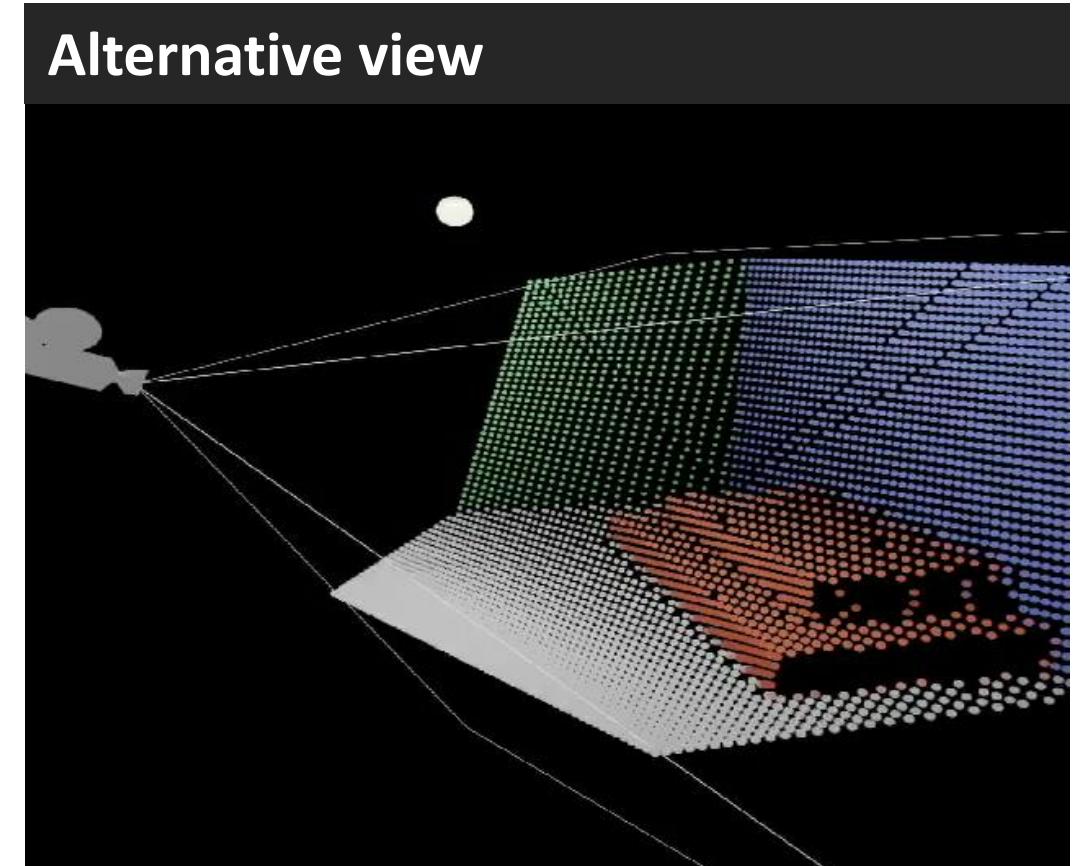
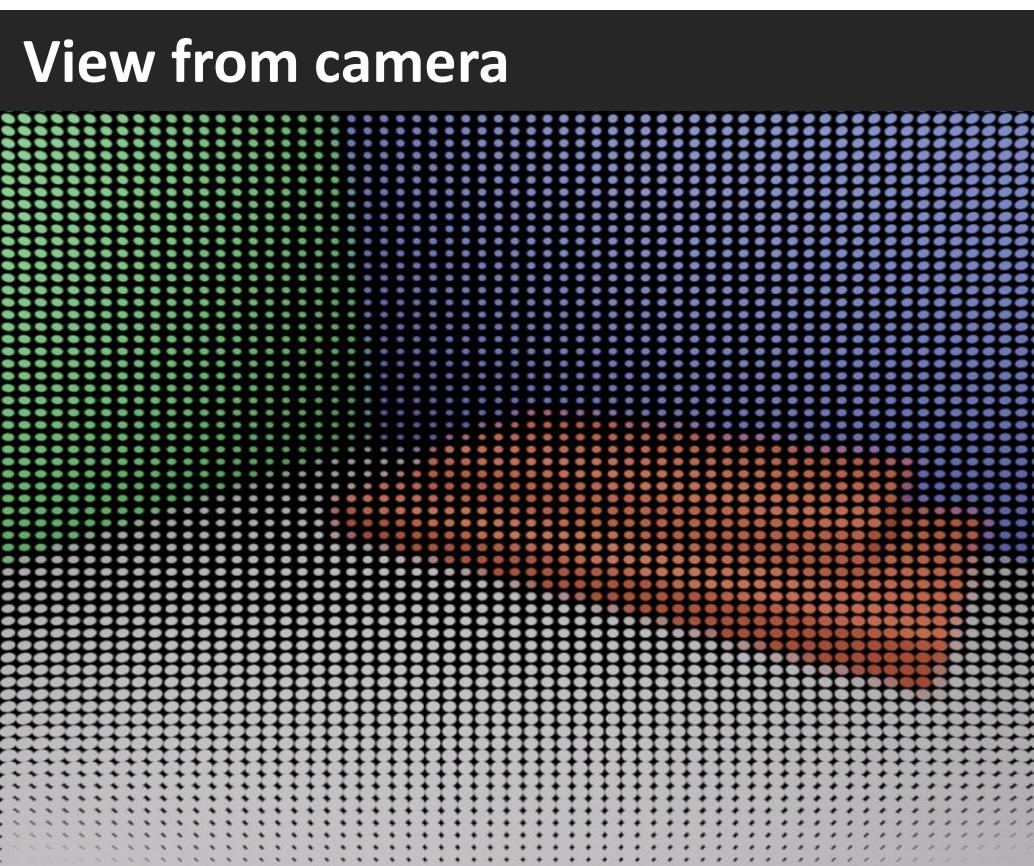
- Classification algorithm
  - Test if tile lies inside shadow volume frustum
    - True: all samples in shadow
      - Mark the shadow buffer bit
      - Abandon the tile
    - False: no samples in shadow
      - Discard the tile
    - Cannot be classified:
      - Go for the lower level
  - Implemented in CUDA

# Per-Triangle Shadow Volumes

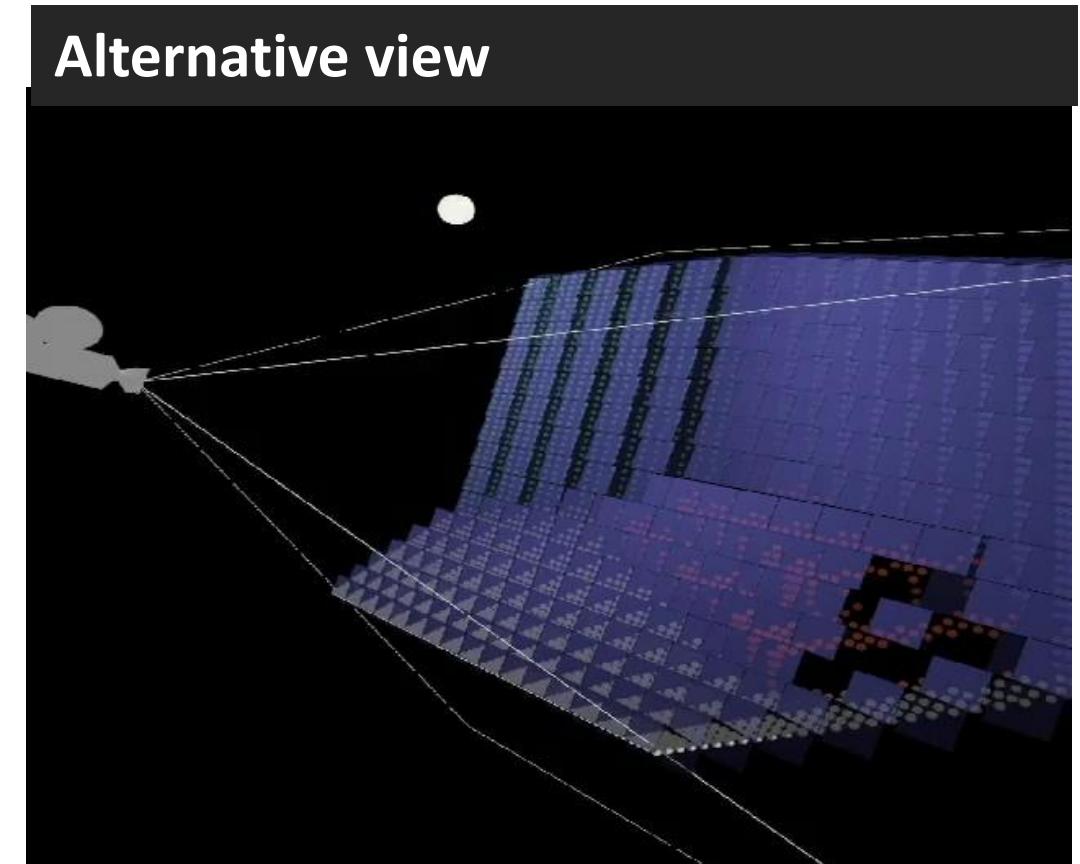
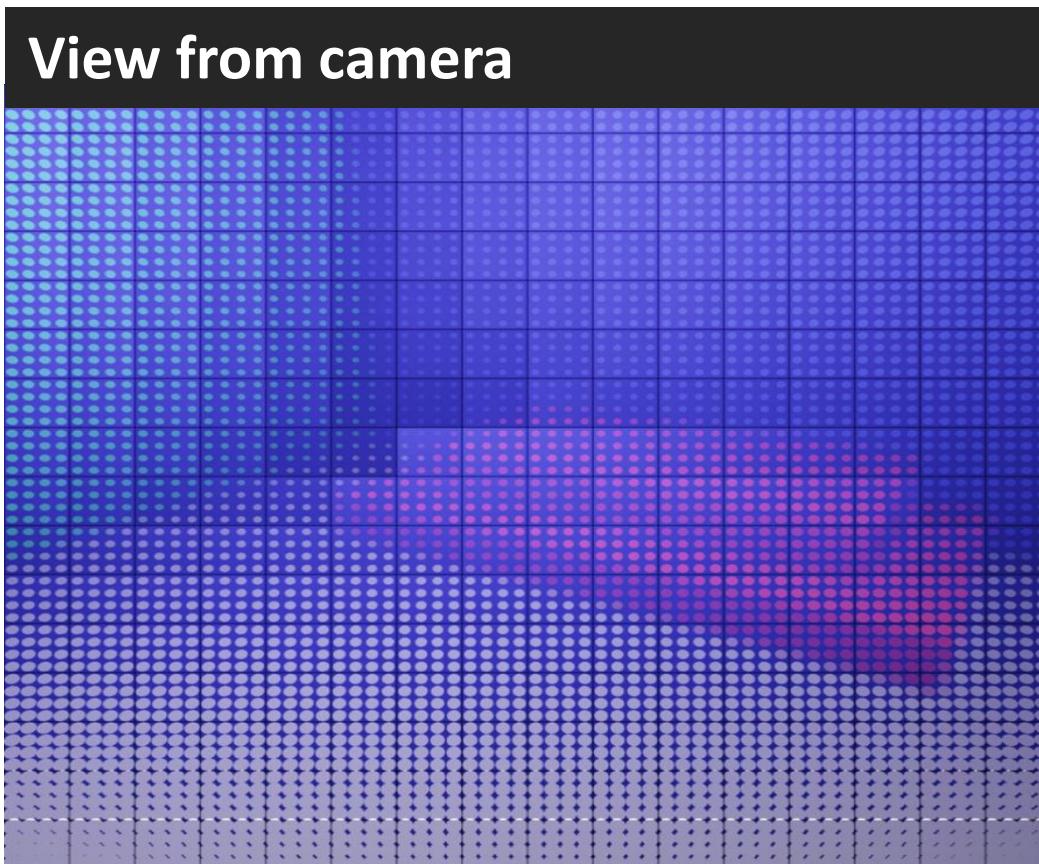
---



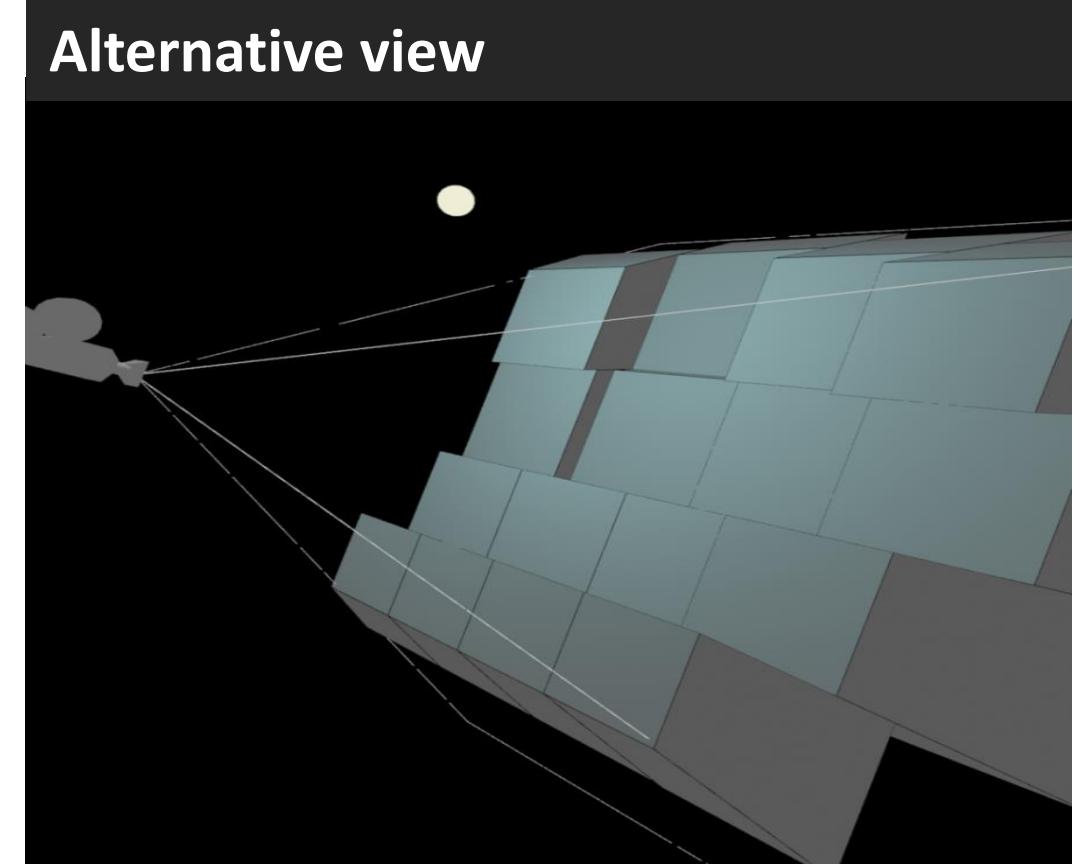
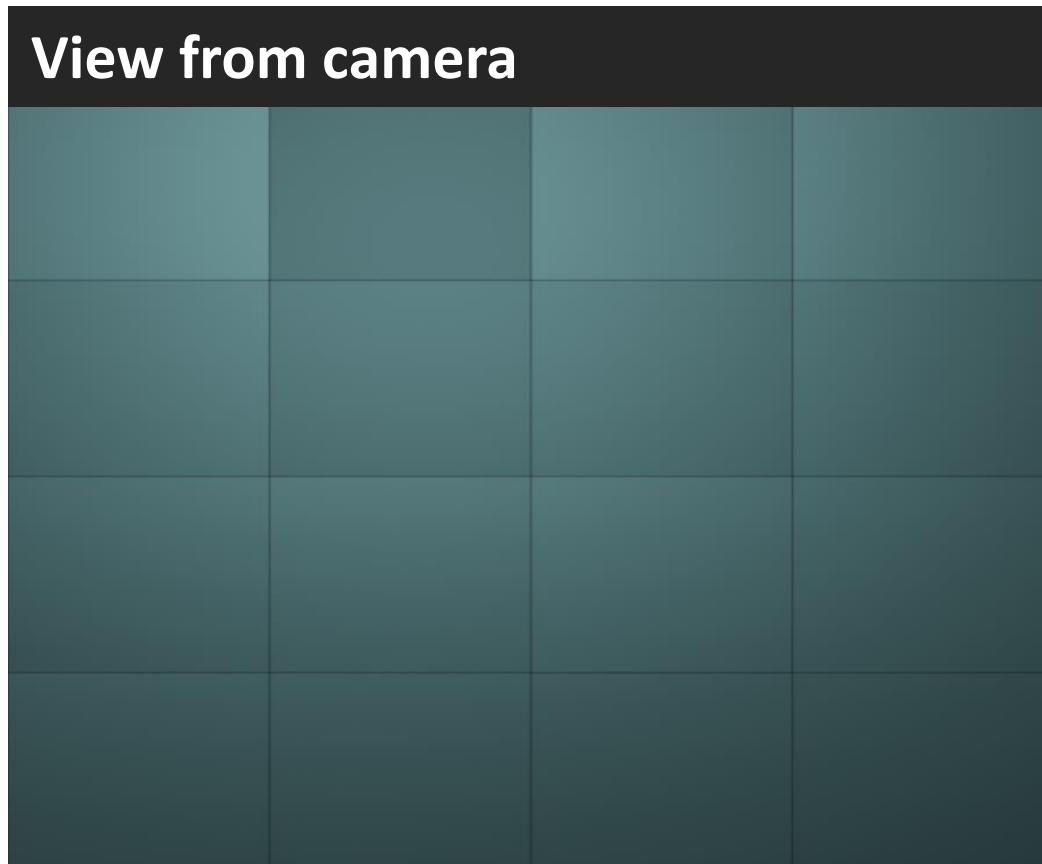
# Per-Triangle Shadow Volumes. Algorithm. Building the depth hierarchy



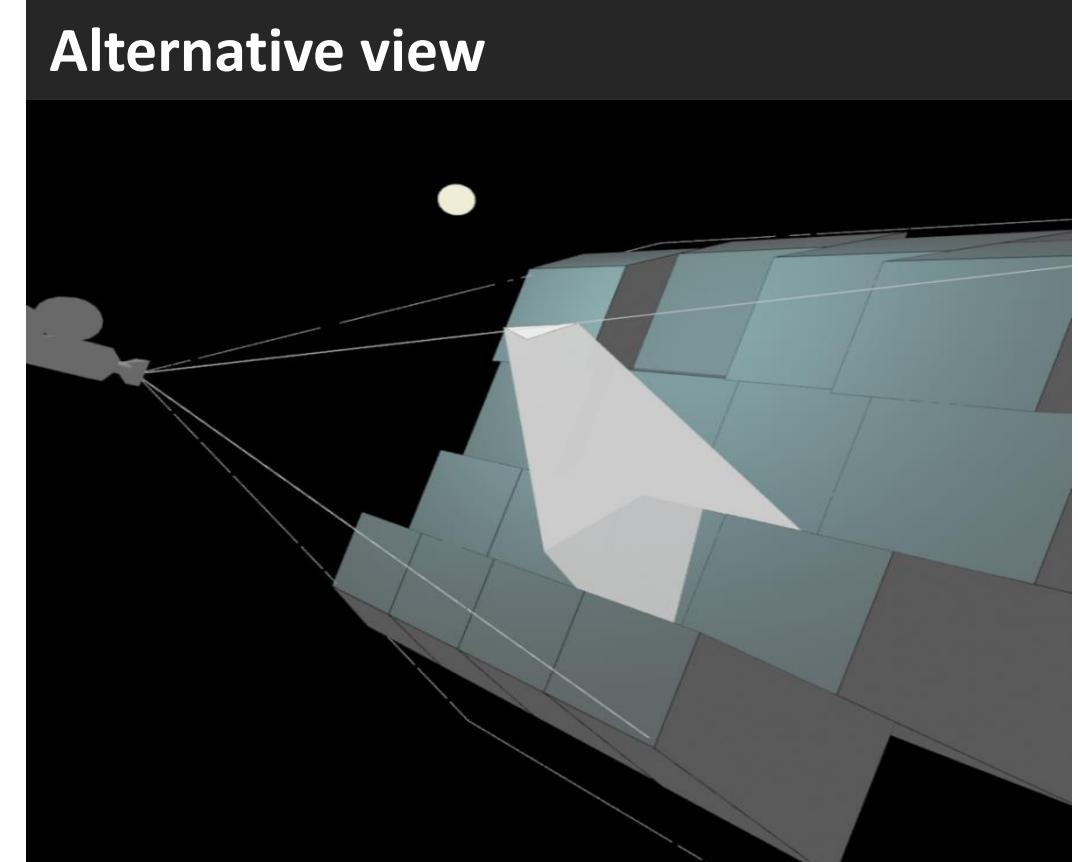
# Per-Triangle Shadow Volumes. Algorithm. Building the depth hierarchy



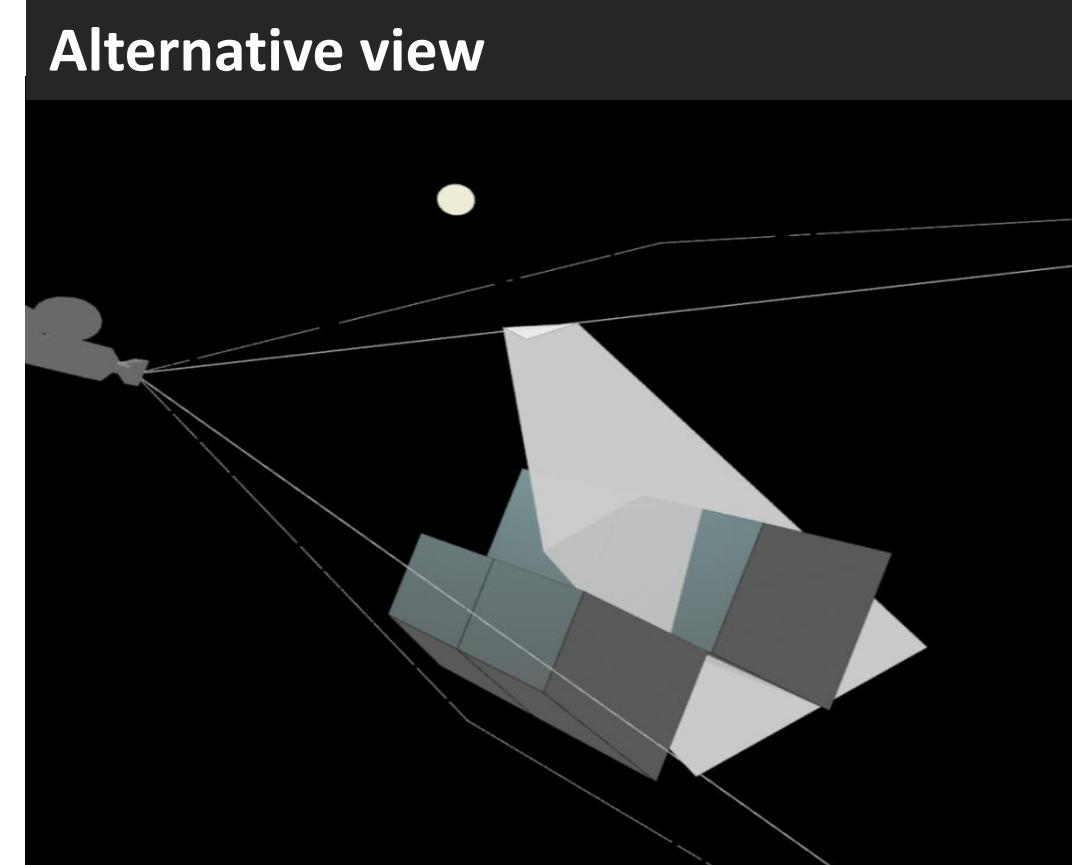
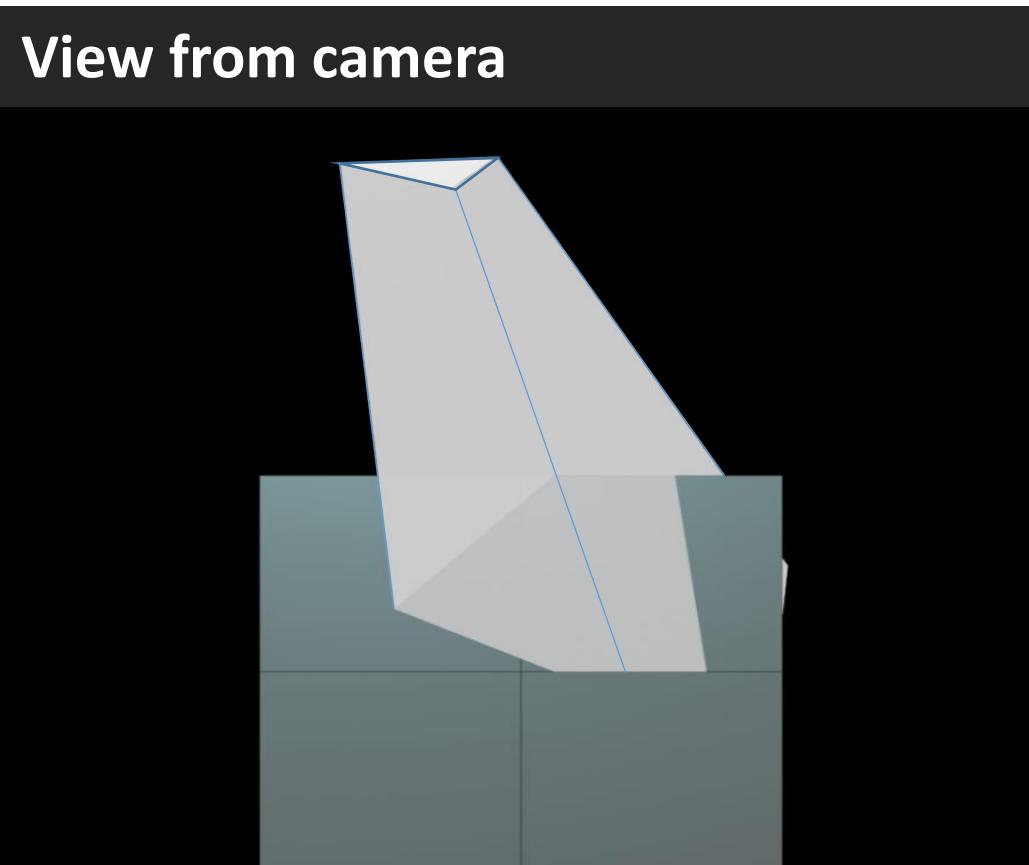
# Per-Triangle Shadow Volumes. Algorithm. Building the depth hierarchy



# Per-Triangle Shadow Volumes. Algorithm. Traversing the hierarchy

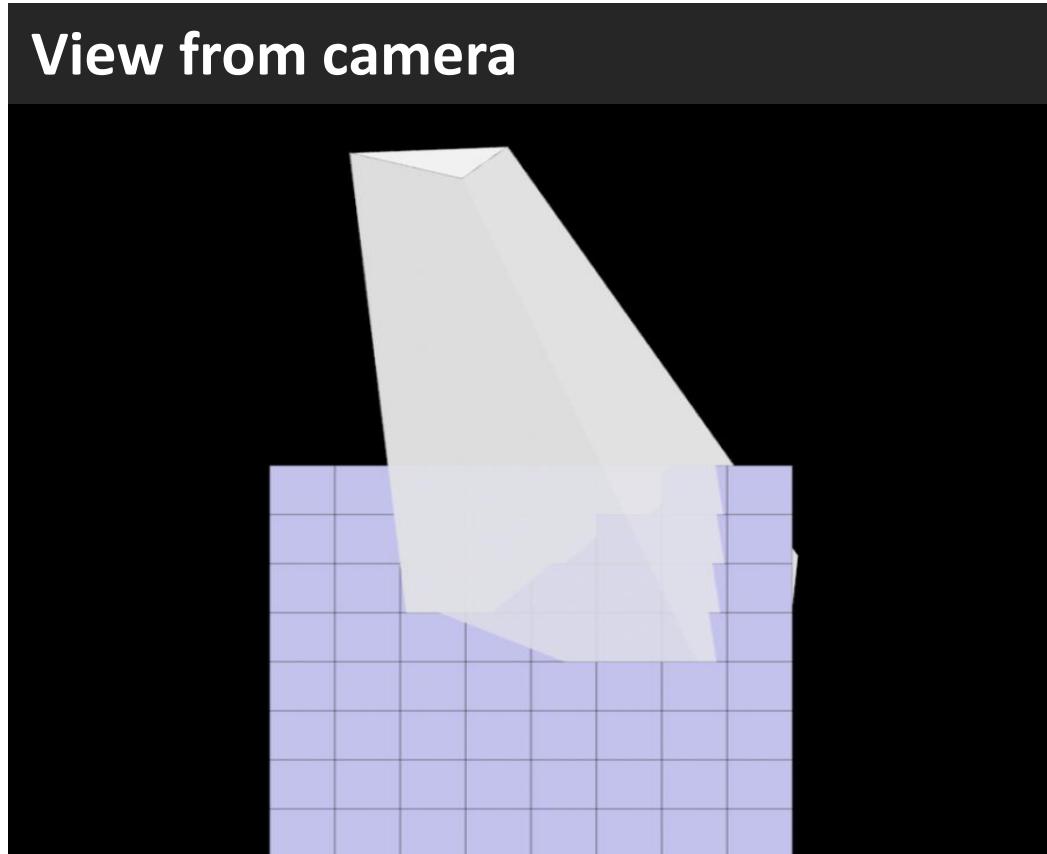


# Per-Triangle Shadow Volumes. Algorithm. Traversing the hierarchy – Cull 4x4 tiles

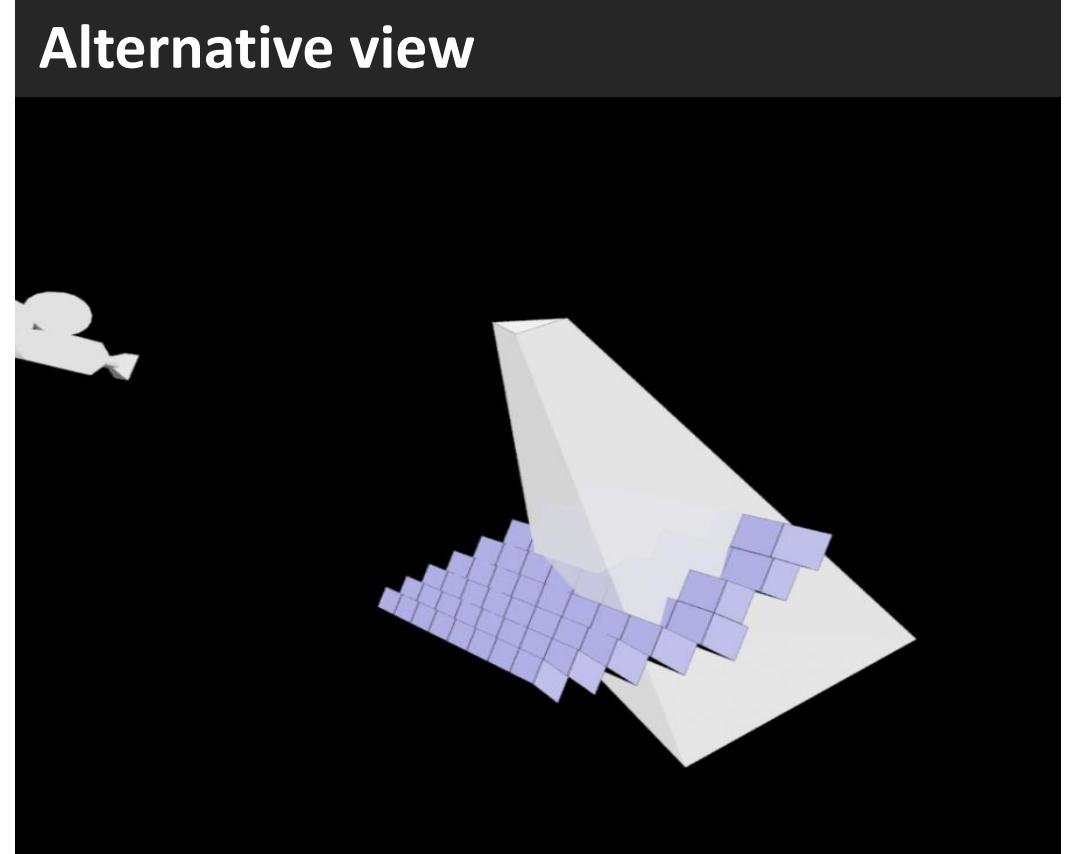


# Per-Triangle Shadow Volumes. Algorithm. Traversing the hierarchy – Cull 4x4 tiles

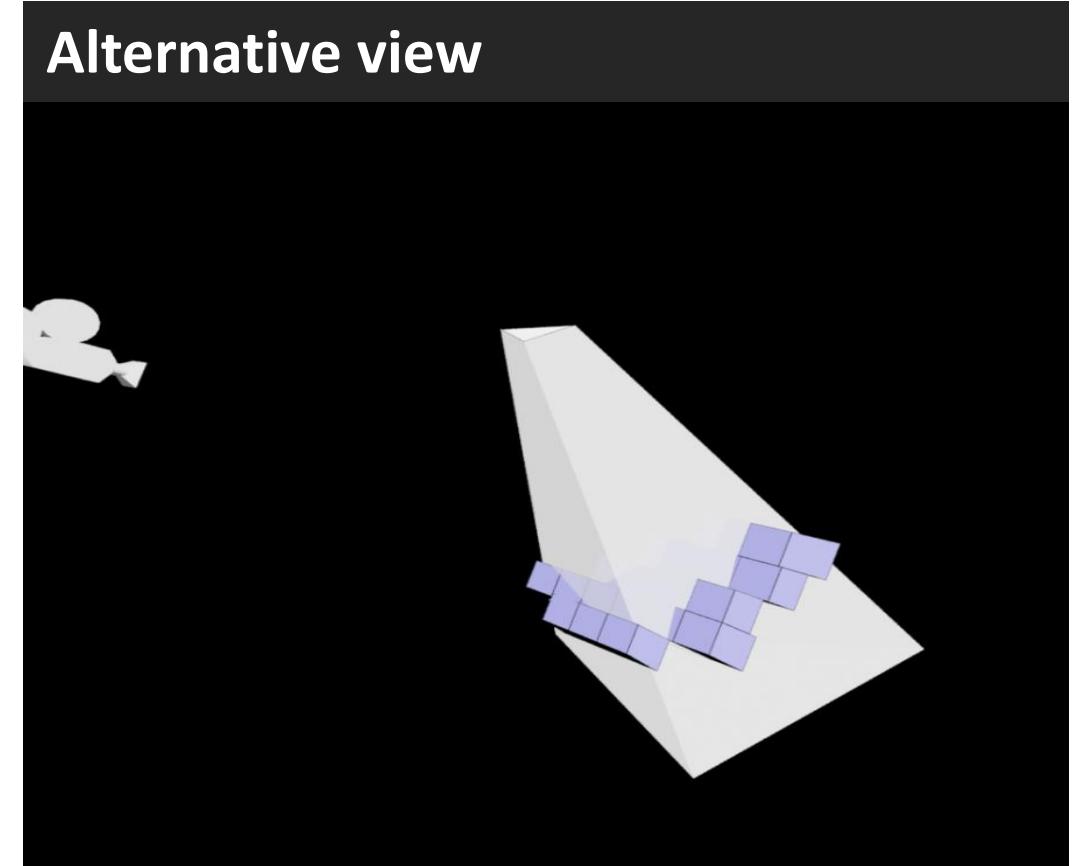
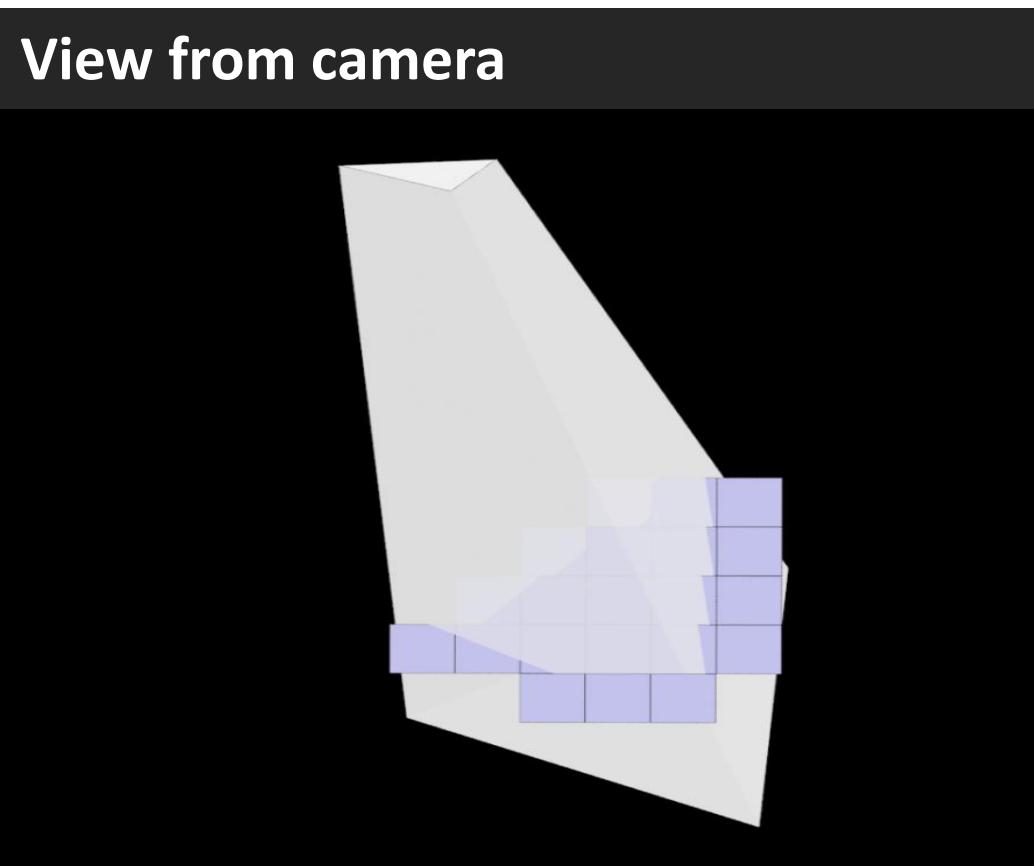
**View from camera**



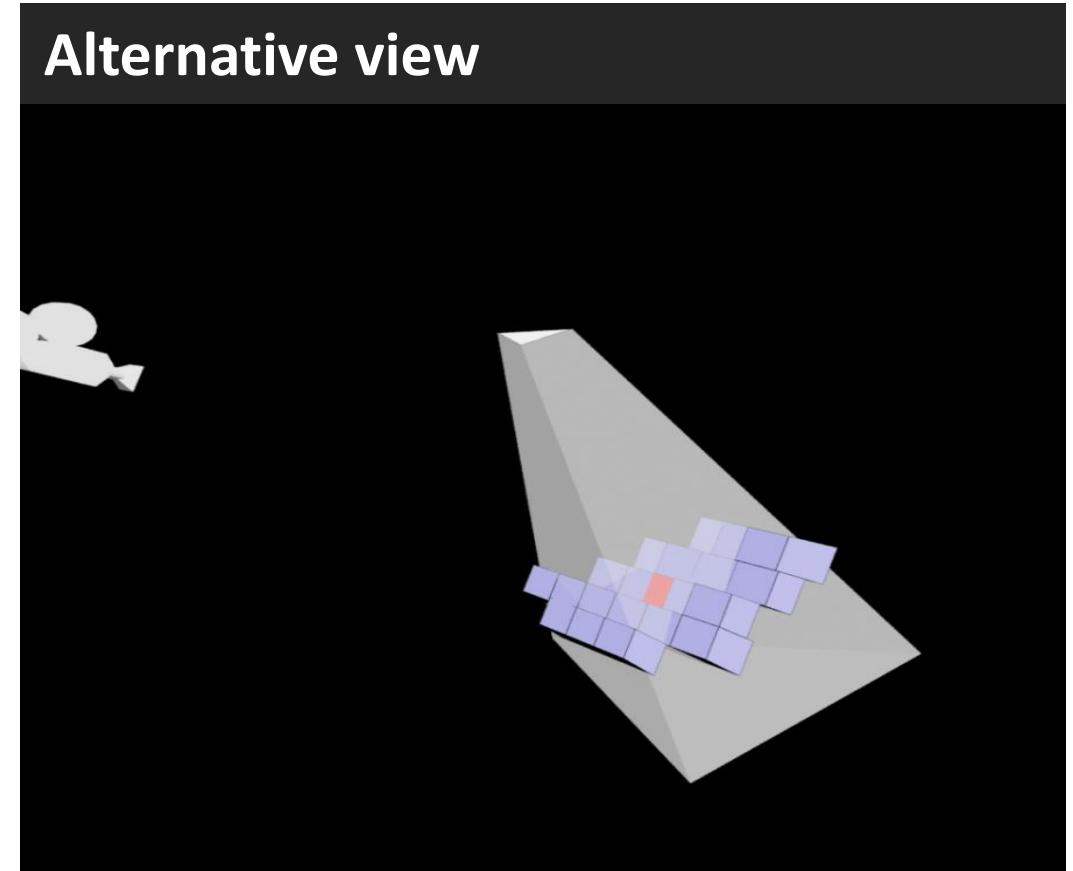
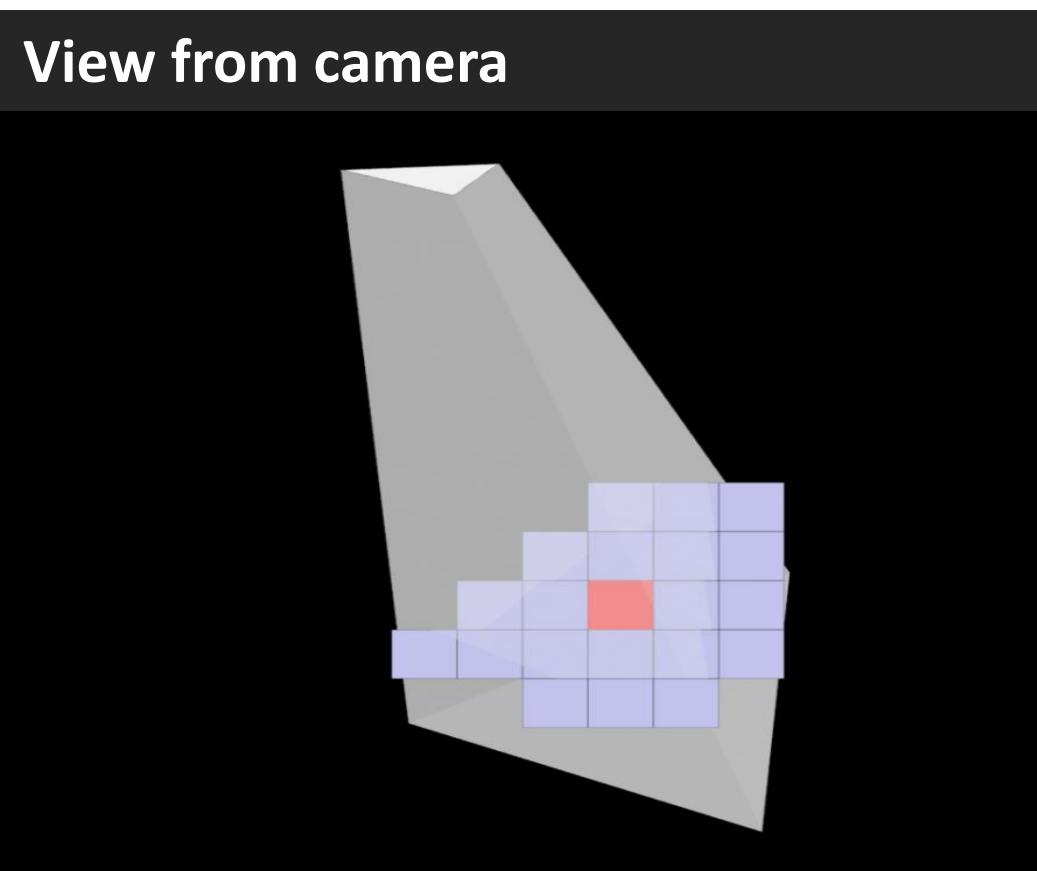
**Alternative view**



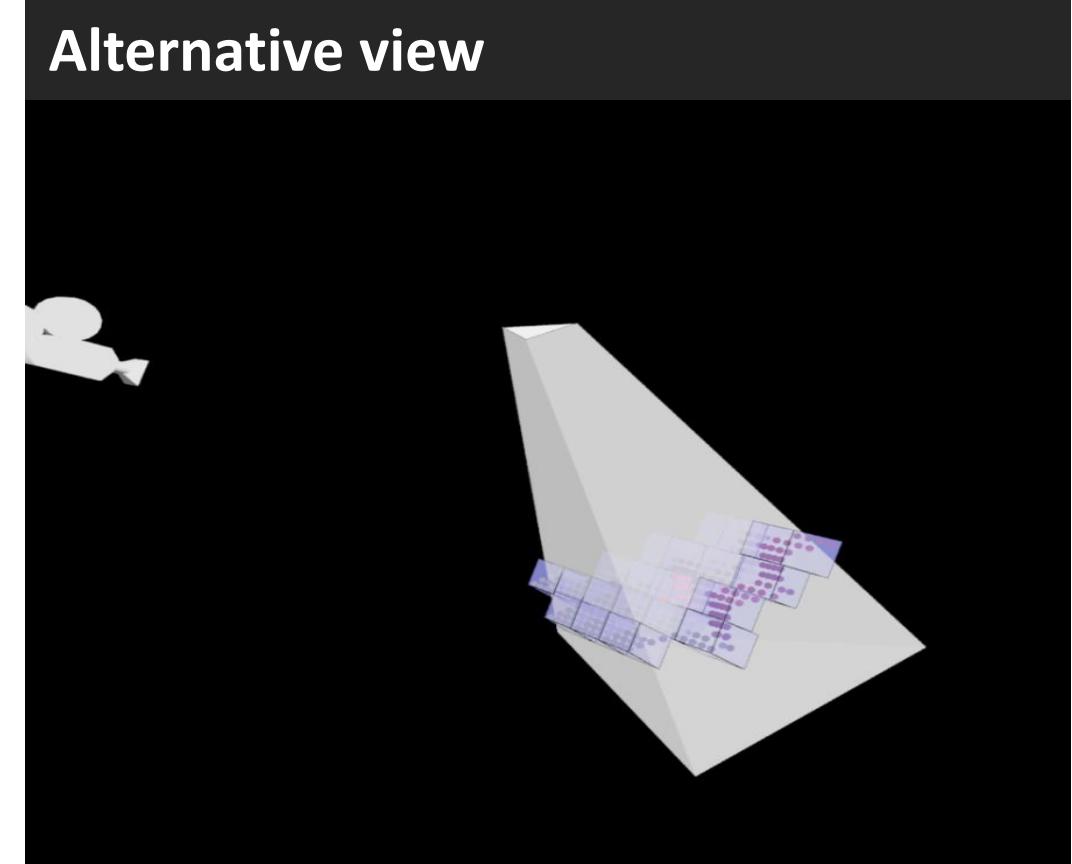
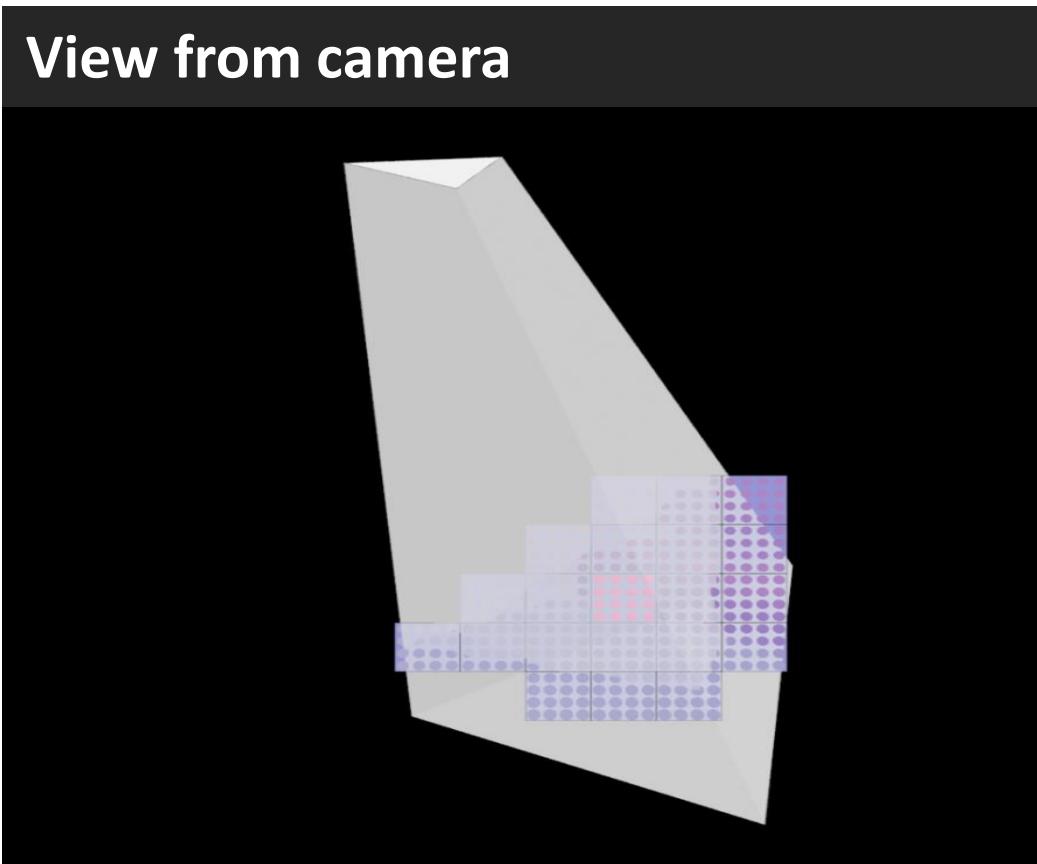
# Per-Triangle Shadow Volumes. Algorithm. Traversing the hierarchy – Cull 4x4 tiles



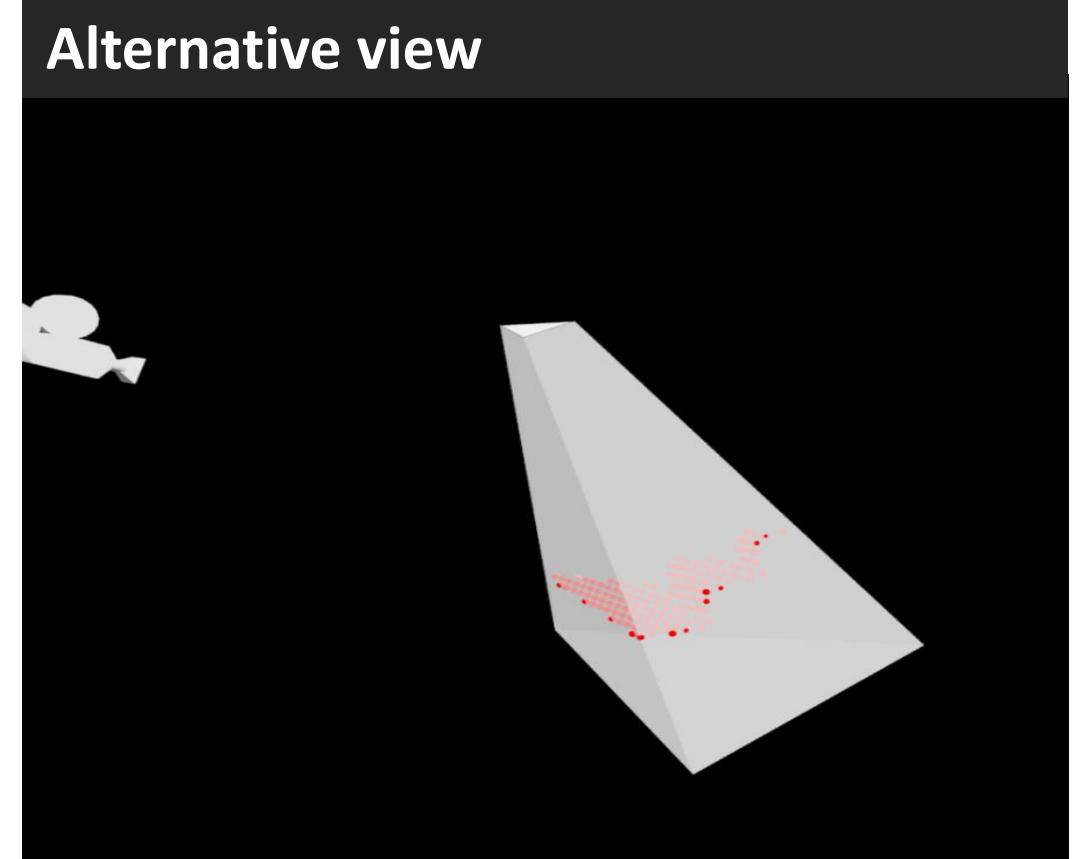
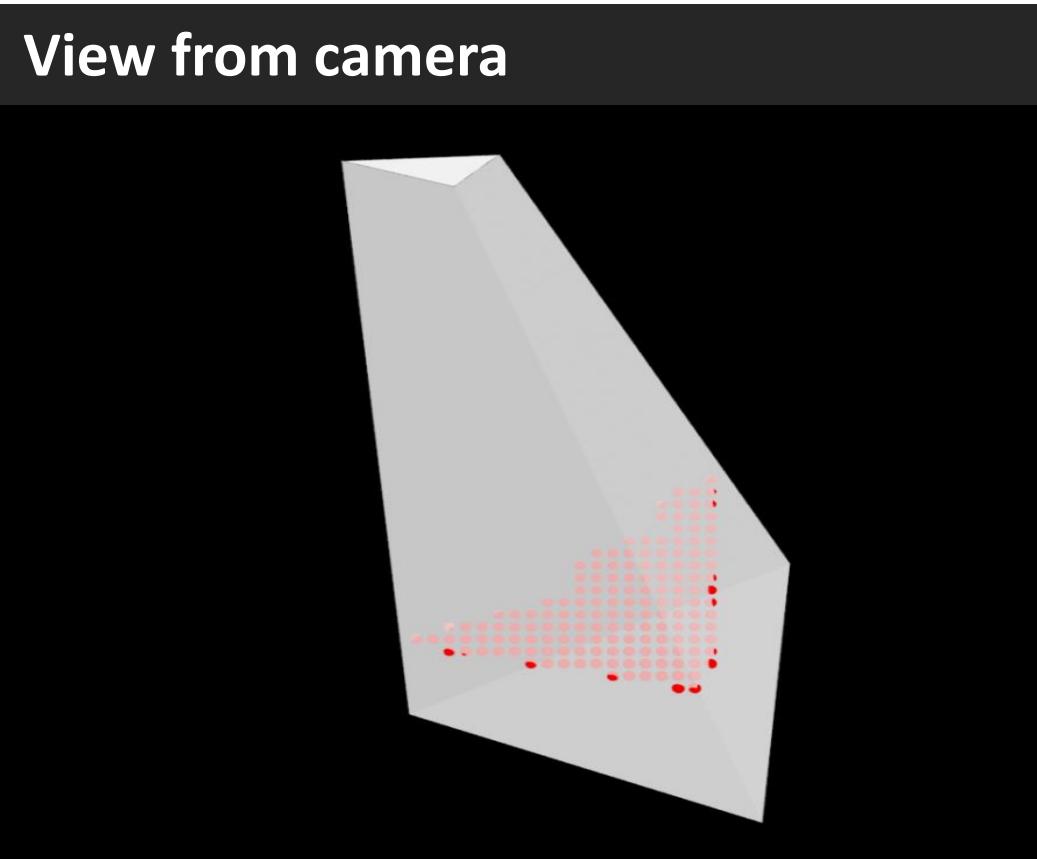
# Per-Triangle Shadow Volumes. Algorithm. Traversing the hierarchy – Trivial Accept



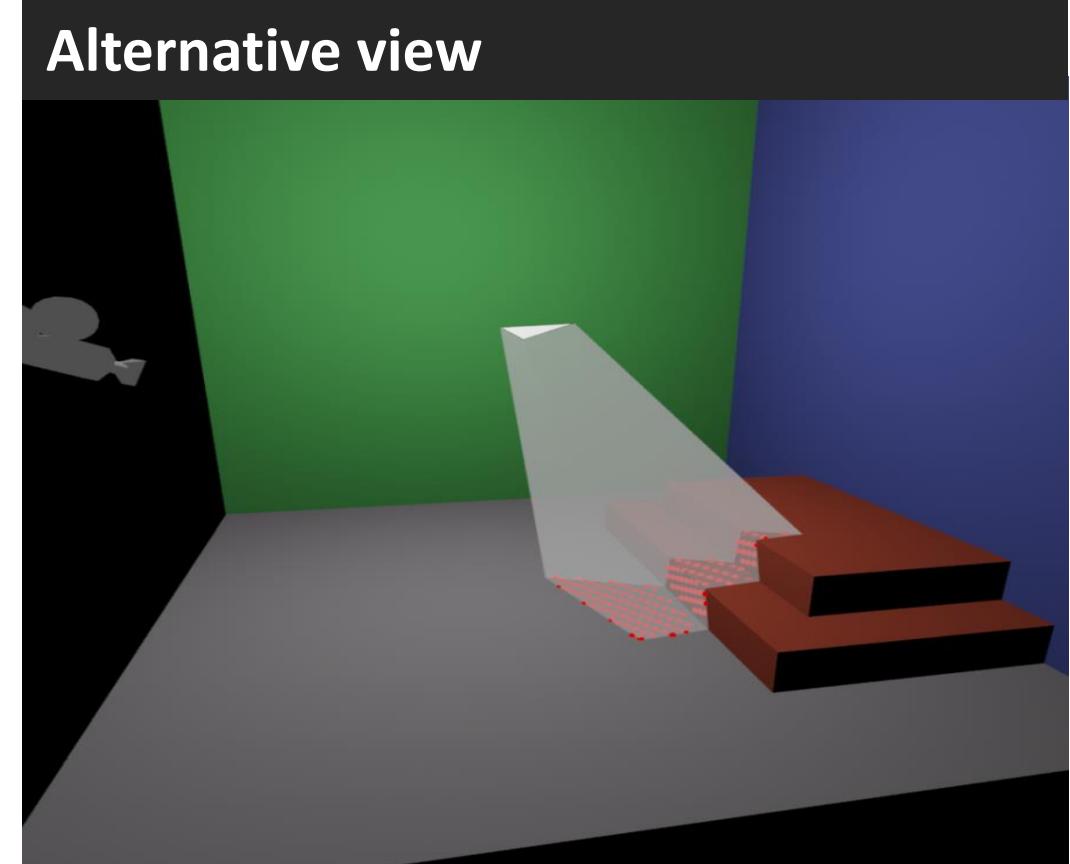
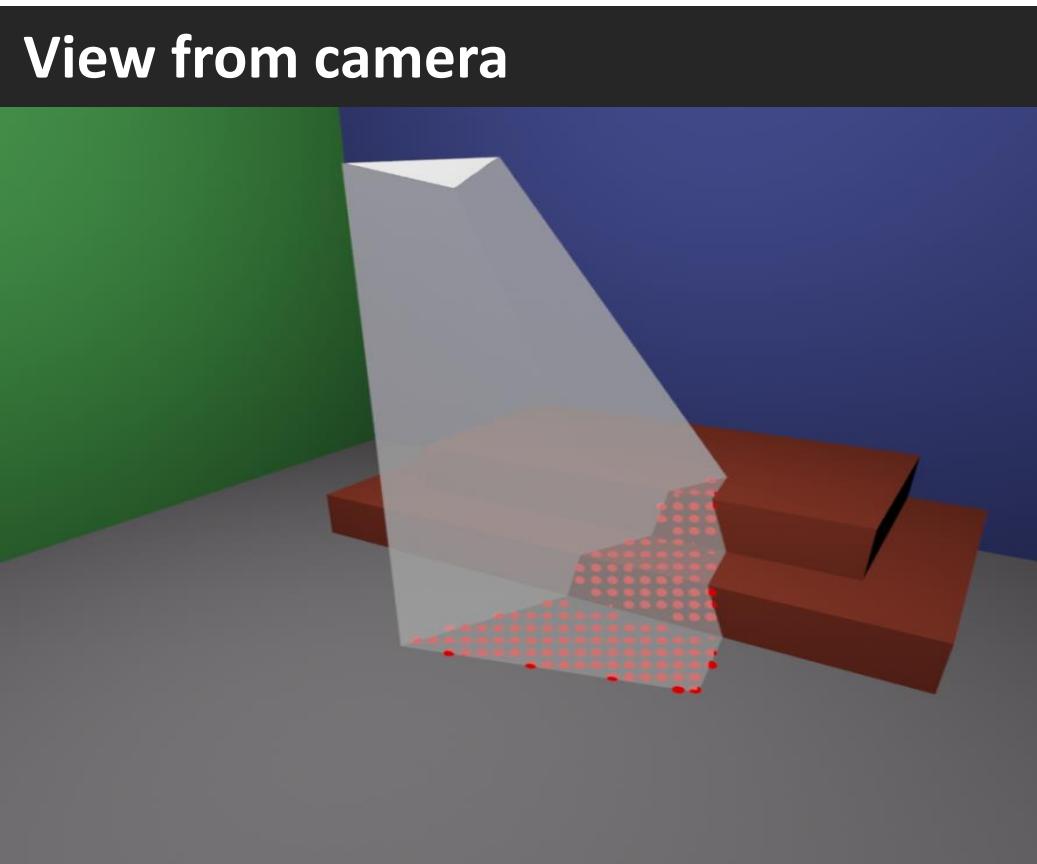
# Per-Triangle Shadow Volumes. Algorithm. Traversing the hierarchy – Trivial Accept



# Per-Triangle Shadow Volumes. Algorithm. Traversing the hierarchy – Test samples



# Per-Triangle Shadow Volumes. Algorithm

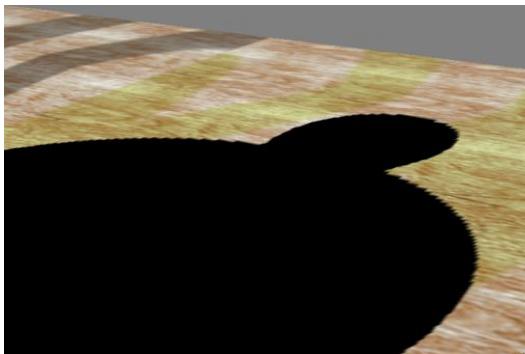


# Shadow Maps vs Shadow Volumes

---

## Shadow Maps

- *Good:* Handles any rasterizable geometry, lookup: **constant cost** regardless of complexity, map can sometimes be reused. **Very fast.**
- *Bad:* Frustum limited. **Jagged shadows** if res too low, **biasing** headaches.
  - Solution: next talk...

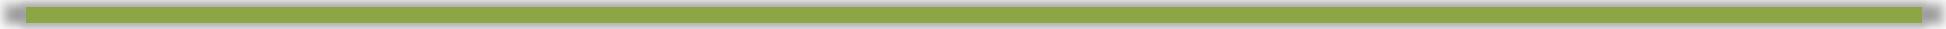


## Shadow Volumes

- *Good:* shadows are **sharp**. Handles omnidirectional lights.
- *Bad:* **3 passes.** Shadow polygons must be generated and rendered → lots of polygons & **fill**
  - Solution: Culling & Clamping or Per-triangle SV using CUDA.



# Thanks



- Used slides from:
  - Ulf Assarsson, Lukai Lan, and others

# Hard shadows

---

Pere-Pau Vázquez  
ViRVIG – UPC