

Pere-Pau Vázquez – ViRVIG Group, UPC

Percentage Closer Filtering and Percentage Closer Soft Shadows

Overview of Shadow Mapping



- Introduced by Williams in 1978
- Advantages compared to shadow volumes:
 - Cost less sensitive to geometric complexity
 - Can be queried at arbitrary locations
 - Often easier to implement
- Disadvantages:
 - Aliasing

Shadow Mapping Algorithm

- Render scene from light's point of view
 - Store depth of each pixel
- When shading a surface:
 - Transform surface point into light coordinates
 - Compare current surface depth to stored depth
 - If $\text{depth} > \text{stored depth}$, the pixel is in shadow; otherwise the pixel is lit


Aliasing Artifacts

- Magnification artifacts




Aliasing Artifacts

- Magnification artifacts
 - Can be addressed to some extent by increasing shadow map resolution. However maximum texture size limits are reached quickly, and are still inadequate (and needlessly expensive) for a large scene.
 - This is a 512x512 shadow map



Aliasing Artifacts

- Minification artifacts



- Typically encountered when viewed from a distance
- Produces ugly and distracting “swimming” effect along shadow edges

Aliasing Artifacts

- Anisotropic artifacts
 - A mix of minification and magnification
 - Encountered at shallow angles



Solutions?

- Also encountered with colour textures
- Reduce aliasing by hardware filtering
 - Magnification artifacts => linear interpolation
 - Minification artifacts => trilinear, mipmapping
 - Anisotropic artifacts => anisotropic filtering

Solutions?

- Can we apply these to shadow maps?
 - Not at the moment
- Interpolating depths is incorrect
 - Gives $\text{depth} < \text{average}(\text{occluder_depth})$
 - Want $\text{average}(\text{depth} < \text{occluder_depth})$

Percentage Closer Filtering

- Proposed by Reeves et al. in 1987
- Filter result of the depth comparison
 - Sample surrounding shadow map pixels
 - Do a depth comparison for each pixel
 - Percentage lit is the percentage of pixels that pass the depth comparison (i.e. are “closer” than the nearest occluder)
- NVIDIA hardware support for bilinear PCF
- Good results, but can be expensive!

Percentage Closer Filtering


- Approaches
 - Percentage-Closer Filtering
 - Jittered Percentage-Closer Filtering

Percentage Closer Filtering

- Percentage-Closer Filtering:
 - Normal shadow maps present aliasing
 - Shadow map textures cannot be prefiltered
 - Filtering can be achieved averaging multiple shadow map comparisons per pixel
 - It calculates the percentage of the surface that is closer to the light:
 - Not in shadow

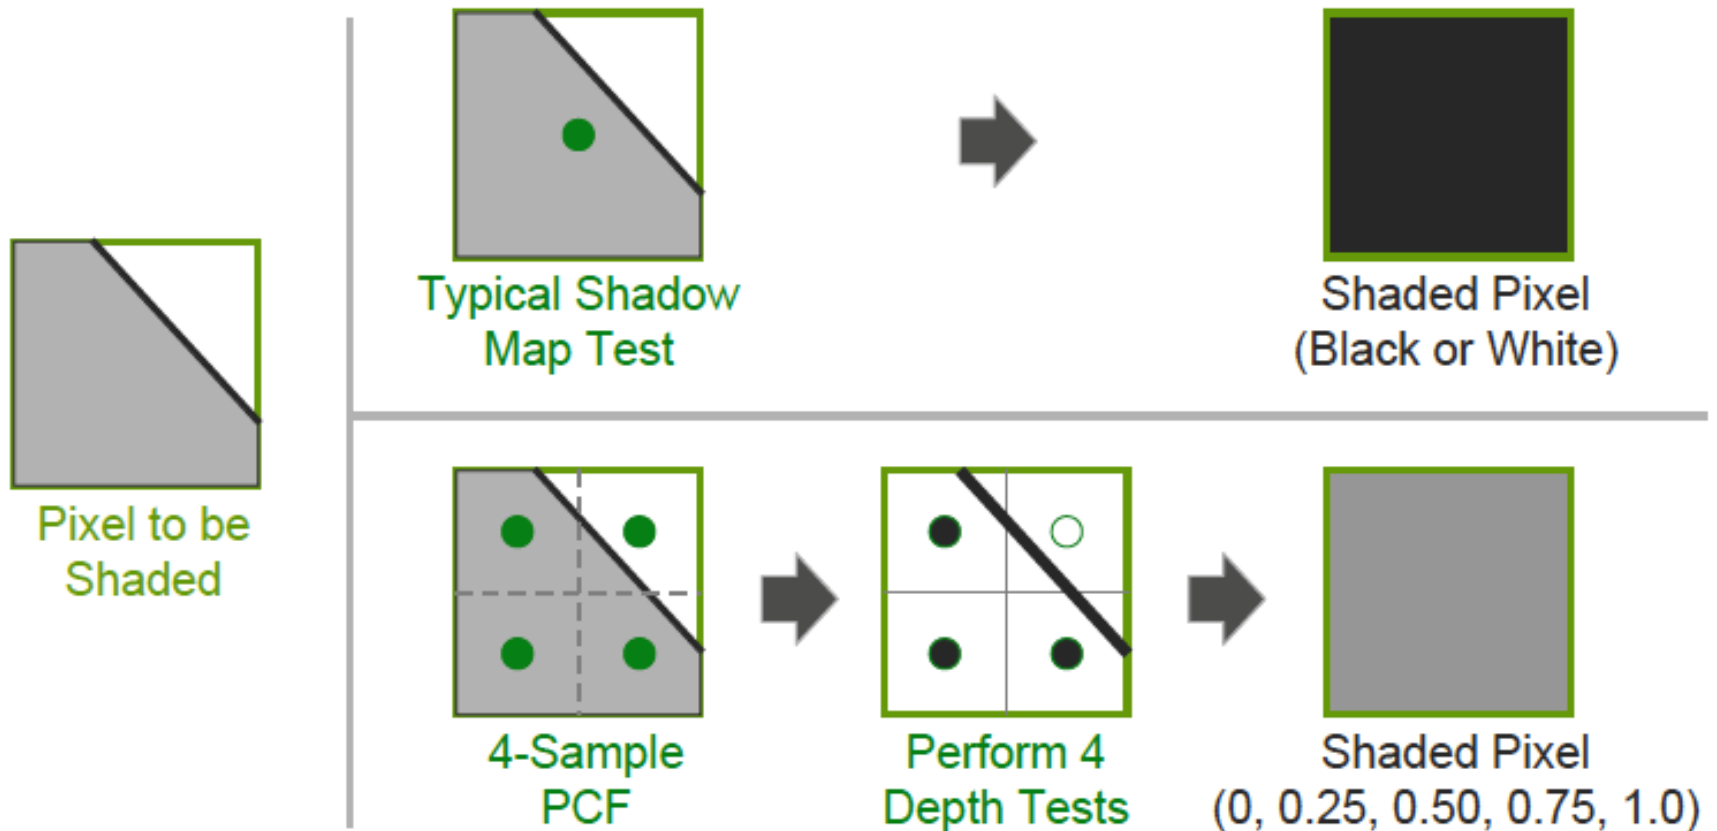
Percentage Closer Filtering

- Percentage-Closer Filtering:
 - Original PCF [Reeves et al. 1987] sampled region to be shaded stochastically (\approx randomly)
 - First implemented using the REYES rendering engine



Percentage Closer Filtering

- Percentage-Closer Filtering:



Percentage Closer Filtering

- Percentage-Closer Filtering implementation:
 - Size of the sampled region passed through an uniform parameter (*fwidth*).
 - The larger *fwidth* the larger the penumbra region (though computed with the same number of samples)
 - Number of samples now is a constant

Percentage Closer Filtering

```
#define SAMPLES_COUNT 32
#define INV_SAMPLES_COUNT (1.0f / SAMPLES_COUNT)

uniform sampler2D decal; // decal texture
uniform sampler2D spot;  // projected spotlight image
uniform sampler2DShadow shadowMap; // shadow map

uniform float fwidth;
uniform vec2 offsets[SAMPLES_COUNT];

// these are passed down from vertex shader
varying vec4 shadowMapPos;
varying vec3 normal;
varying vec2 texCoord;
varying vec3 lightVec;
varying vec3 view;
```

Percentage Closer Filtering

```
#define SAMPLES_COUNT 32
#define INV_SAMPLES_COUNT (1.0f / SAMPLES_COUNT)

uniform sampler2D decal; // decal texture
uniform sampler2D spot; // projected spotlight image
uniform sampler2DShadow shadowMap; // shadow map

uniform float fwidth;
uniform vec2 offsets[SAMPLES_COUNT];

// these are passed down from vertex shader
varying vec4 shadowMapPos;
varying vec3 normal;
varying vec2 texCoord;
varying vec3 lightVec;
varying vec3 view;
```

Percentage Closer Filtering

```
#define SAMPLES_COUNT 32
#define INV_SAMPLES_COUNT (1.0f / SAMPLES_COUNT)

uniform sampler2D decal; // decal texture
uniform sampler2D spot;  // projected spotlight image
uniform sampler2DShadow shadowMap; // shadow map

uniform float fwidth;
uniform vec2 offsets[SAMPLES_COUNT];

// these are passed down from vertex shader
varying vec4 shadowMapPos;
varying vec3 normal;
varying vec2 texCoord;
varying vec3 lightVec;
varying vec3 view;
```

Percentage Closer Filtering

```
void main(void)
{
    float shadow = 0;
    float fsize = shadowMapPos.w * fwidth,
    vec4 smCoord = shadowMapPos;

    for (int i = 0; i < SAMPLES_COUNT; i++) {
        smCoord.xy = offsets[i] * fsize + shadowMapPos;
        shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;
    }
}
```

```
vec3 N = normalize(normal);
```

Percentage Closer Filtering

```
void main(void)
{
    float shadow = 0;
    float fsize = shadowMapPos.w * fwidth;
    vec4 smCoord = shadowMapPos;

    for (int i = 0; i < SAMPLES_COUNT; i++) {
        smCoord.xy = offsets[i] * fsize + shadowMapPos;
        shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;
    }

    vec3 N = normalize(normal);
```

Percentage Closer Filtering

```
void main(void)
{
    float shadow = 0;
    float fsize = shadowMapRes.x * fwidth;
    vec4 smCoord = shadowMapPos;

    for (int i = 0; i < SAMPLES_COUNT; i++) {
        smCoord.xy = offsets[i] * fsize + shadowMapPos;
        shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;
    }
}
```

```
vec3 N = normalize(normal);
```

Percentage Closer Filtering

```
void main(void)
{
    float shadow = 0;
    float fsize = shadowMapPos.w * fwidth;
    vec4 smCoord = shadowMapPos;

    for (int i = 0; i < SAMPLES_COUNT; i++) {
        smCoord.xy = offsets[i] * fsize + shadowMapPos,
        shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;
    }
}
```

```
vec3 N = normalize(normal);
```

Percentage Closer Filtering

```
void main(void)
{
    float shadow = 0;
    float fsize = shadowMapPos.w * fwidth;
    vec4 smCoord = shadowMapPos;

    for (int i = 0; i < SAMPLES_COUNT; i++) {
        smCoord.xy = offsets[i] * fsize + shadowMapPos;
        shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;
    }
}
```

```
vec3 N = normalize(normal);
```


Percentage Closer Filtering

```
for (int i = 0; i < SAMPLES_COUNT; i++) {  
    smCoord.xy = offsets[i] * fsize + shadowMapPos;  
    shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;  
}
```

```
vec3 N = normalize(normal);  
vec3 L = normalize(lightVec);  
vec3 V = normalize(view);  
vec3 R = reflect(-V, N);  
  
// calculate diffuse dot product  
float NdotL = max(dot(N, L), 0);
```

```
// modulate lighting with the computed shadow value  
vec3 color = texture2D(decal, texCoord).xyz;
```

```
gl_FragColor.xyz = (color * NdotL + pow(max(dot(R, L), 0), 64)) *  
    shadow * texture2DProj(spot, shadowMapPos) +  
    color * 0.1;
```

```
}
```

Percentage Closer Filtering

```
for (int i = 0; i < SAMPLES_COUNT; i++) {  
    smCoord.xy = offsets[i] * fsize + shadowMapPos;  
    shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;  
}
```

```
vec3 N = normalize(normal);  
vec3 L = normalize(lightVec);  
vec3 V = normalize(view);  
vec3 R = reflect(-V, N);
```

```
// calculate diffuse dot product  
float NdotL = max(dot(N, L), 0);
```

```
// modulate lighting with the computed shadow value  
vec3 color = texture2D(decals, texCoord).xyz;
```

```
gl_FragColor.xyz = (color * NdotL + pow(max(dot(R, L), 0), 64)) *  
    shadow * texture2DProj(spot, shadowMapPos) +  
    color * 0.1;
```

```
}
```

Percentage Closer Filtering

```
for (int i = 0; i < SAMPLES_COUNT; i++) {  
    smCoord.xy = offsets[i] * fsize + shadowMapPos;  
    shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;  
}
```

```
vec3 N = normalize(normal);  
vec3 L = normalize(lightVec);  
vec3 V = normalize(view);  
vec3 R = reflect(-V, N);
```

```
// calculate diffuse dot product  
float NdotL = max(dot(N, L), 0);
```

```
// modulate lighting with the computed shadow value  
vec3 color = texture2D(decals, texCoord).xyz;
```

```
gl_FragColor.xyz = (color * NdotL + pow(max(dot(R, L), 0), 64)) *  
    shadow * texture2DProj(spot, shadowMapPos) +  
    color * 0.1;
```

```
}
```

Percentage Closer Filtering

- Results:




Percentage Closer Filtering

- Percentage-Closer Filtering.
 - Fast due to the spatial coherence of texture texels
 - Many queries benefit from texture cache
 - Modern GPUs inherently implement a 2x2 PCF for shadow map queries
 - Improvement in shadow mapping-based silhouettes



Percentage Closer Filtering

- Percentage-Closer Filtering. Modified region:





Percentage Closer Filtering

- Recap: hard shadows vs soft shadows



Percentage Closer Filtering

- Recap: hard shadows vs soft shadows




Percentage Closer Filtering

- Recap: hard shadows vs soft shadows
 - Shadow hardening on contact




Percentage Closer Filtering

- Recap: soft shadows is a point-region visibility calculation
 - For each receiver sample (point)
 - determine visible fraction of light source(region)



Percentage Closer Filtering

- Original PCF generates a soft-shadow-like appearance
 - But ignoring penumbra width



Percentage Closer Filtering

- PCF implementation shown adapts penumbra width through the use of *fwidth* parameter

```
void main(void)
{
    float shadow = 0;
    float fsize = shadowMapPos.w * fwidth;
    vec4 smCoord = shadowMapPos;


    for (int i = 0; i < SAMPLES_COUNT; i++) {
        smCoord.xy = offsets[i] * fsize + shadowMapPos;
        shadow += texture2DProj(shadowMap, smCoord) * INV_SAMPLES_COUNT;
    }
}
```

Percentage Closer Filtering

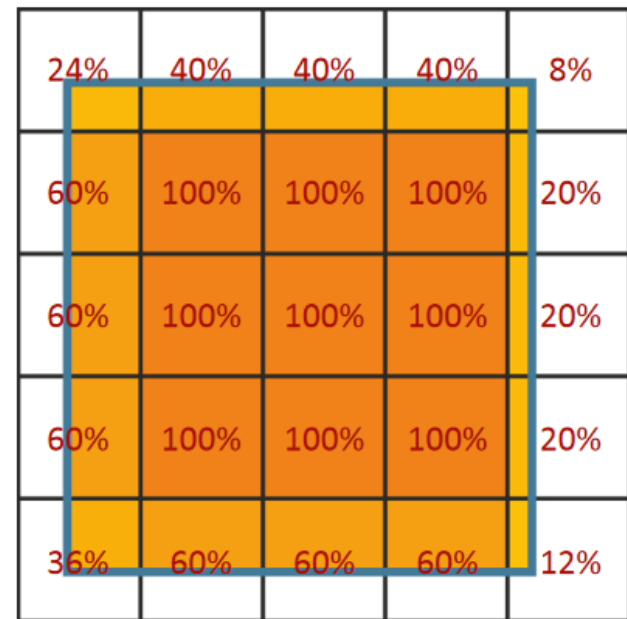
- PCF + large kernel + width dependency = Percentage Closer Soft Shadows (PCSS)

Percentage Closer Filtering

- PCF issues: Oversimplified sampling



PCF: samples weighted equally

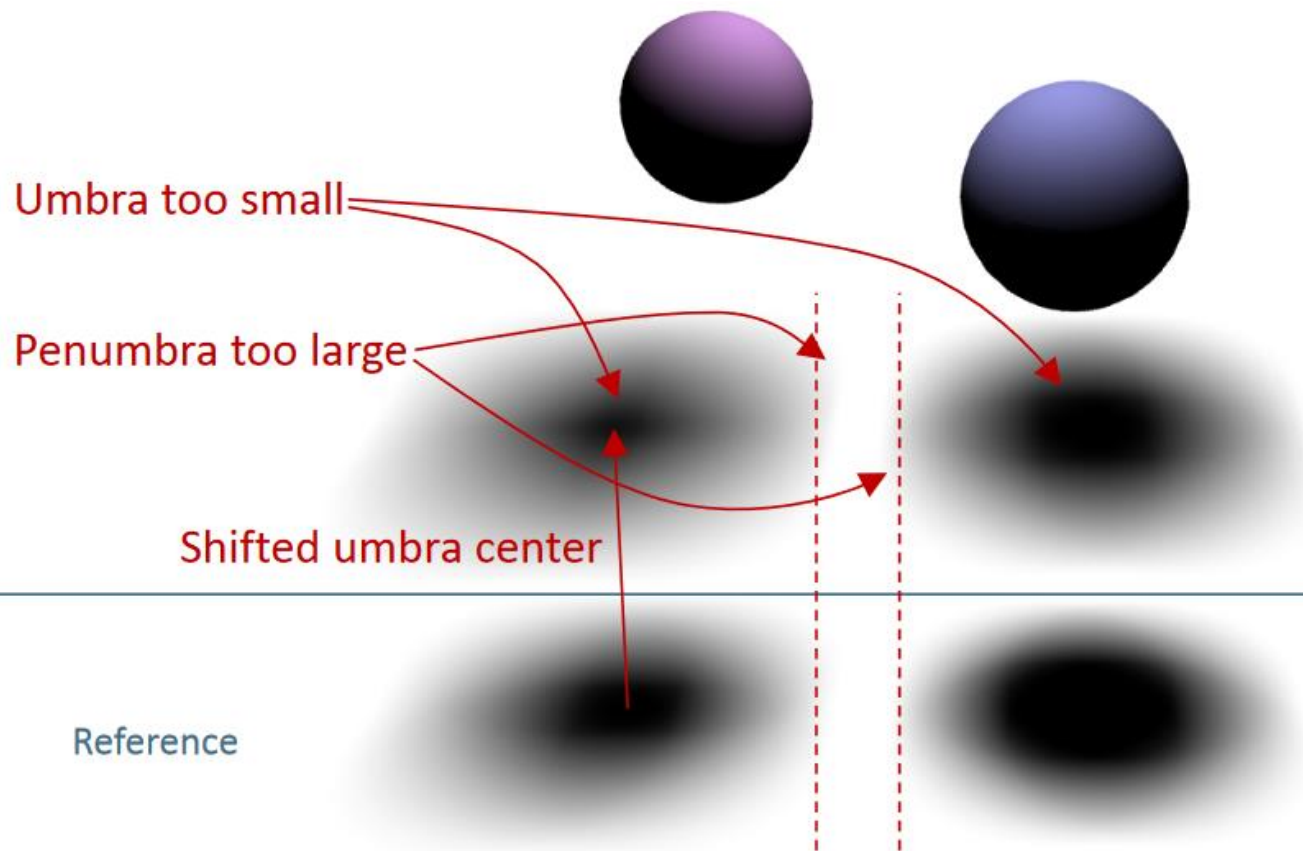


Analytic: actual coverage

[Shen et al., 2011]

Percentage Closer Filtering

■ PCSS issues




Main sources of incorrectness

- Single planar occluder assumption
- Classification as light blocking solely based on depth test

Percentage Closer Filtering

■ PCSS issues



Main sources of incorrectness

- Single planar occluder assumption
- Classification as light blocking solely based on depth test

Percentage Closer Filtering

- PCSS conclusions:
 - Simple and reasonably fast
 - Often visually pleasing results (for smaller light sources)
 - Not really physically correct
 - Only accounts for occluders visible from light source's center

Percentage Closer Filtering

- **Jittered Percentage-Closer Filtering. Motivation:**
 - Though PCF achieves soft shadows, banding artifacts still present
 - An arbitrarily large kernel will soften the shadow even further
 - Might require a high number of samples
 - Still might present banding artifacts due to regular sampling

Percentage Closer Filtering

- Jittered Percentage-Closer Filtering:
 - Sample larger regions in an stochastic manner
 - Use fewer samples taking advantage of:
 - PCF hardware
 - Stochastically placing the samples (i. e. dependent on the image position) reduces banding artifacts
 - If texture is magnified the sampling regions of neighboring pixels will overlap:
 - Smooth transition from unlit to lit
 - No banding artifacts

Percentage Closer Filtering


- Jittered Percentage-Closer Filtering. Implementation:
 - Jitter map passed as a 3D texture

```
vec4 smCoord = ShadowCoord;
vec3 jcoord = vec3 (gl_FragCoord.xy*jxyscale, 0.);
float fsize = ShadowCoord.w * fwidth;

for (int i = 0; i < 4; i++)
{
    vec4 offset = texture3D(jitterMap, jcoord)*2. -1.;
    jcoord.z+=1.0f/SAMPLES_COUNT_DIV_2; //0.03125
    smCoord.xy = offset.xy * fsize + ShadowCoord.xy;
    shadow+= lookup_shadowMap(smCoord);
    smCoord.xy = offset.zw * fsize + ShadowCoord.xy;
    shadow+= lookup_shadowMap(smCoord);
}
```

Percentage Closer Filtering

- Jittered Percentage-Closer Filtering. Results:



Percentage Closer Filtering

- Jittered Percentage-Closer Filtering. PCF vs JPCF: Note the banding artifacts in PCF

PCF

vs

JPCF



Percentage Closer Filtering

- Sample the result of $(d < z)$ around projected point
 - Filter the binary results in a given kernel

0	1
1	0

- Bilinear PCF
 - NVIDIA and AMD GPUs implement 2x2 PCF in one fetch
 - Using the same location and weights for bilinear filtering

Percentage Closer Filtering

```
Texture2D<float> tDepthMap;
```

```
SamplerComparisonState ShadowSampler
```

```
{
```

```
    ComparisonFunc = LESS;
```

```
    Filter = COMPARISON_MIN_MAG_LINEAR_MIP_POINT;
```

```
};
```

```
// ...
```

```
sum += tDepthMap.SampleCmpLevelZero(ShadowSampler,  
                                     uv + offset, z);
```


Percentage Closer Filtering

- Texture configuration

```
// Create the FBO
glGenFramebuffers(1, &m_fbo);

// Create the depth buffer
glGenTextures(1, &m_shadowMap);
glBindTexture(GL_TEXTURE_2D, m_shadowMap);
glTexImage2D(GL_TEXTURE_2D, 0, GL_DEPTH_COMPONENT32, WindowWidth, WindowHeight, 0, GL_DEPTH_COMPONENT, GL_FLOAT, NULL);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, GL_LINEAR);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, GL_LINEAR);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_COMPARE_MODE, GL_COMPARE_REF_TO_TEXTURE);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_COMPARE_FUNC, GL_LEQUAL);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_S, GL_CLAMP_TO_EDGE);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_T, GL_CLAMP_TO_EDGE);

glBindFramebuffer(GL_FRAMEBUFFER, m_fbo);
glFramebufferTexture2D(GL_FRAMEBUFFER, GL_DEPTH_ATTACHMENT, GL_TEXTURE_2D, m_shadowMap, 0);
```

Percentage Closer Filtering

- Increasing the number of PCF taps increases the softness of the shadows



1 tap



9x9 taps



17x17 taps

Percentage Closer Filtering

- PCF with large kernels requires many samples
 - Using irregular sampling
 - Trades banding for noise




regular sampling



irregular sampling




Percentage Closer Filtering. Self-shadowing issue

- Traditional depth bias: write $(w + \text{bias})$ in shadow map
 - Bias = constant bias + slope-based bias
 - Issue: huge depth biases may be required for large PCF kernels




Percentage Closer Filtering

- Use depth gradient = $\text{float2}(dz/du, dz/dv)$
 - Make depth d follow tangent plane
 - $d = d_0 + \text{dot}(\text{uv_offset}, \text{gradient})$
 - [Schuler06] and [Isidoro 06]
- Render midpoints into shadow map
 - Midpoint $z = (z_0 + z_1)/2$
 - Requires two rasterization passes
 - Depth peel two depth layers
 - Still requires a depth bias for thin objects
- Render back faces into shadow map
 - Only works for closed objects
 - Light bleeding for large PCF kernels



Percentage Closer Filtering

- Rendering back faces into shadow map generates light bleeding for large PCF kernels
 - Not due to FP precision or shadow map resolution
 - But reverse of the surface acne issue



Percentage Closer Filtering

- PCF cannot be prefiltered as is
 - $\text{Average}(d < z) \neq (\text{Average}(z) < d)$
 - Filtering the depth buffer would smooth the heightfield of the shadow map
 - Does not generate soft shadows
 - May introduce artifacts
- Solutions: Approximate shadow test by a linear function which can be prefiltered
 - Goal: blurring the shadow map to generate realistic soft shadows

Fast Percentage Closer Soft Shadows using Temporal Coherence

Michael Schwärzler*
VRVis Research Center,
Austria

Christian Luksch†
VRVis Research Center,
Austria

Daniel Scherzer‡
Max-Planck-Institut für
Informatik, Germany

Michael Wimmer§
Vienna University of
Technology, Austria

Outline

- Related work
 - Real-time soft shadow mapping
 - Data catching / Temporal coherence
- Their approach
 - Shadow reprojection
 - Detecting moving objects
 - Reconstruction error
- Implementation and evaluation
- Conclusion and future work

Related Work


Real-time soft shadow mapping

- ***Percentage Closer Soft Shadows (PCSS)***

[Fernando 2005]

- Based on *Percentage Closer Filtering* [Reeves et al. 1987]


- Average comparison results, not depth values



- But supporting variable kernel sizes
 - To simulate varying penumbra sizes

Percentage Closer Soft Shadows (PCSS)

- PCSS [Fernando05]
- Assume a square light centered at the shadow map center
- Assuming some parallel blocker to receiver
 - Compute penumbra width using similar triangles




$$w_{Penumbra} = \frac{(d_{Receiver} - d_{Blocker}) \cdot w_{Light}}{d_{Blocker}}$$

Percentage Closer Soft Shadows (PCSS)

- **Step 1: Blocker search**
 - Sample the depth buffer using point sampling
 - Average all blockers with $(\text{depth} + \text{bias} < \text{receiver})$ in search region / kernel
 - Early out if no blocker found
- **Step 2: Filtering**
 - Use filter radius from step 1
 - Clamp filter width to be $\geq \text{MinRadius}$ for antialiasing
 - Filter the shadow map with PCF or VSM/CSM/ESM

Percentage Closer Soft Shadows (PCSS)


- Where to find blockers?
- Conservative search radius using similar triangles



$$\text{LightRadius} / d = \text{SearchRadius} / (d - \text{znear})$$

Percentage Closer Soft Shadows (PCSS)


- Why not doing just one sample?



Percentage Closer Soft Shadows (PCSS)

- Step 2: Penumbra estimation
 - assumes that blocker and receiver are planar and in parallel

$$w_{Penumbra} = \frac{(d_{Receiver} - d_{Blocker}) \cdot w_{Light}}{d_{Blocker}}$$



Percentage Closer Soft Shadows (PCSS)

- The more samples in the blocker search, the less noisy artifacts in the soft shadows
 - In practice, 4x4 or 5x5 samples is sufficient



Blocker Search with 3x3 taps



Blocker Search with 5x5 taps

Percentage Closer Soft Shadows (PCSS)

PCSS In Hellgate: London



PCSS




PCF

16 POINT taps for the blocker search
16 PCF taps for the PCF filtering

Percentage Closer Soft Shadows (PCSS)

- Step 3: Percentage-Closer Filtering



Related Work

Data catching / Temporal coherence


- ***Reverse reprojection***

[Nehab et al. 2007], [Scherzer et al. 2007]

- storing per pixel info in a *history buffer*
- compare stored depth with the current one
 - in order to reuse it
 - safely (e.g. not when lights moved)
- calculate information in disoccluded regions or new areas into frustum
- some error may be introduced during reprojection

Shadow reprojection

- *History buffer* stores
 - PCSS map
 - Depth map
 - In a ping-pong style
- Depth test only in
 - Disoccluded regions
 - New border regions




Detecting moving objects

- Two type shadows on moving objects
 - Shadows cast **on** it
 - Detect areas with depth test
 - Shadows cast **by** it
 - *Movement Map*
 - “1” moving object visible
 - “0” otherwise





Detecting moving objects

- Movement Map



Detecting moving objects

- Movement map problem:
 - Only have a single hard shadow




Detecting moving objects

- Solution for the inner penumbra static objects
 - Render moving objects first in both maps
 - Releasing movement map as render target
- Solution for the outer penumbra dynamic objects
 - *Pixel Mipmap* generation procedure to create a pyramid of movement map
 - Mipmap selection equal to occluder search radius PCSS

$$r_{search} = \frac{w_{light} * (z_{receiver} - d_{Nearplane})}{z_{receiver}}$$

Reconstruction error

- With camera movement, reprojecting the history buffer from the previous frame introduces error
 - Bilinear interpolation



Reconstruction error

- Accumulated projection error cannot be too large
 - Avoid “oversmoothing”
 - Depends on the scene parameters
 - Two options:
 - Bicubic texture sampling (Catmull-Rom interpolation)
 - Dividing the screen into a groups in a grid
 - Update periodically
 - Parts with too much error recalculated and blended with the old ones

Implementation and evaluation

- Implementation


- in a C++, with shaders using DirectX10

- Tests


- Intel Core i7 920 with 6GB RAM and GeForce GTX 580
 - Screen resolution 1920x1080
 - Three different scenarios and three different methods

Implementation and evaluation


Static Scene









Fully Dynamic Scene



Mixed Scene (dynamic & static objects)



Implementation and evaluation



Conclusions and Future Work

- New improvement method for PCSS
 - Easy to integrate into an existing rendering framework
 - Can be used for all kinds of different scenes
 - The achievable performance gain comes at the cost of memory consumption
- Not for a moving light sources
- FPS variable: may be a problem
- Future work
 - real-time calculation of physically correct soft shadows in dynamic scenes

Fast Percentage Closer Soft Shadows using Temporal Coherence

Michael Schwärzler*

VRVis Research Center,
Austria

Christian Luksch[†]

VRVis Research Center,
Austria

Daniel Scherzer[‡]

Max-Planck-Institut für
Informatik, Germany

Michael Wimmer[§]

Vienna University of
Technology, Austria

Pere-Pau Vázquez – ViRVIG Group, UPC

Percentage Closer Filtering and Percentage Closer Soft Shadows
